



Machine Learning and Data Mining Models Applied to Social Analysis, Healthcare, Financial Fraud Analysis

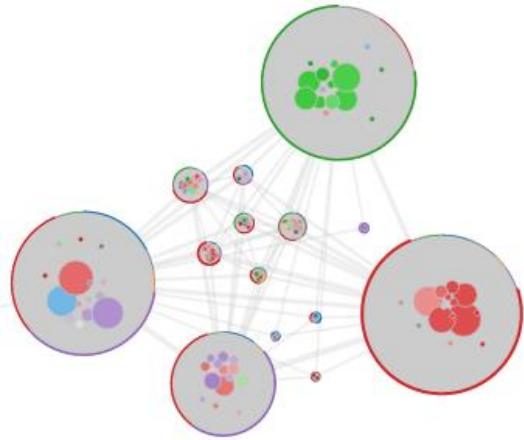
Presenter

Lecheng Zheng

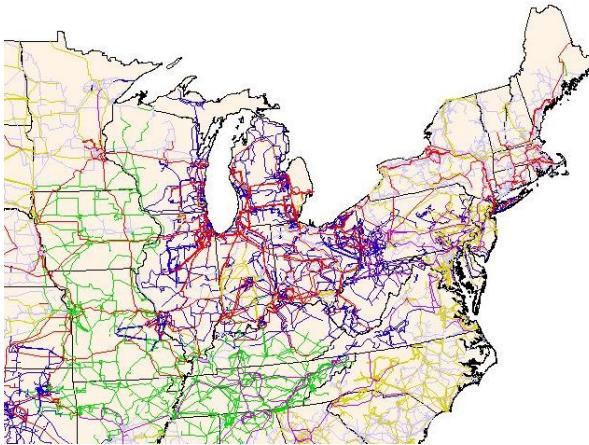
University of Illinois at Urbana-Champaign,

Email: lecheng4@illinois.edu

Observation: Graphs Are Everywhere!



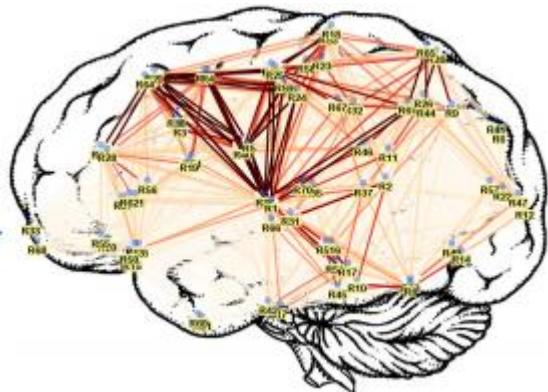
Collaboration Networks



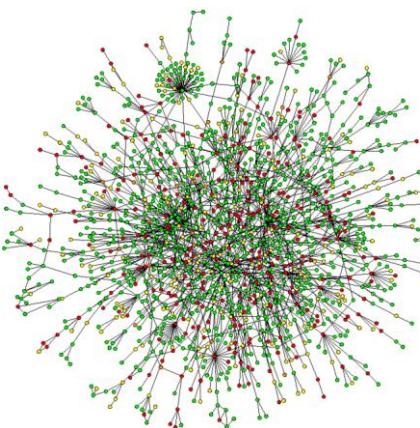
US Power Grid



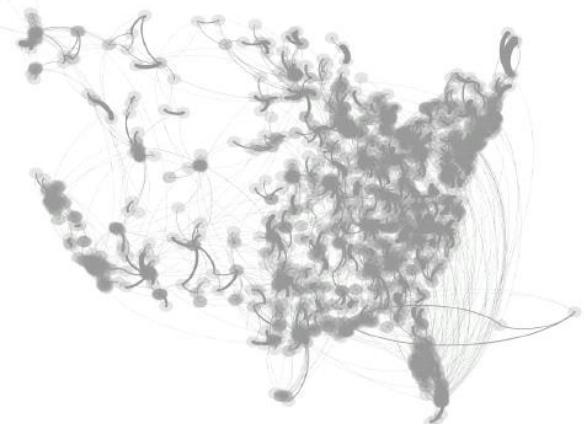
Traffic Network



Brain Networks

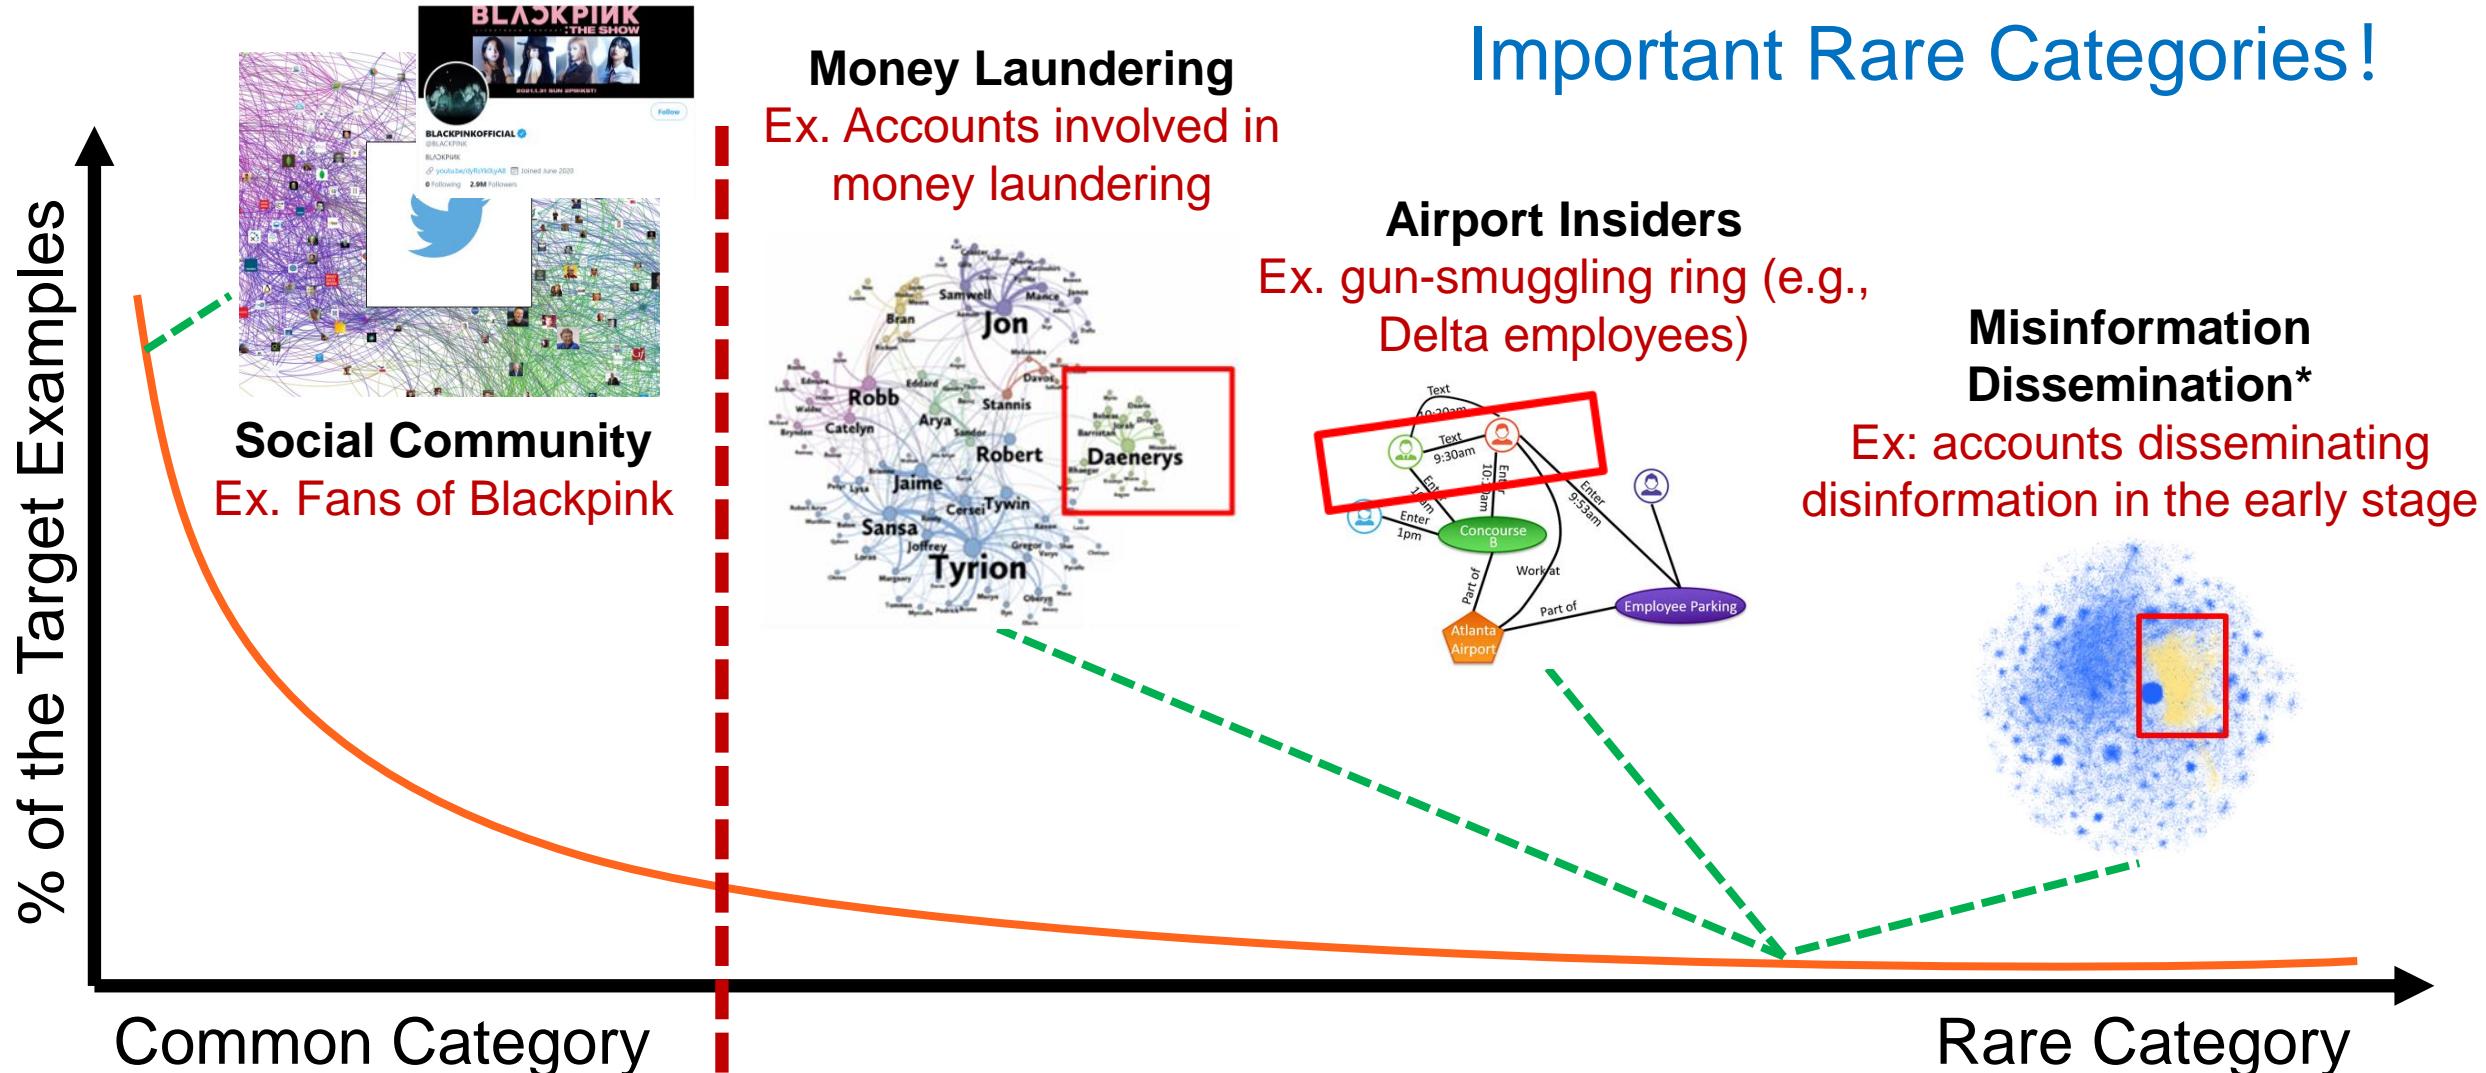


Biological Networks



Hospital Networks

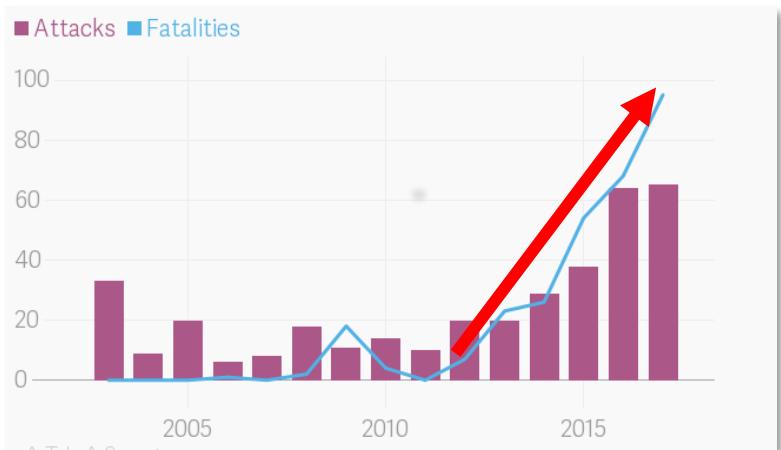
Examples of Rare Categories on Graphs



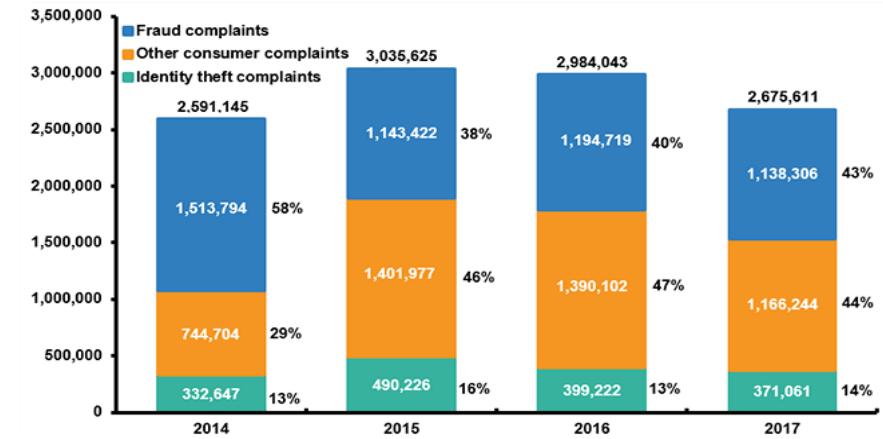
- Wang, Yu-Xiong, et al. "Learning to model the tail." NeurIPS. 2017.
- Laite, Ralph, et al. "Terrorist Group Classification of Historic Terrorist Attacks from The Global Terrorism Database." ICCSE. 2019.
- *Picture from: <https://cnets.indiana.edu/blog/2016/12/21/hoaxy/>

Why Do We Care (Security Domain)?

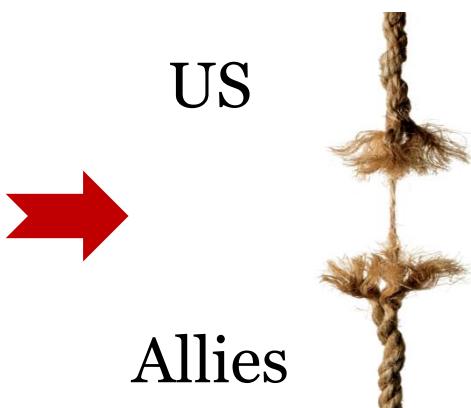
□ Terrorist Events Worldwide



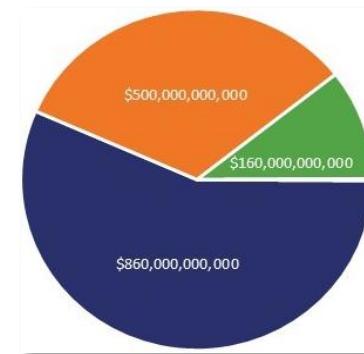
□ Financial Fraud



□ PR Crisis



□ Cybercrime



- Illegal online Markets
- Trade secret, IP theft, data trading

Cost over \$3.5 billion in 2019

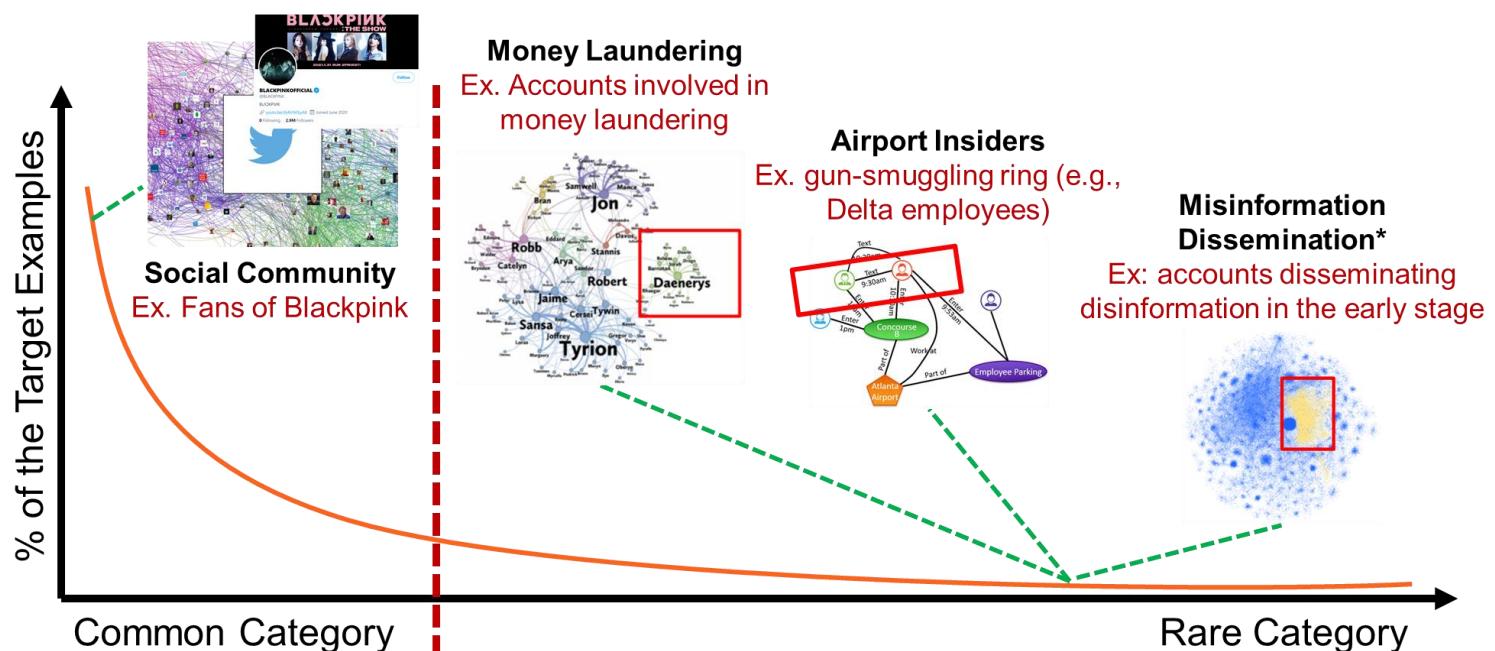
- West, Jarrod, et al. "Intelligent financial fraud detection: a comprehensive review." *Computers & security*. 2016.
- Hasisi, Badi, et al. "Crime and terror: Examining criminal risk factors for terrorist recidivism." *Journal of Quantitative Criminology*. 2020.
- Massie, Justin, et al. "America's Allies and the Decline of US Hegemony." Routledge." 2019.

Rare Category Analysis

□ Problem definition

- Rare category analysis (RCA) refers to the problem of analyzing the underrepresented minority classes in an imbalanced data set.

□ The nature of rare categories



Previous Studies

- Rarity [Lin et al., 2017]
- Non-separability [He et al., 2010]
- Label scarcity [He et al., 2010]

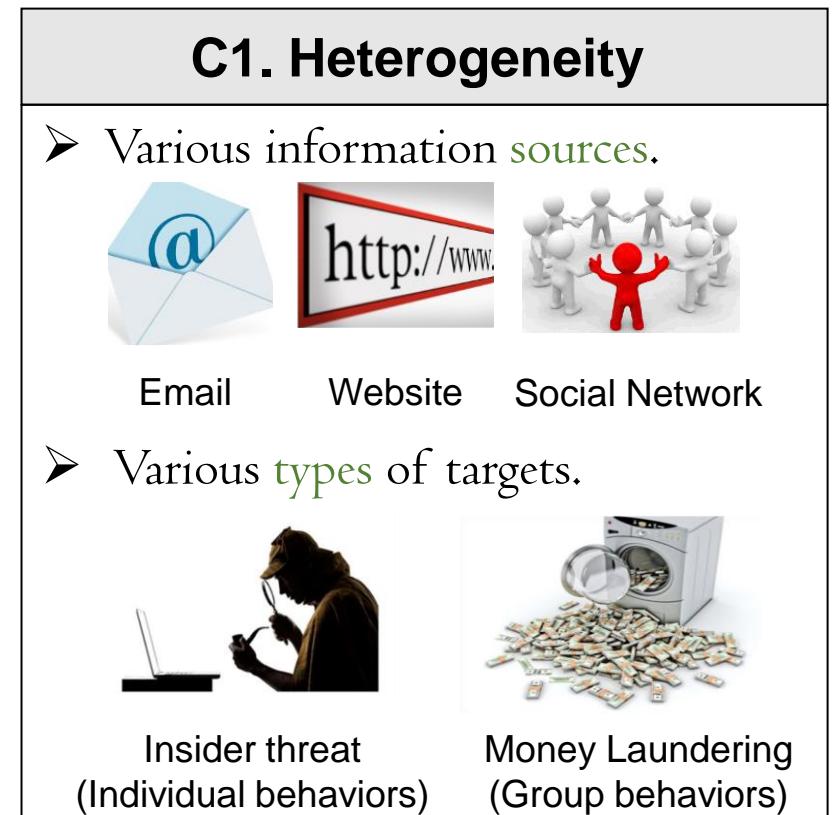
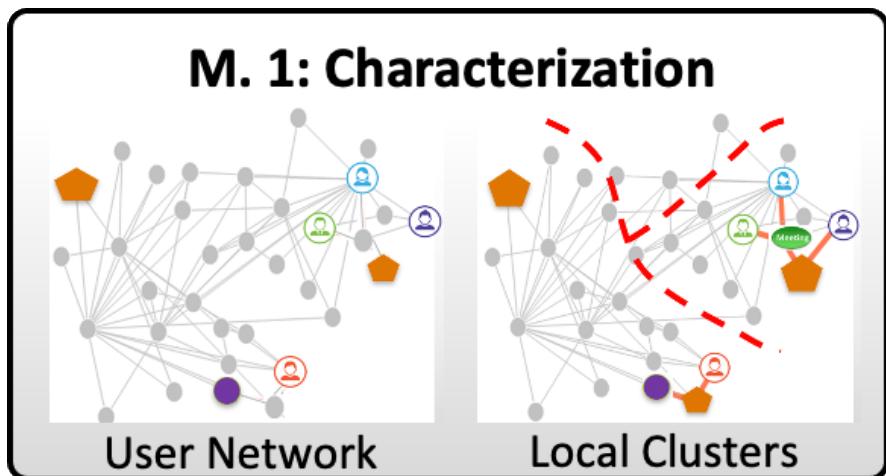
Rare Category Analysis – Challenges

□ Problem definition

- Rare category analysis (RCA) refers to the problem of analyzing the underrepresented minority classes in an imbalanced data set.

□ Emerging challenges in the era of big data

Q1: How to **characterize** rare categories with data and task heterogeneity?



- Zhou, Dawei, et al. "MUVIR: Multi-View Rare Category Detection." IJCAI. 2015.
- Zhou, Dawei, et al. "A local algorithm for structure-preserving graph cut." ACM SIGKDD. 2017.
- Ji, Heng, et al. "Refining event extraction through cross-document inference." ACL. 2008.

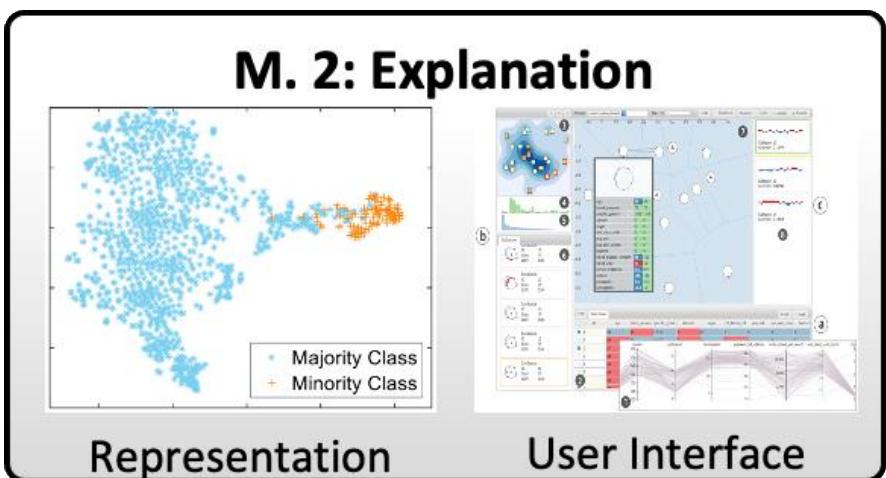
Rare Category Analysis – Challenges

□ Problem definition

- Rare category analysis (RCA) refers to the problem of analyzing the underrepresented minority classes in an imbalanced data set.

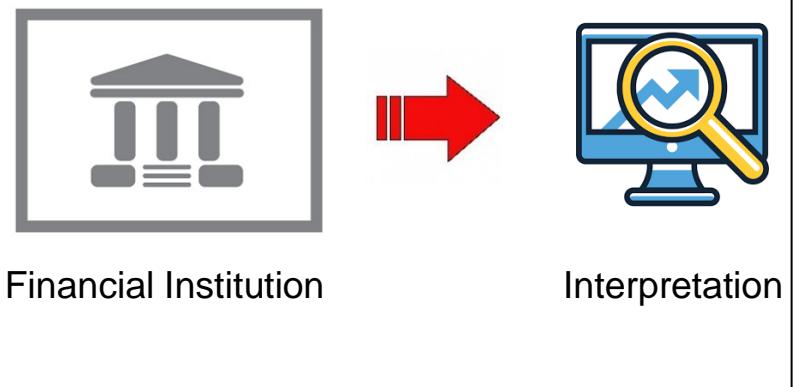
□ Emerging challenges in the era of big data

Q2: How to **interpret** the prediction results and provide **relevant clues** for the end-users?



C2. Interpretability

- Many machine learning models remain “**black-box**” in nature.
- Traditional industries require models to be **interpretable**.



- Zhou, Dawei, et al. "Spac: Self-paced network representation for few-shot rare category characterization." ACM. 2018.
- Pan, Jia-cheng, et al. "RCAnalyzer: visual analytics of rare categories in dynamic networks." Frontiers of Information Technology & Electronic Engineering. 2020.
- Yang, Carl, et al. "Neural Concept Map Generation for Effective Document Classification with Interpretable Structured Summarization." ACM SIGIR. 2020.

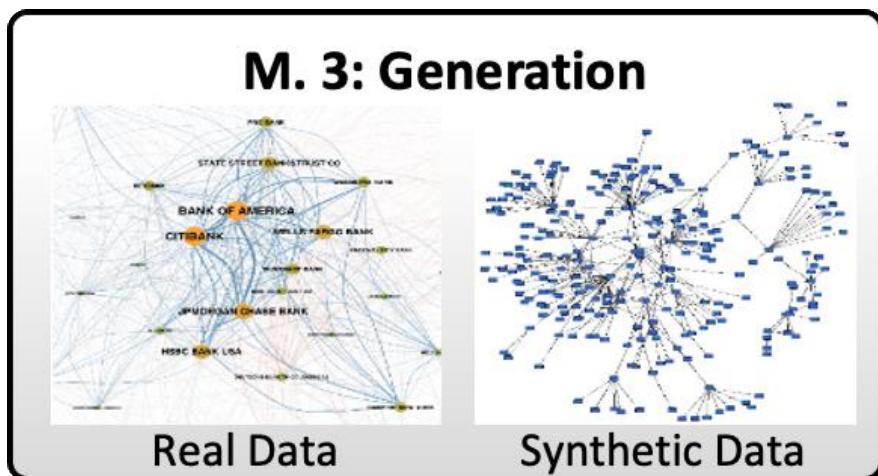
Rare Category Analysis – Challenges

□ Problem definition

- Rare category analysis (RCA) refers to the problem of analyzing the underrepresented minority classes in an imbalanced data set.

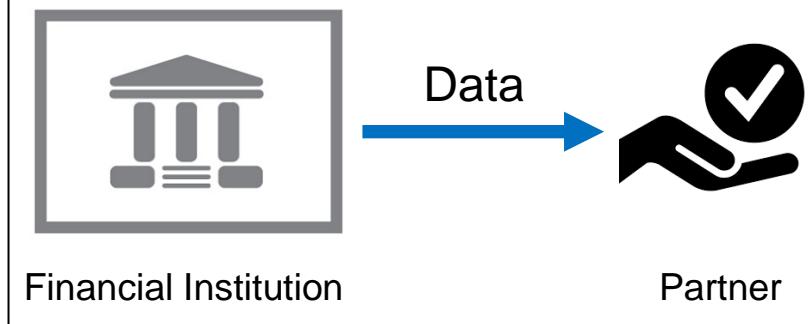
□ Emerging challenges in the era of big data

Q3: How to alleviate the **data scarcity** and enable **data augmentation**?



C3. Data Scarcity

- Massive data containing **sensitive** information (e.g., sex, age).
- Data **security** and **privacy**.
- Data **proprietary**.



• Zhou, Dawei, et al. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." ACM. 2020.

• Zhou, Dawei, et al. "Misc-GAN: A Multi-scale Generative Model for Graphs." Frontiers in Big Data. 2019.

• Akoglu, Leman, and Christos Faloutsos. "RTG: a recursive realistic graph generator using random typing." ECML PKDD. 2009.

Comparison with Imbalanced Classification

□ Imbalanced classification

- Labeled examples from **all the classes**.
- Models for imbalanced classification focus on **overall performance** of all classes.
- **Methodology:** [Kubat & Matwin, ICML1997]; [Chawla et al, JAIR2002]; [Wu & Chang, ICML2003]; [Huang et al., CVPR2016]; ...

□ Rare category analysis

- **None/one/few-shot** labeled example from the rare categories.
- Models for rare category analysis put heavy emphasis on **learning minority classes** with a good performance.
- **Methodology:** [Fine & Mansour, COLT2006]; [Dasgupta & Hsu, ICML2008]; [Vatturi & Wong, KDD2009]; [Zhou et al., KDD2018]; ...

• Branco, Paula, et al. "A survey of predictive modeling on imbalanced domains." ACM Computing Surveys. 2016.

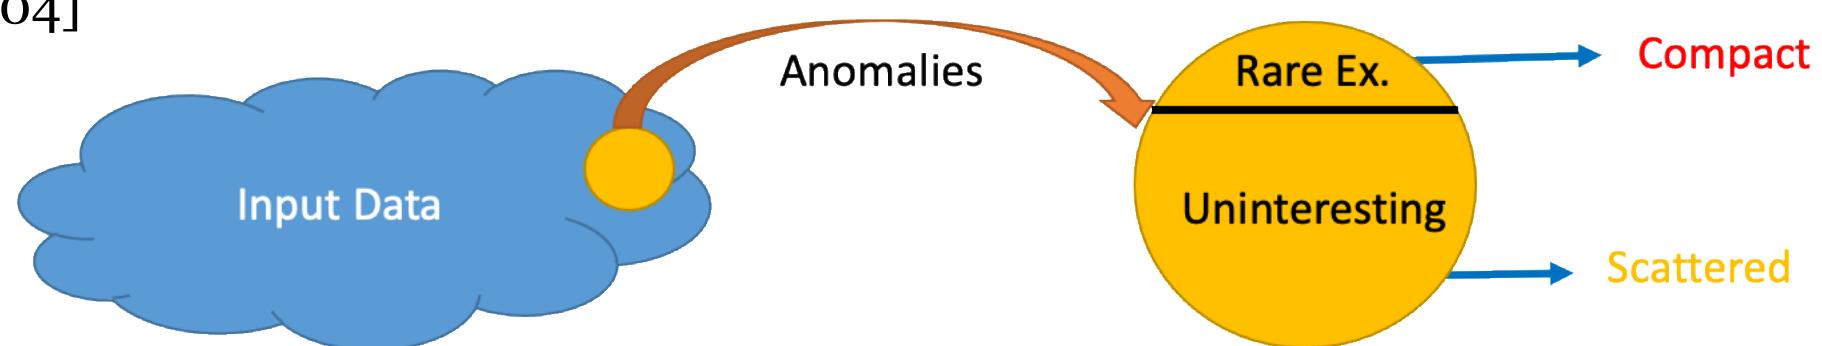
Comparison with Outlier/Anomaly Detection

□ Outliers/anomalies

- Patterns in data that do not conform to expected behavior [Chandola, Banerjee & Kumar, 2009].
- Typically separable from normal examples and **scattered** in the feature space.

□ Rare categories

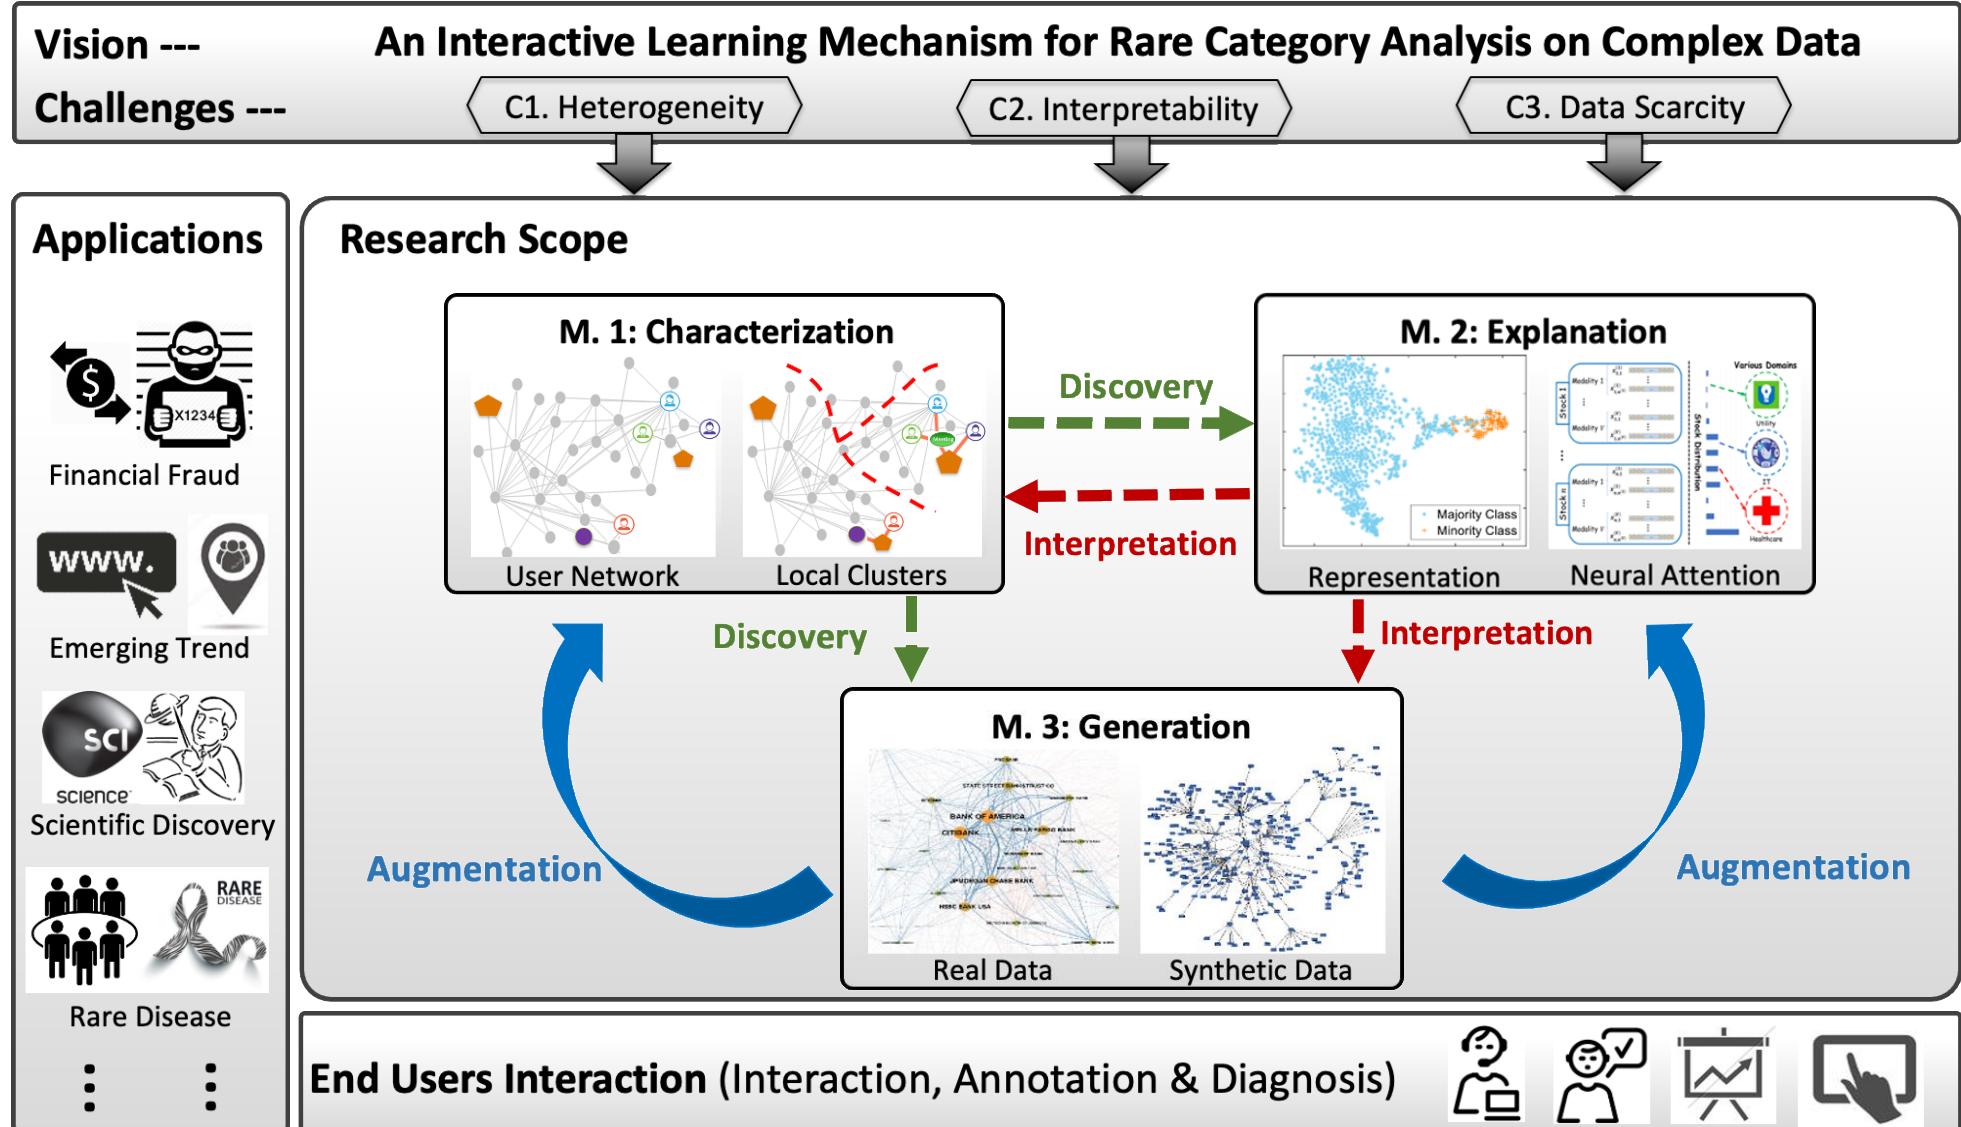
- **Compact** in the feature space.
- “Most of the objects (i.e., **majority classes**) are well explained by current theories and remainder are anomalies, but 99% of these anomalies are uninteresting (i.e., **boring anomalies**), and only 1% of them (i.e., **rare categories**) are useful” [Pelleg & Moore, 2004]



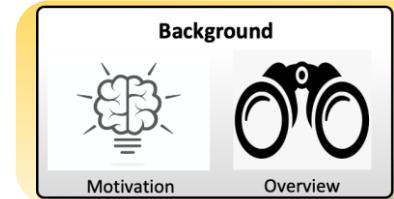
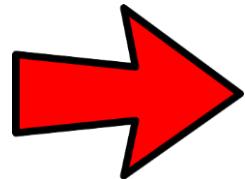
- Gupta, Manish, et al. "Outlier detection for temporal data: A survey." IEEE Transactions on Knowledge and data Engineering, 2013.
- Akoglu, Leman, et al. "Graph based anomaly detection and description: a survey." Data mining and knowledge discovery. 2015.



Overview of Rare Category Trinity

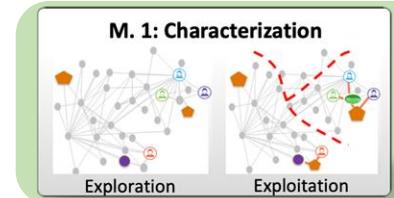


Roadmap



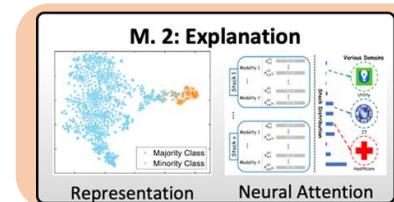
Background

- Motivation
- Research overview



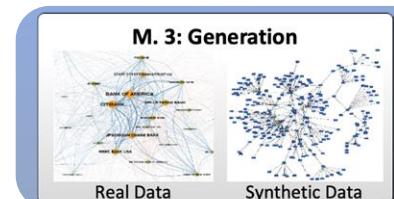
I. Rare Category Characterization

- Rare category characterization on homogeneous graphs
- Rare category characterization on heterogeneous graphs



II. Rare Category Explanation

- Data insights
- Model insights



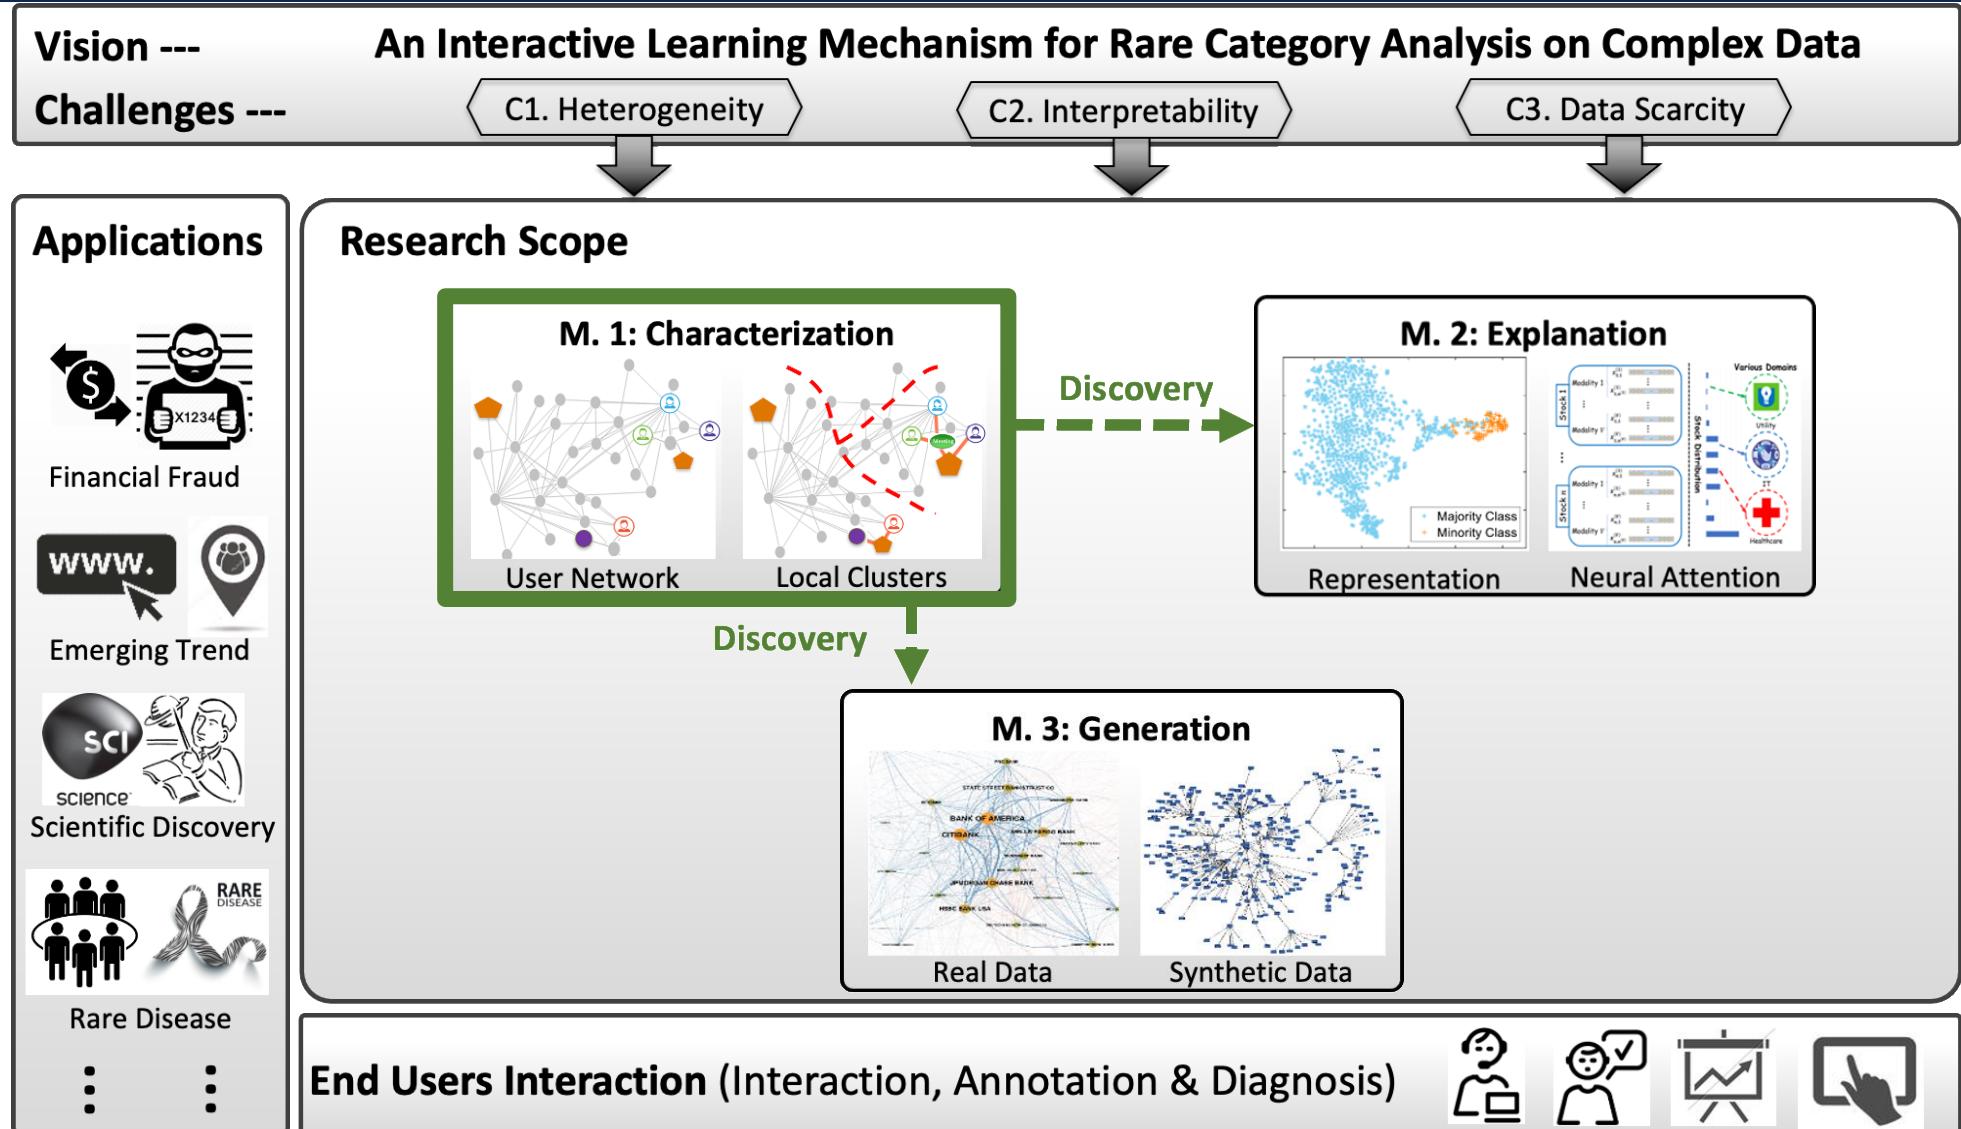
III. Rare Category Generation

- Unsupervised rare category generation
- Supervised rare category generation



IV. Real-world Application

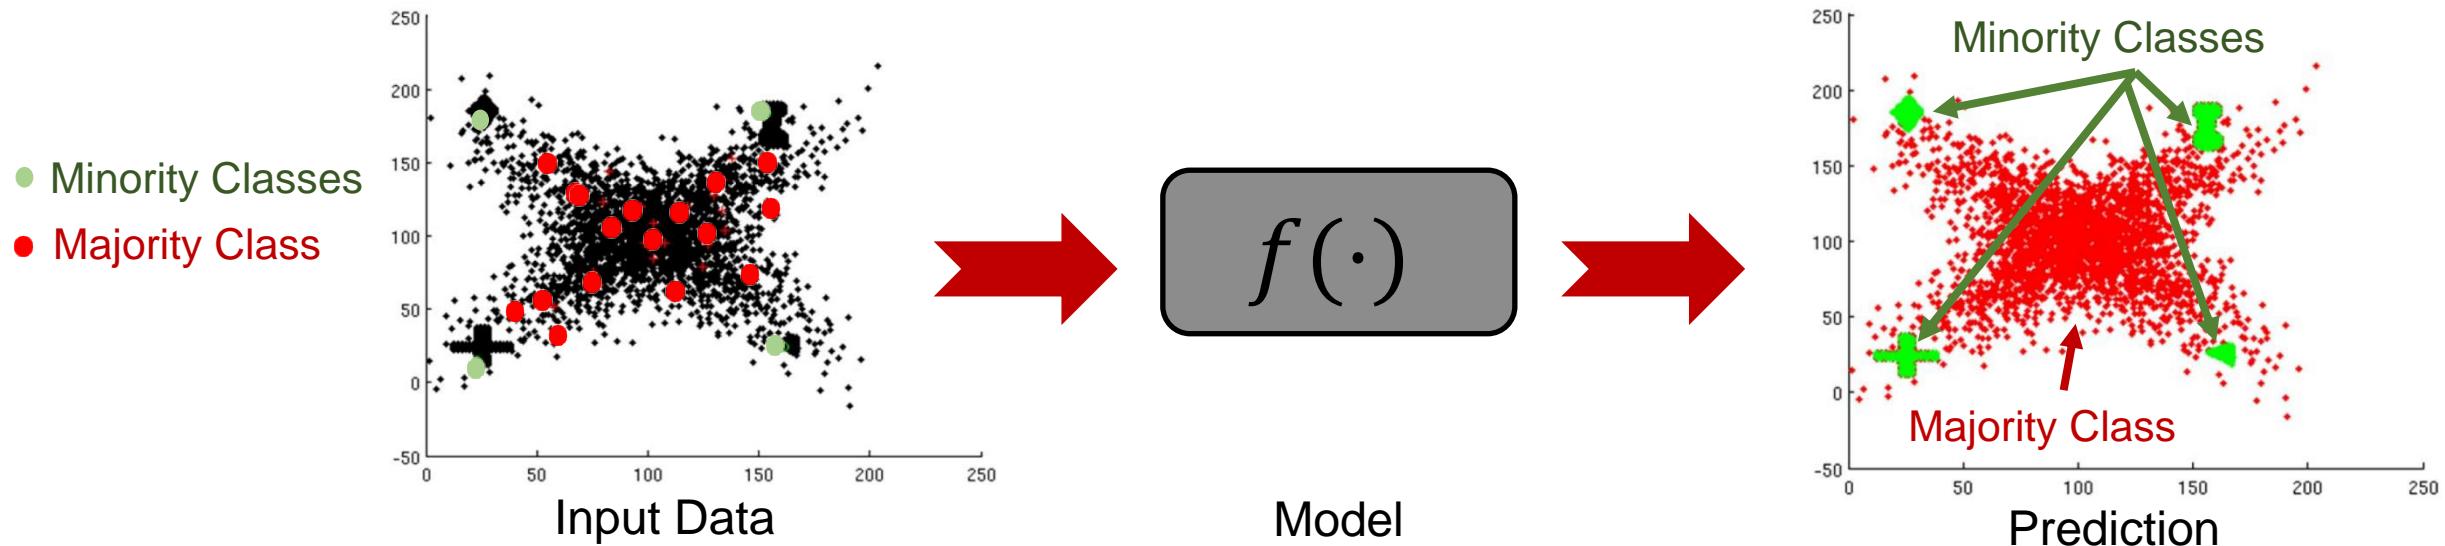
Rare Category Characterization



Rare Category Characterization

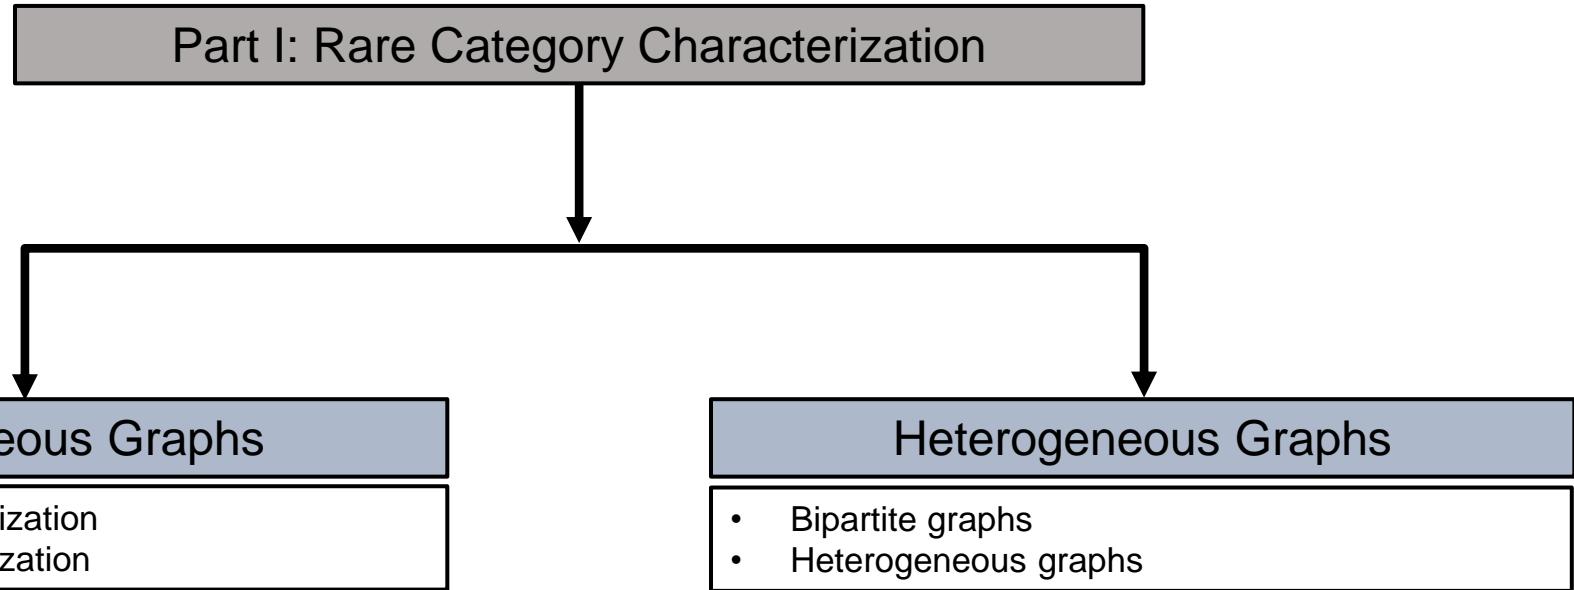
□ Problem 1.1. Rare category characterization

- **Given:** (1) A data set $D = \{x_1, \dots, x_n\}$, (2) a handful of labeled examples $L = \{x_1, \dots, x_{|L|}\}$ from both minority classes and majority classes, (3) unlabeled examples U .
- **Find:** A **compact representation** that characterizes minority class examples.



• He, Jingrui, et al. "Rare category characterization." IEEE ICDM. 2010.

Rare Category Characterization



Node-Level Characterization on Static Graphs

□ Problem 1.2. Rare category characterization on weighted graphs

- **Given:** a weighted graph $G = (V, E)$.
- **Find:** a list of potential rare examples, which are **strange, abnormal, extremely rare**.

□ Challenges

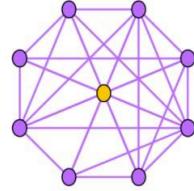
- How to extract relevant features?
- How to spot the nodes that are anomalous (**rare examples**) ?



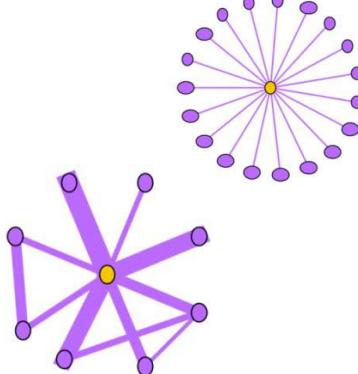
OddBall Algorithm

□ Feature extraction for egonet

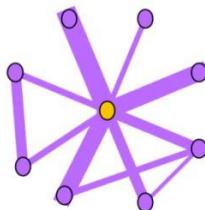
- N_i : number of neighbors (degree) of egonet i



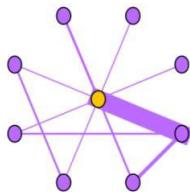
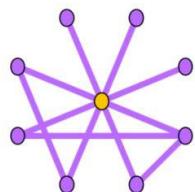
- E_i : number of edges in egonet i



- W_i : total weight of egonet i



- $\lambda_{W,i}$: principal eigenvalue of the weighted adjacency matrix of egonet i

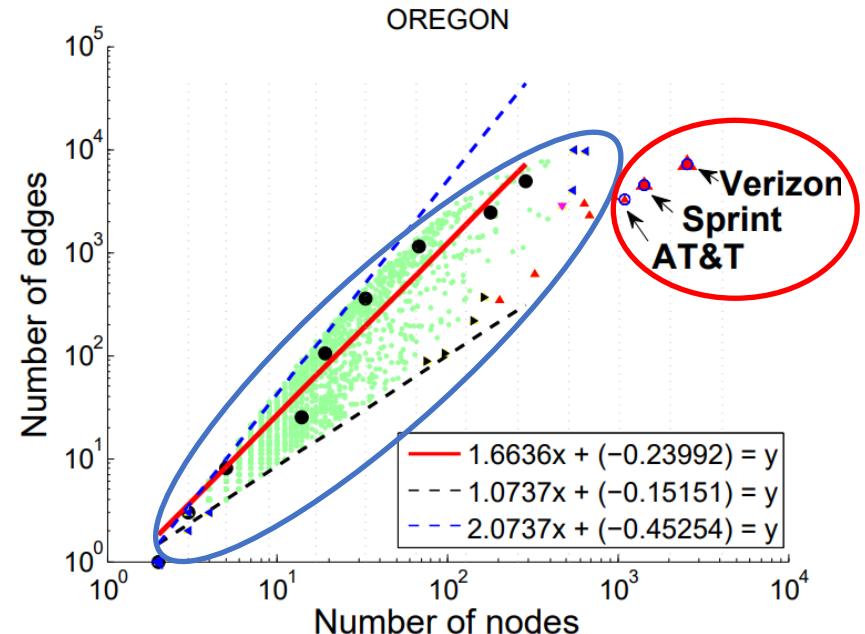


• Akoglu, Leman, et al. “OddBall: Spotting Anomalies in Weighted Graphs.” PAKDD, 2010.

OddBall Algorithm

□ Proposed method

- For each node,
 - Extract “ego-net” (=1-step neighborhood)
 - Extract features (#edges, total weight, etc.)
 - features that could yield “laws”
 - features fast to compute and interpret
- Characterize regular patterns
 - Examples from **majority classes**
 - Compute the distribution of regular patterns
- Detect irregular patterns
 - Examples from **minority classes**
 - Intuition: far from the fitting line → anomaly



• Akoglu, Leman, et al. “OddBall: Spotting Anomalies in Weighted Graphs.” PAKDD, 2010.

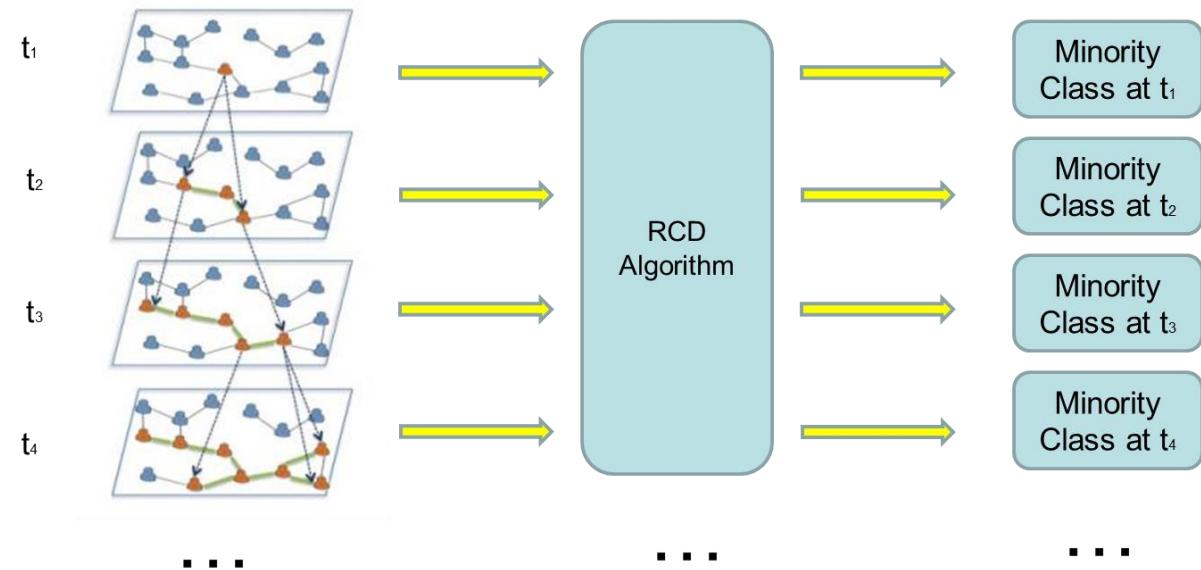
Node-Level Characterization on Dynamic Graphs

□ Problem 1.3. Rare category characterization on Time-evolving Graphs

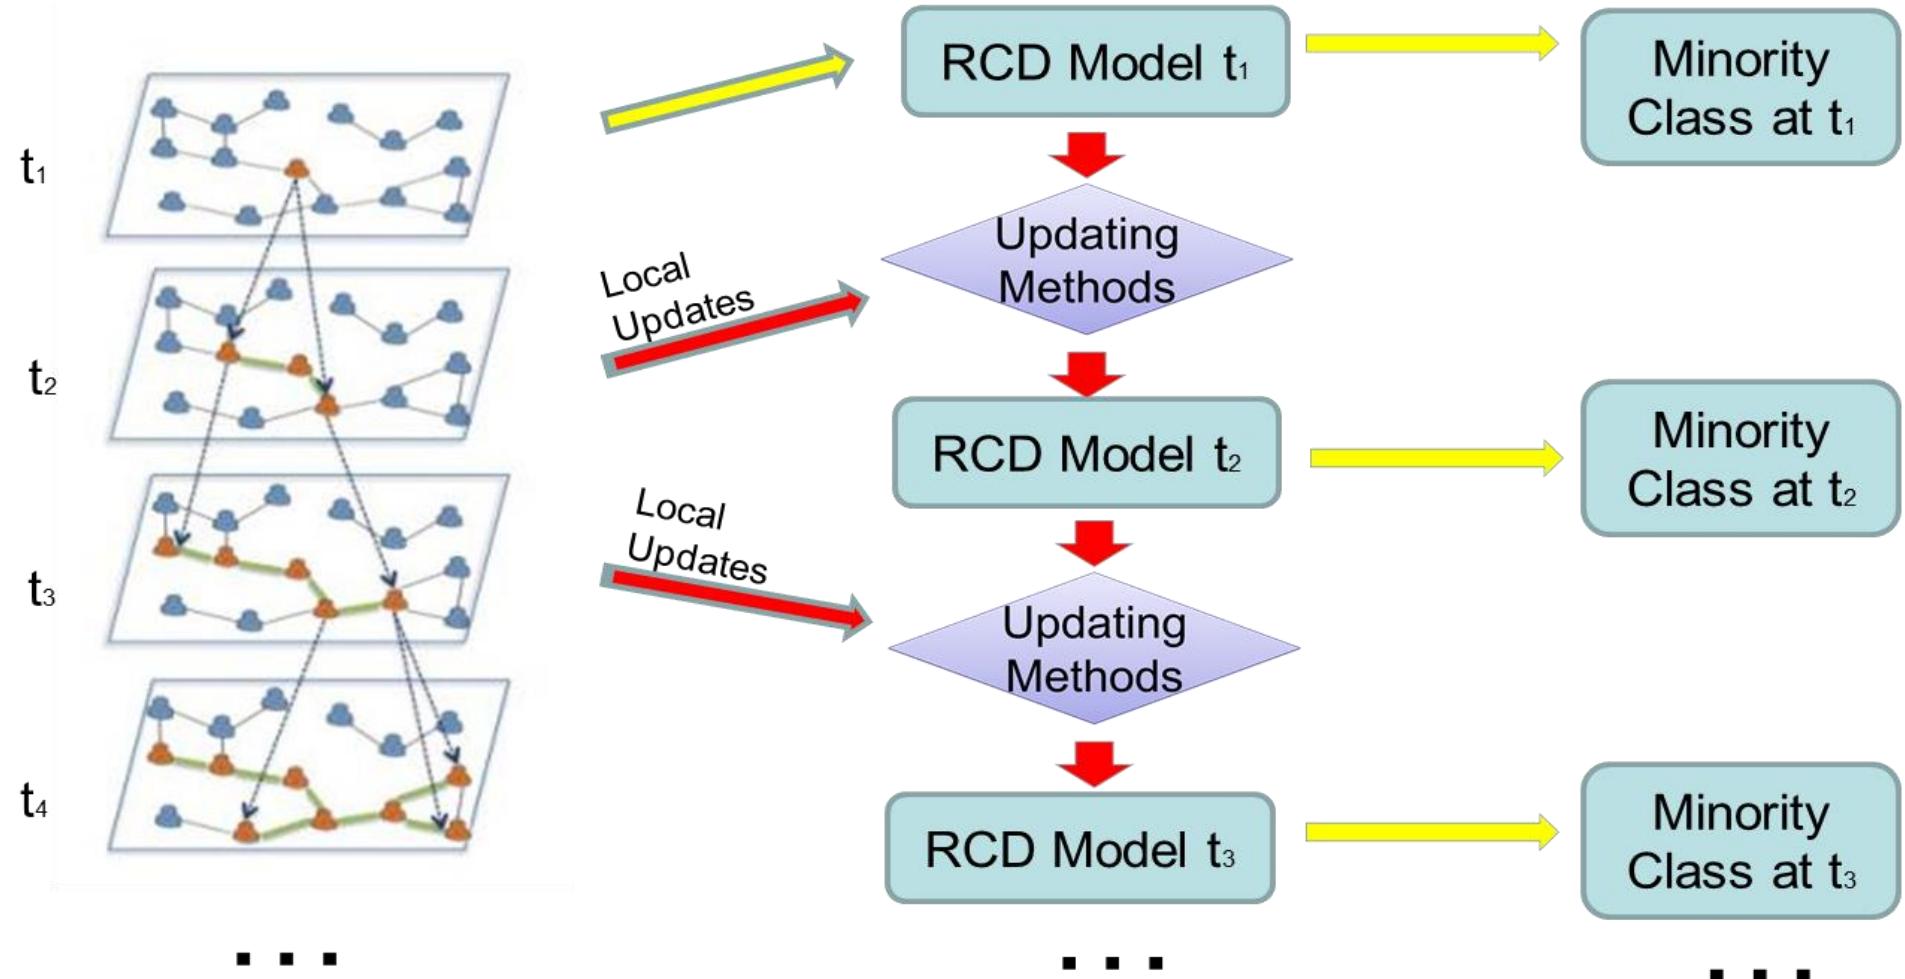
- **Given:** a time-evolving graphs $\tilde{G} = \{G_1, G_2, \dots, G_T\}$
- **Find:** examples from each minority class at each snapshot G_t , $t=1, \dots, T$.

□ Challenges

- **C1:** Graphs are evolving over time
 - New nodes/edges show up/die out
 - Edge weights change
- **C2:** High computational cost
 - Space complexity
 - Time complexity



BIRD Algorithm

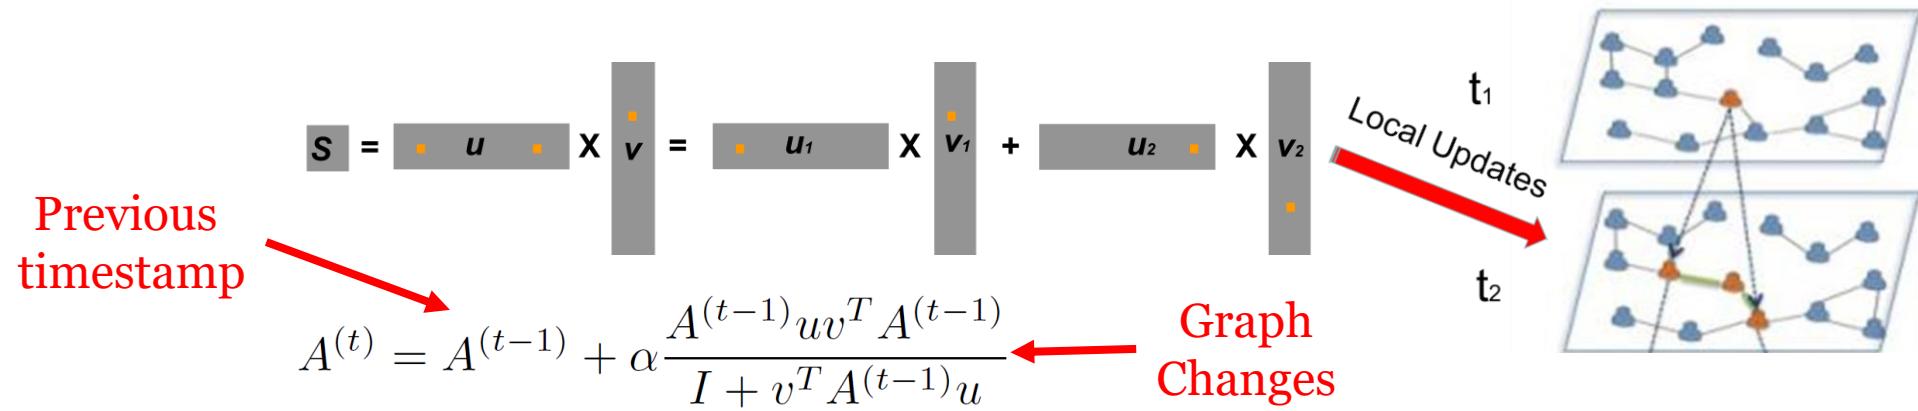


- Zhou, Dawei, et al. "Rare Category Detection on Time-Evolving Graphs." ICDM, 2015.

BIRD Algorithm

□ Proposed methods

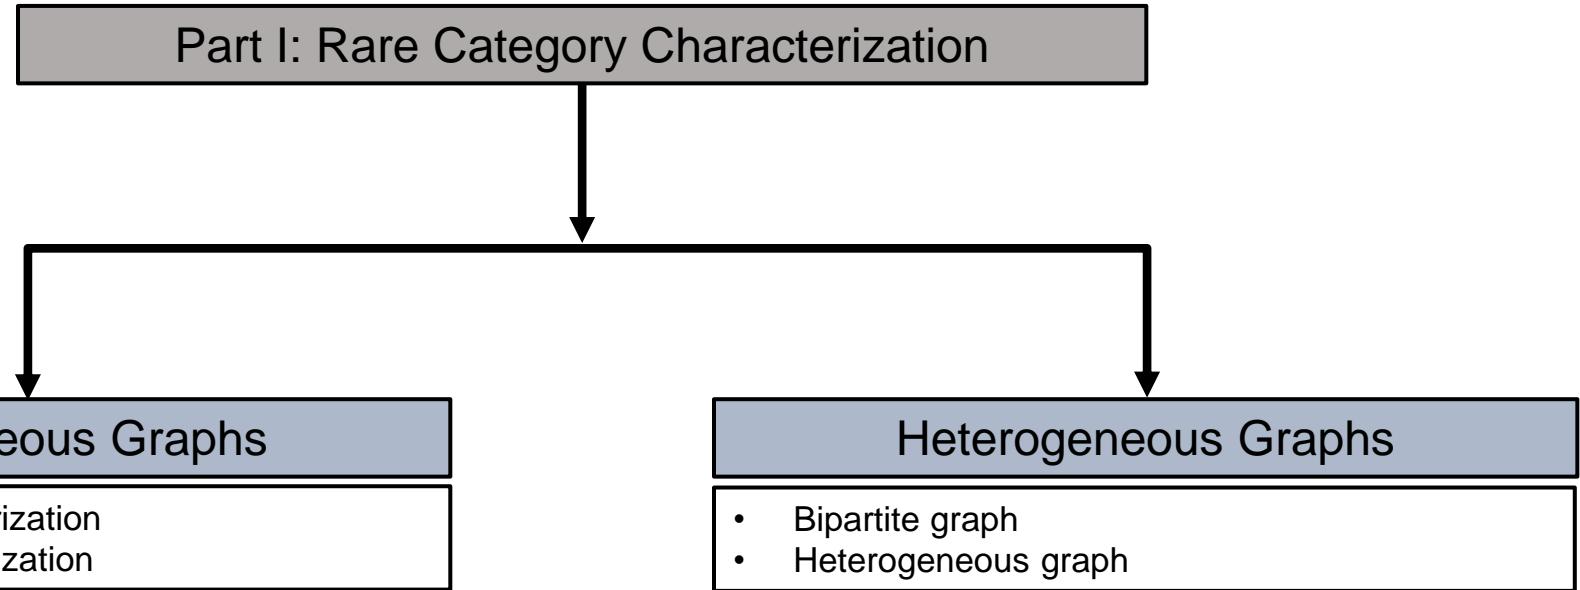
- Extract the updating matrix S from the last timestamp.
- Iteratively update A with each updated edge based on Sherman-Morrison formula.



- Compute K-NN matrix at the current timestamp.
 - Only update the rows in which the order of elements is changed in the evolved graph.
- Compute local density changes of each node.
 - Largest local density changes → potential rare category examples.

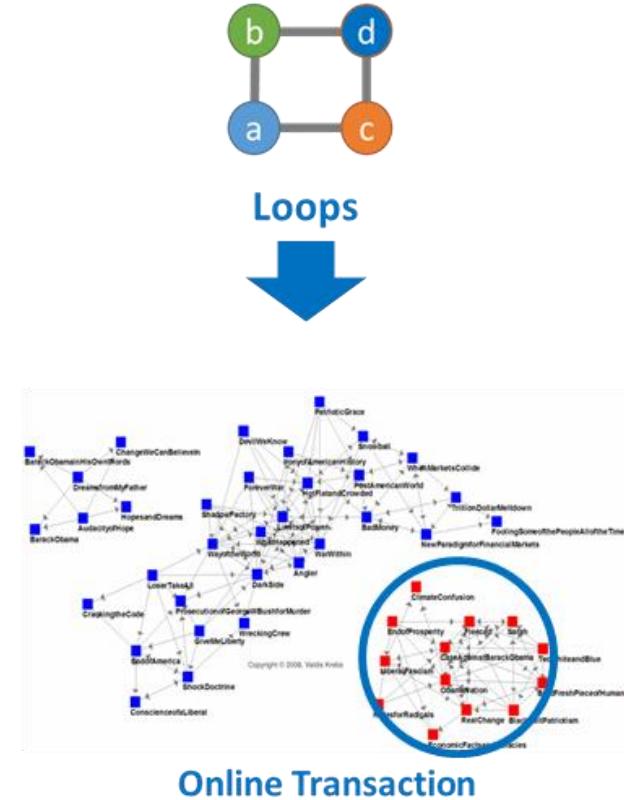
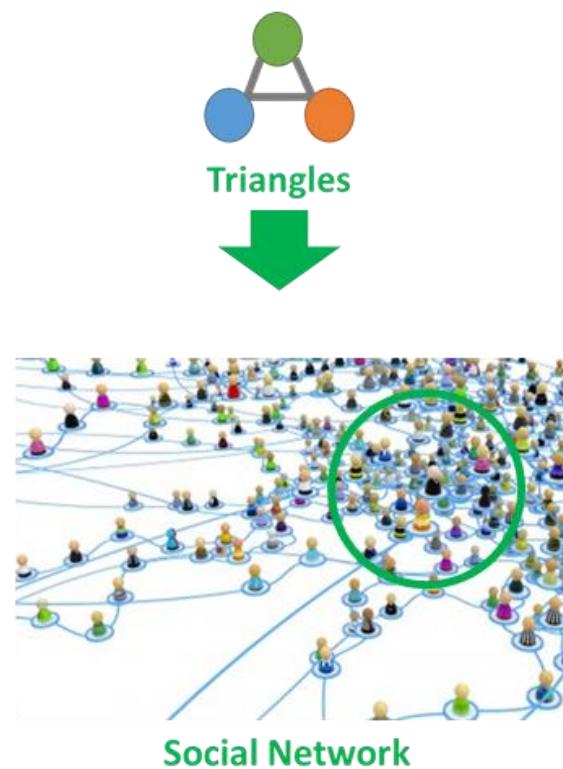
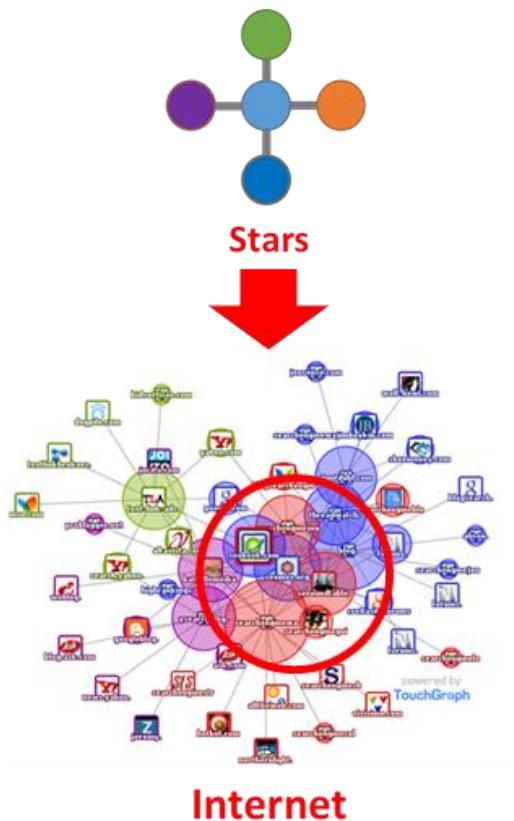
• Zhou, Dawei, et al. "Rare Category Detection on Time-Evolving Graphs." ICDM, 2015.

Rare Category Characterization



Motif-Level Characterization

□ Higher-order connectivity patterns, or **network motifs**, mediate complex networks.



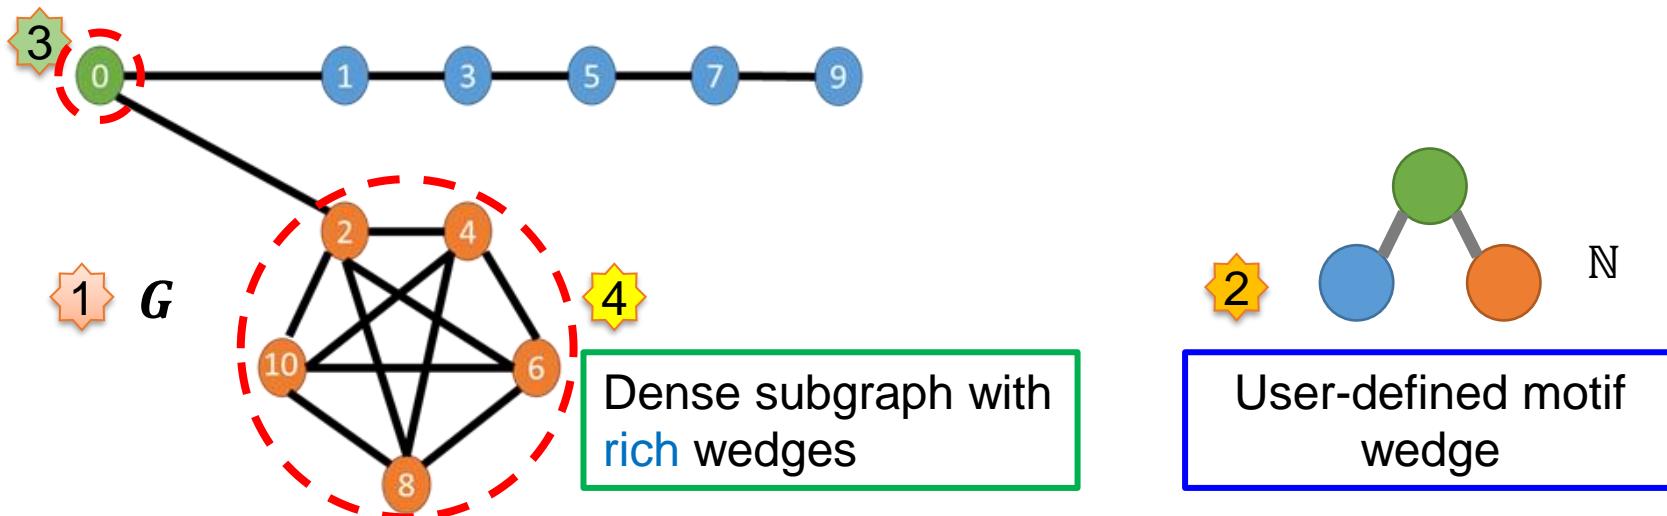
• West, Jarrod, et al. "Intelligent financial fraud detection: a comprehensive review." *Computers & security*. 2016.

- 23 - • Kim, Jooho, and Makarand Hastak. "Social network analysis: Characteristics of online social networks after a disaster." *International Journal of Information Management*. 2018.

Motif-Level Characterization

□ Problem 1.4. Structure-preserving local graph cut

- Given: Graph $G=(V, E)$, a user-defined motif N , an initial node v .
- Find: A structure-rich dense subgraph that (1) includes or near the initial node v and (2) minimizes the motif conductance w.r.t. the user-defined N .

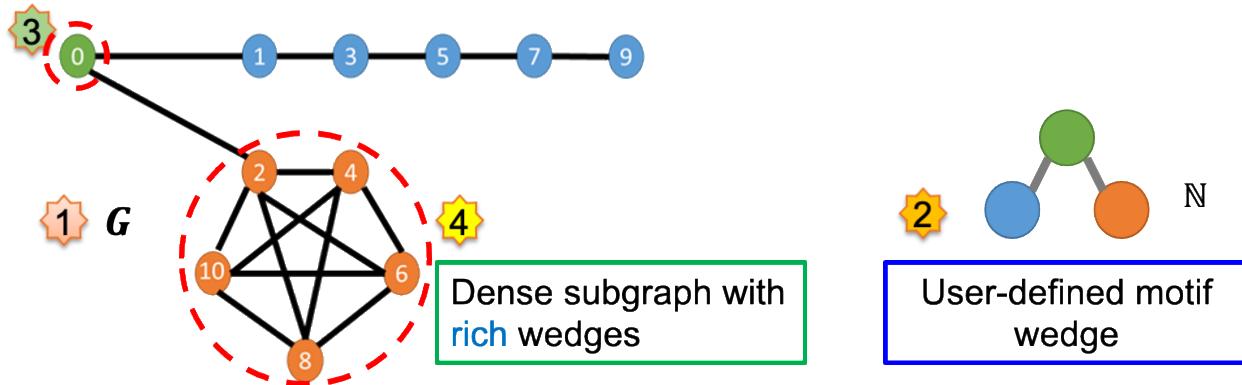


• Zhou, Dawei, et al. "A local algorithm for structure-preserving graph cut." ACM SIGKDD. 2017.

Problem

□ Problem 1.4 is NP-hard

- Any generalization of the conductance minimization problem is **NP-hard** [Wagner and Wagner, 1993].



□ An approximation solution

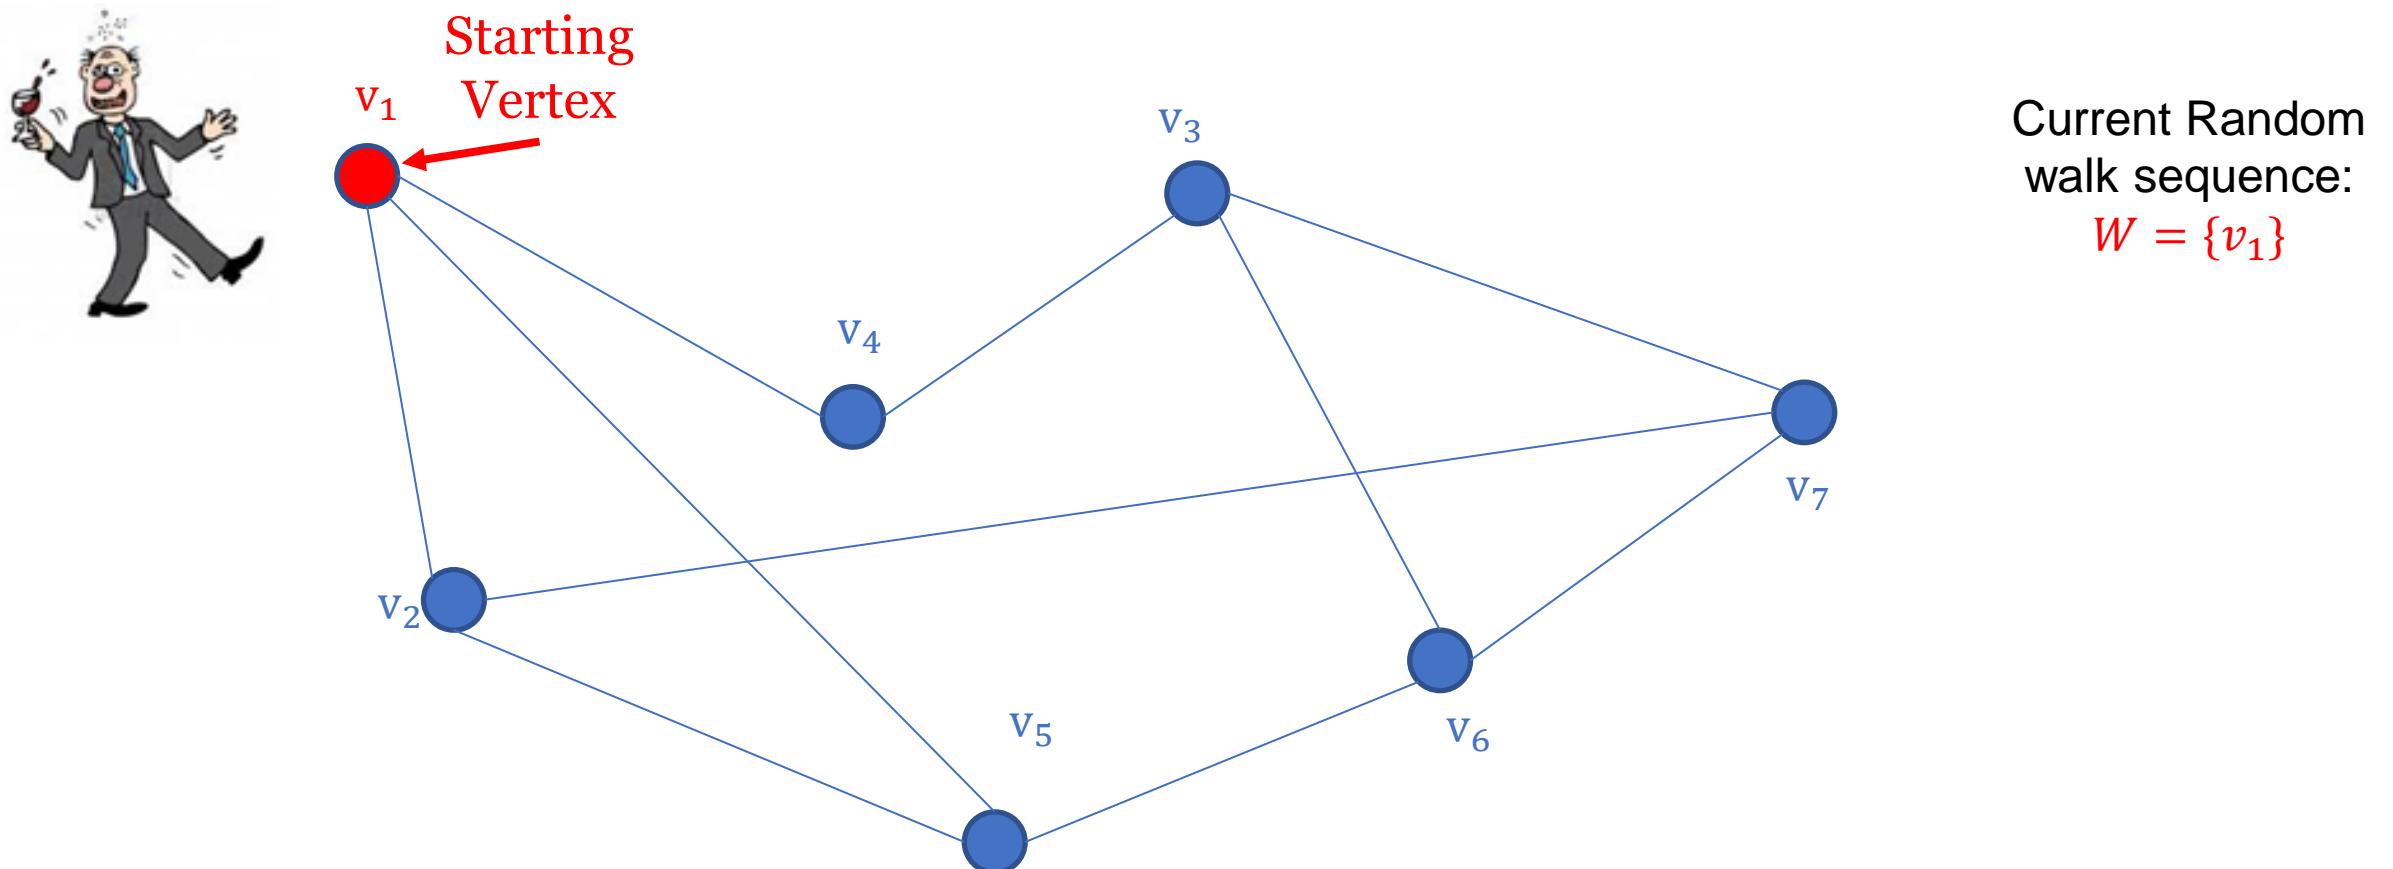
- **HOSPLOC** - High-Order Structure-Preserving LOcal Graph Cut.
- A random walk based method.

• Wagner, Dorothea, and Frank Wagner. "Between min cut and graph bisection." MFCS, 1993.

- 25 - • Spielman, Daniel A., and Shang-Hua Teng. "A local clustering algorithm for massive graphs and its application to nearly linear time graph partitioning." SIAM Journal on computing, 2013.

Graph Random Walk

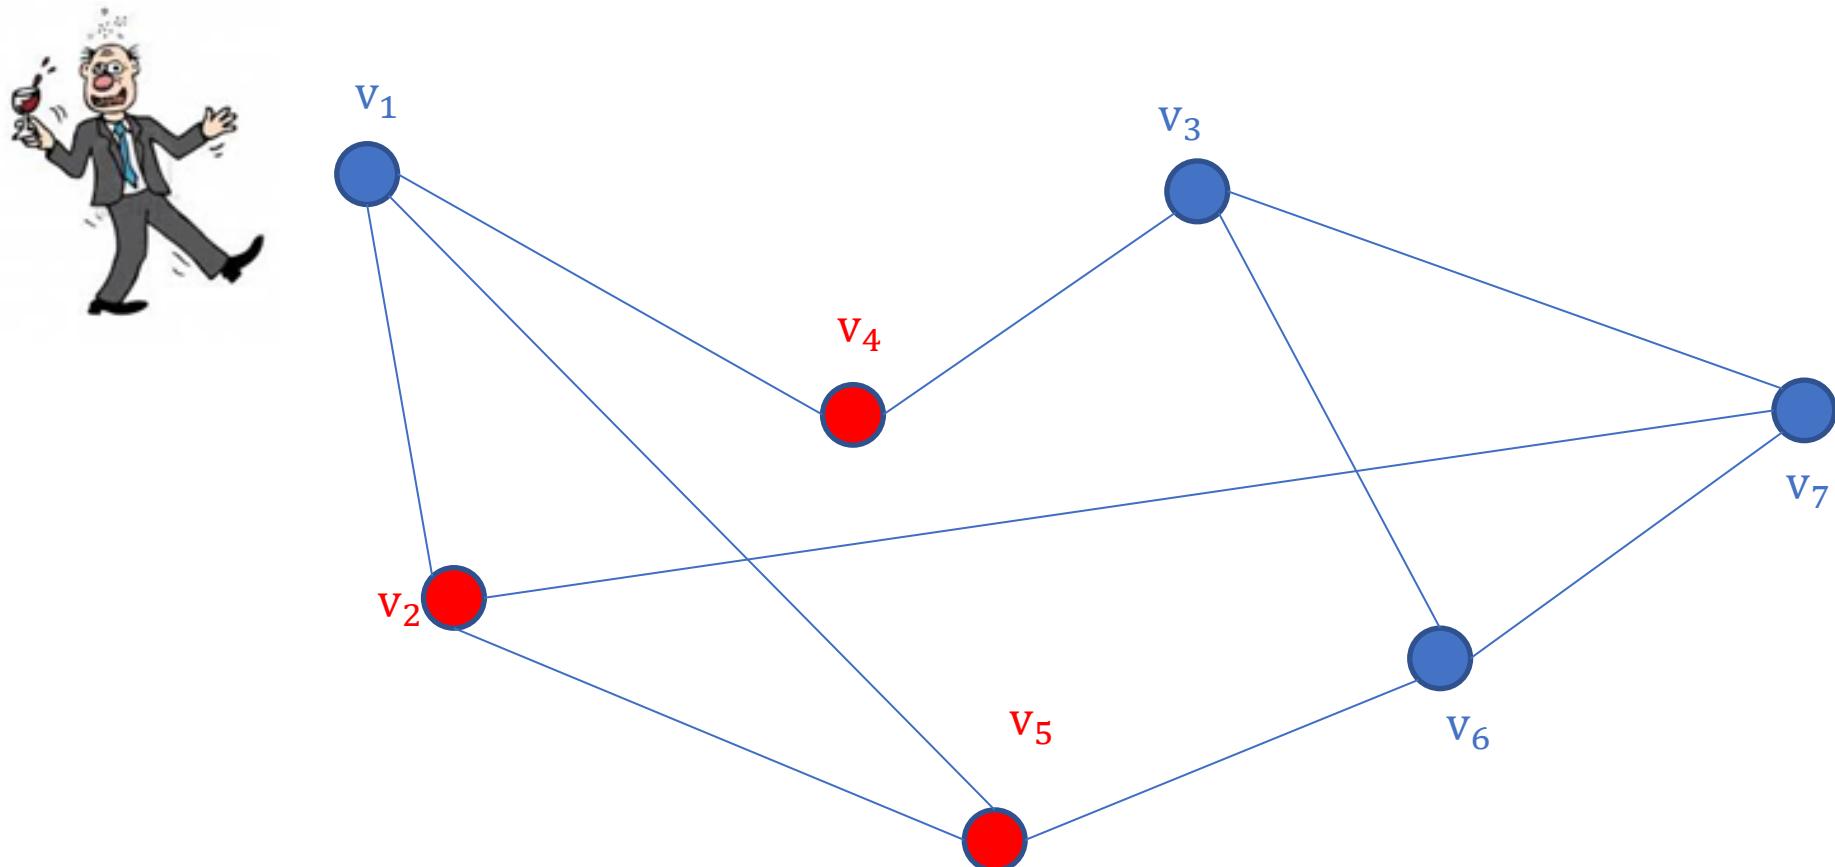
- A random walk on a graph is a process that begins at some vertex, and at each time step moves to another vertex.



• Tong, Hanghang, Christos Faloutsos, and Jia-Yu Pan. "Fast random walk with restart and its applications." *Sixth international conference on data mining (ICDM'06)*. IEEE, 2006.

Graph Random Walk

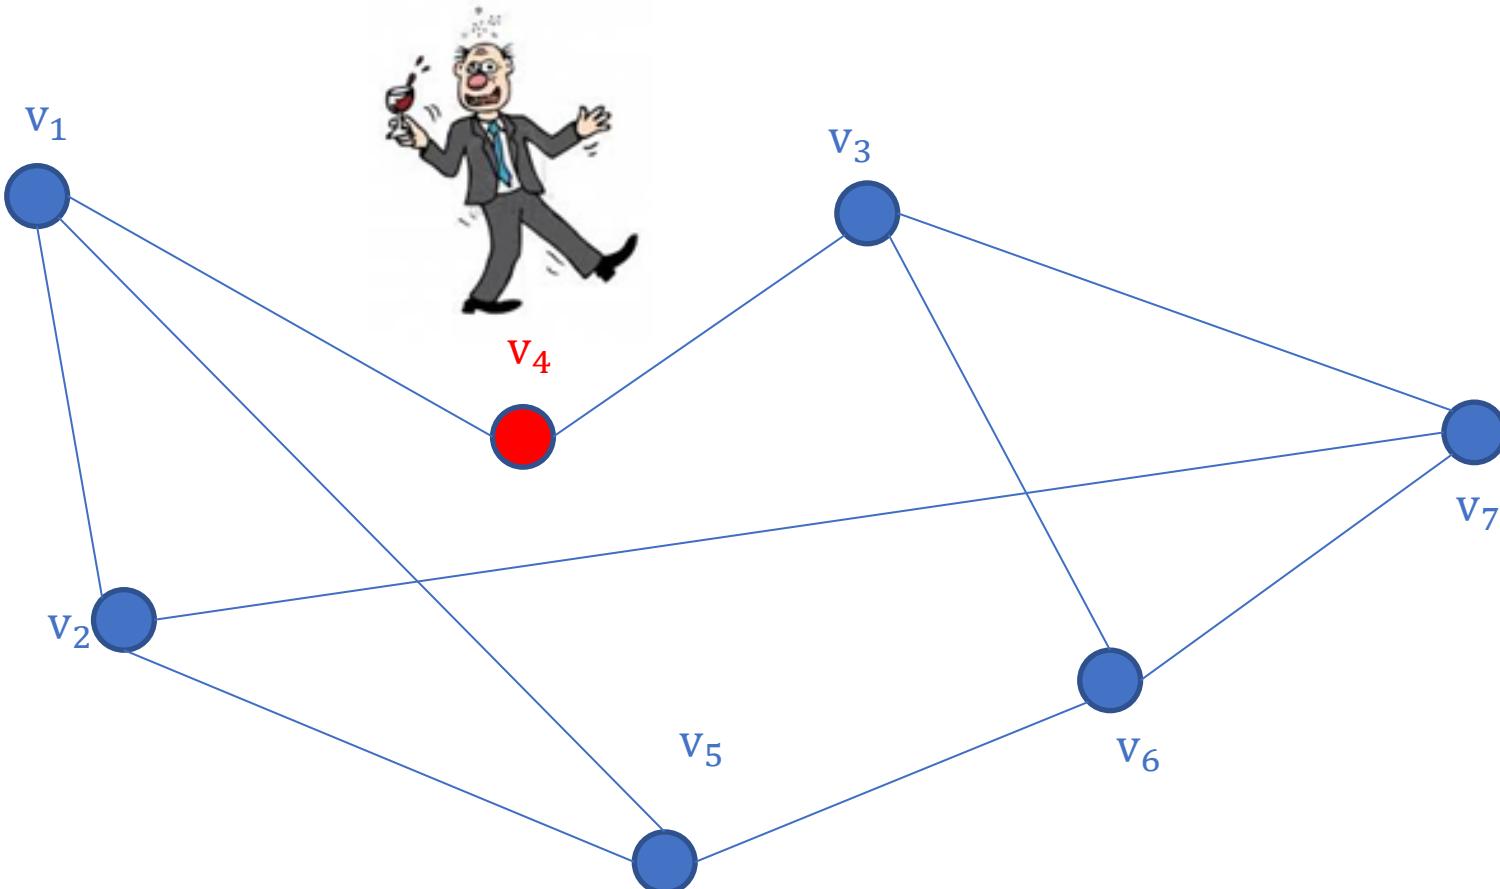
- A random walk on a graph is a process that begins at some vertex, and at each time step moves to another vertex.



Current Random
walk sequence:
 $W = \{v_1\}$

Graph Random Walk

- A random walk on a graph is a process that begins at some vertex, and at each time step moves to another vertex.

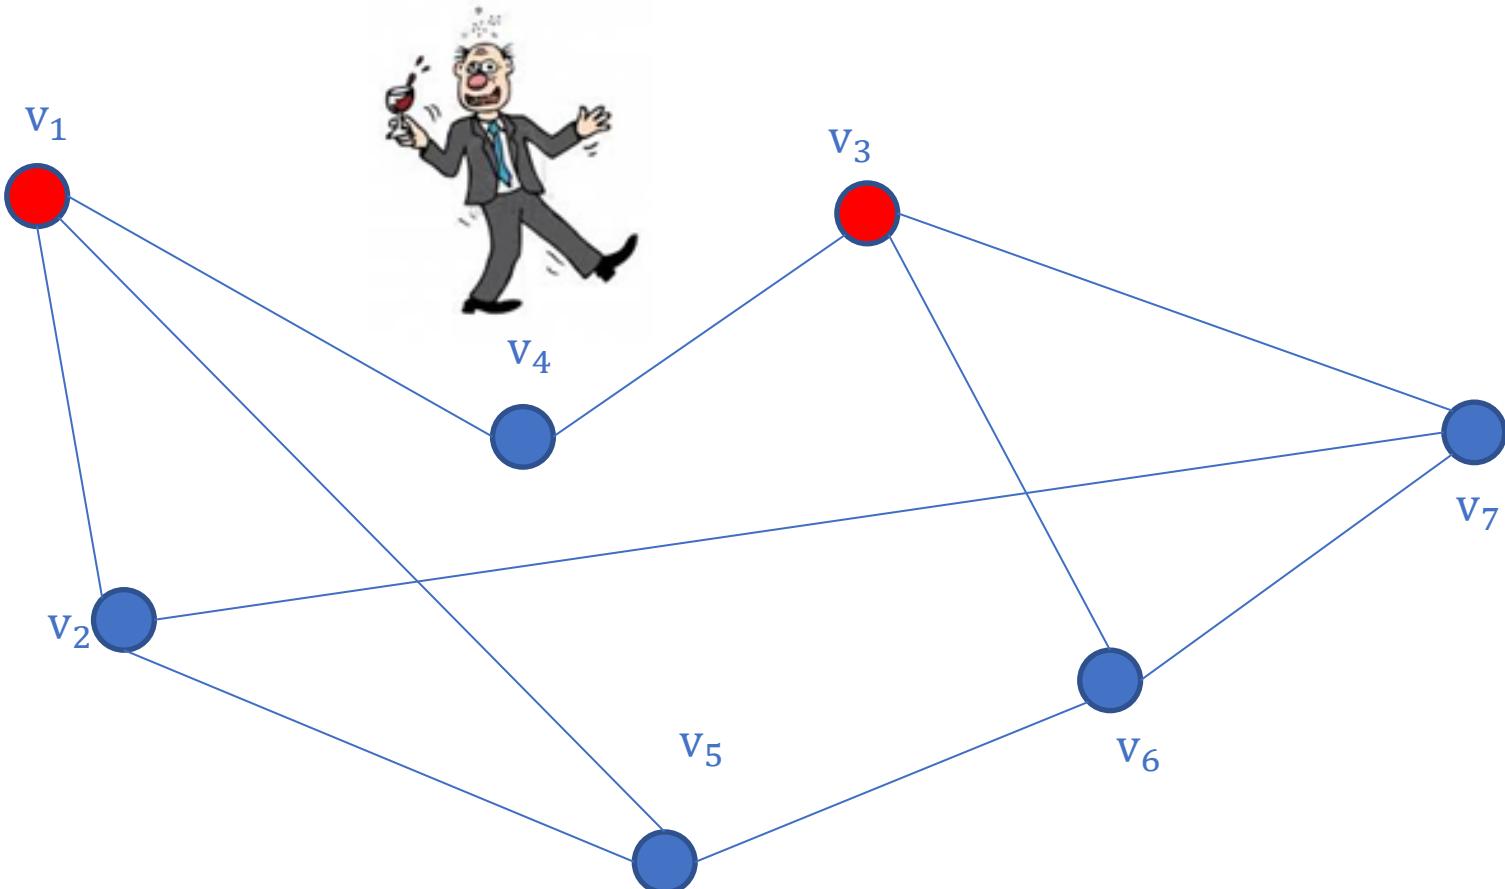


Current Random
walk sequence:
 $W = \{v_1, v_4\}$

- Tong, Hanghang, Christos Faloutsos, and Jia-Yu Pan. "Fast random walk with restart and its applications." *Sixth international conference on data mining (ICDM'06)*. IEEE, 2006.

Graph Random Walk

- A random walk on a graph is a process that begins at some vertex, and at each time step moves to another vertex.

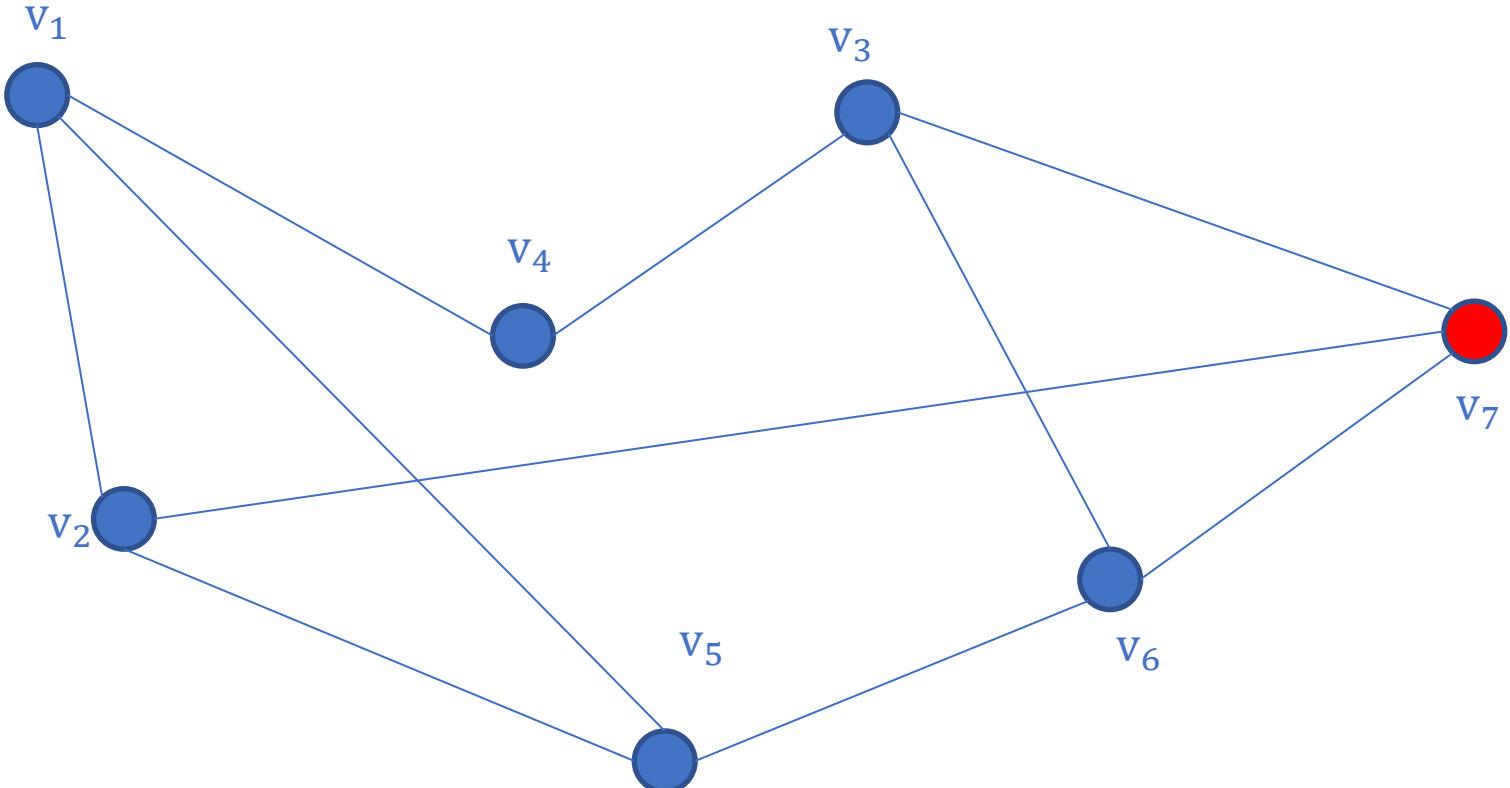


Current Random
walk sequence:
 $W = \{v_1, v_4\}$

• Tong, Hanghang, Christos Faloutsos, and Jia-Yu Pan. "Fast random walk with restart and its applications." *Sixth international conference on data mining (ICDM'06)*. IEEE, 2006.

Graph Random Walk

- A random walk on a graph is a process that begins at some vertex, and at each time step moves to another vertex.



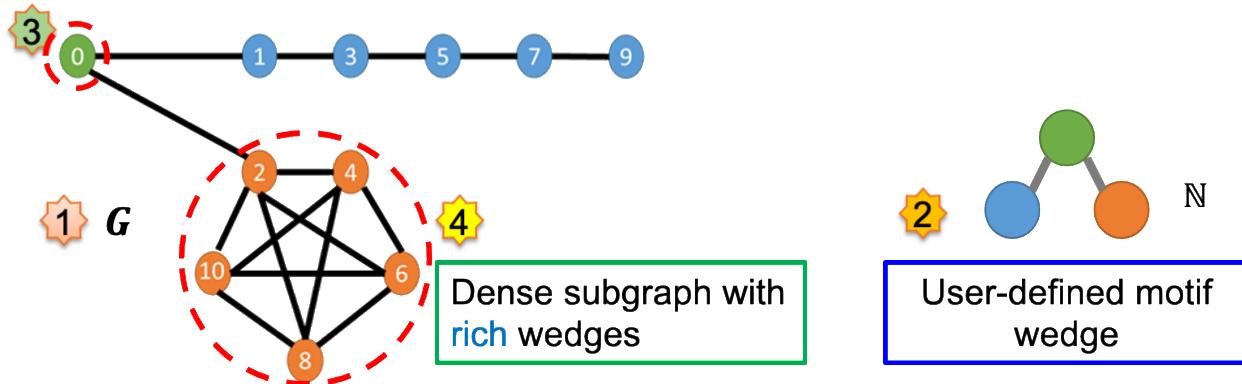
k -length Random walk sequence:
$$W = \{v_1, v_4, v_3, \dots, v_7\}$$



Problem

□ Problem 1.4 is NP-hard

- Any generalization of the conductance minimization problem is **NP-hard** [Wagner and Wagner, 1993].



□ An approximation solution

- **HOSPLOC - High-Order Structure-Preserving LOcal Graph Cut.**
- A random walk based method.
- **Space and time complexity** could be greatly reduced.

• Wagner, Dorothea, and Frank Wagner. "Between min cut and graph bisection." MFCS, 1993.

- 31 - • Spielman, Daniel A., and Shang-Hua Teng. "A local clustering algorithm for massive graphs and its application to nearly linear time graph partitioning." SIAM Journal on computing, 2013.

HOSPLLOC Algorithm

□ Algorithm overview

- **Step 1.** Modeling high-order motifs by constructing **adjacency tensor** of the original graph.
- **Step 2.** Exploring neighborhood context around initial node via **truncated high-order random walk**.
- **Step 3.** Conducting structure-preserving graph cut via **sweep cut procedure** on the stationary distribution vectors of high-order random walk.

□ Properties

- **Quality guarantee**
 - Finding a **near-optimal** local cluster with minimum motif conductance.
- **Scalability guarantee**
 - Procedure stops upon finding a good cluster, no need to explore the rest of graph.
 - Algorithm runs in **polylogarithmic** time w.r.t. the number of edges of original graph.

• Wagner, Dorothea, and Frank Wagner. "Between min cut and graph bisection." MFCS. 1993.

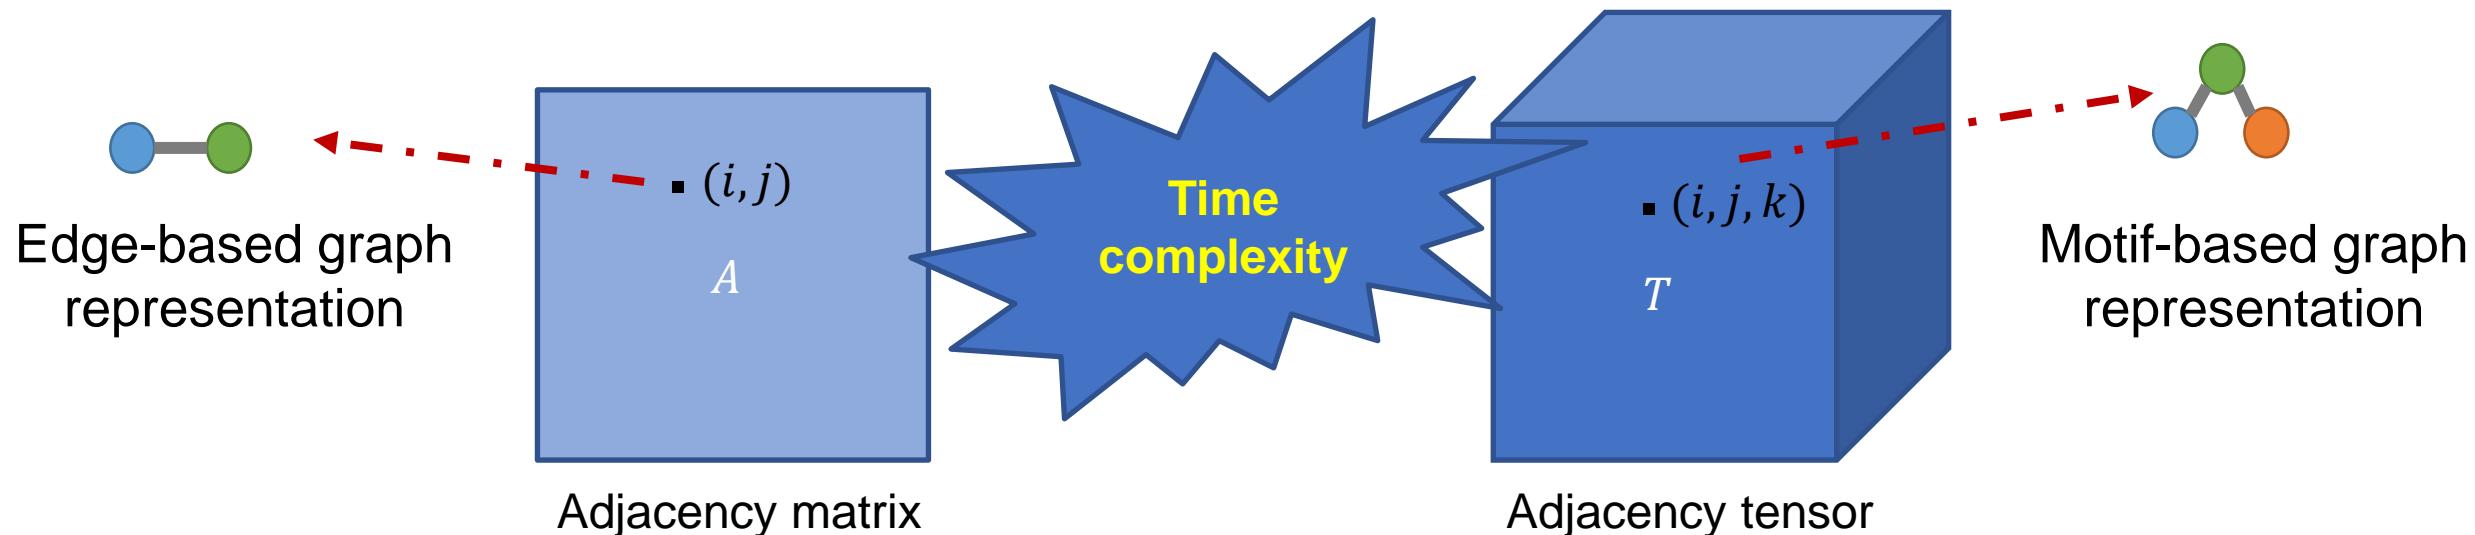
- 32 - • Spielman, Daniel A., and Shang-Hua Teng. "A local clustering algorithm for massive graphs and its application to nearly linear time graph partitioning." SIAM Journal on computing. 2013.

HOSPLOC – Step I: Modeling Network Motifs

□ Adjacency tensor

Definition 1.1. Given a graph $G=(V,E)$, the k^{th} -order network motif \mathbb{N} on G could be represented in a k -dimensional adjacency tensor T as follows

$$T(i_1, i_2, \dots, i_k) = \begin{cases} 1 & (i_1, i_2, \dots, i_k) \in V \text{ and form } \mathbb{N} \\ 0 & \text{Otherwise} \end{cases}$$



- Zhou, Dawei, et al. "A local algorithm for structure-preserving graph cut." ACM SIGKDD. 2017.

HOSPLOC – Step II: Exploring Neighborhood

□ Truncated high-order random walk (THRW)

- Exploring k^{th} -order network structure \mathbb{N} .

Transition tensor

$$P(i_1, i_2, \dots, i_k) = \frac{T(i_1, i_2, \dots, i_k)}{\sum_{i_1=1}^n T(i_1, i_2, \dots, i_k)}$$

- Stationary distribution of HRW with “rank-1” assumption [Li and NG, 2013].

HRW stationary distribution

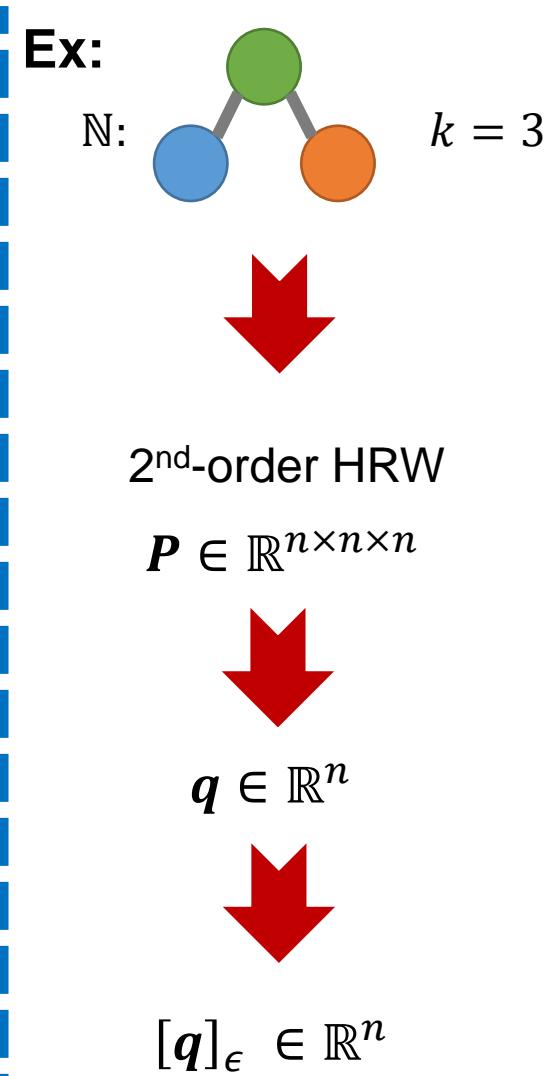
$$q(i_1) = \sum_{(i_2, \dots, i_k)} P(i_1, i_2, \dots, i_k) q(i_2) \dots q(i_k) \quad O(n^k)!$$

- Rounding small values of q to 0 to reduce time complexity.

Cluster indicator vector

$$[q]_\epsilon(u) = \begin{cases} q(u) & \text{If } q(u) > d(u)\epsilon \\ 0 & \text{Otherwise} \end{cases}$$

Sparse computation!



- Li, Wen, and Michael K. Ng. "On the limiting probability distribution of a transition probability tensor." Linear and Multilinear Algebra. 2014.
- Benson, Austin R., et al. "The spacey random walk: A stochastic process for higher-order data." SIAM Review. 2017.

HOSPLLOC – Step III: Conducting Sweep Cut

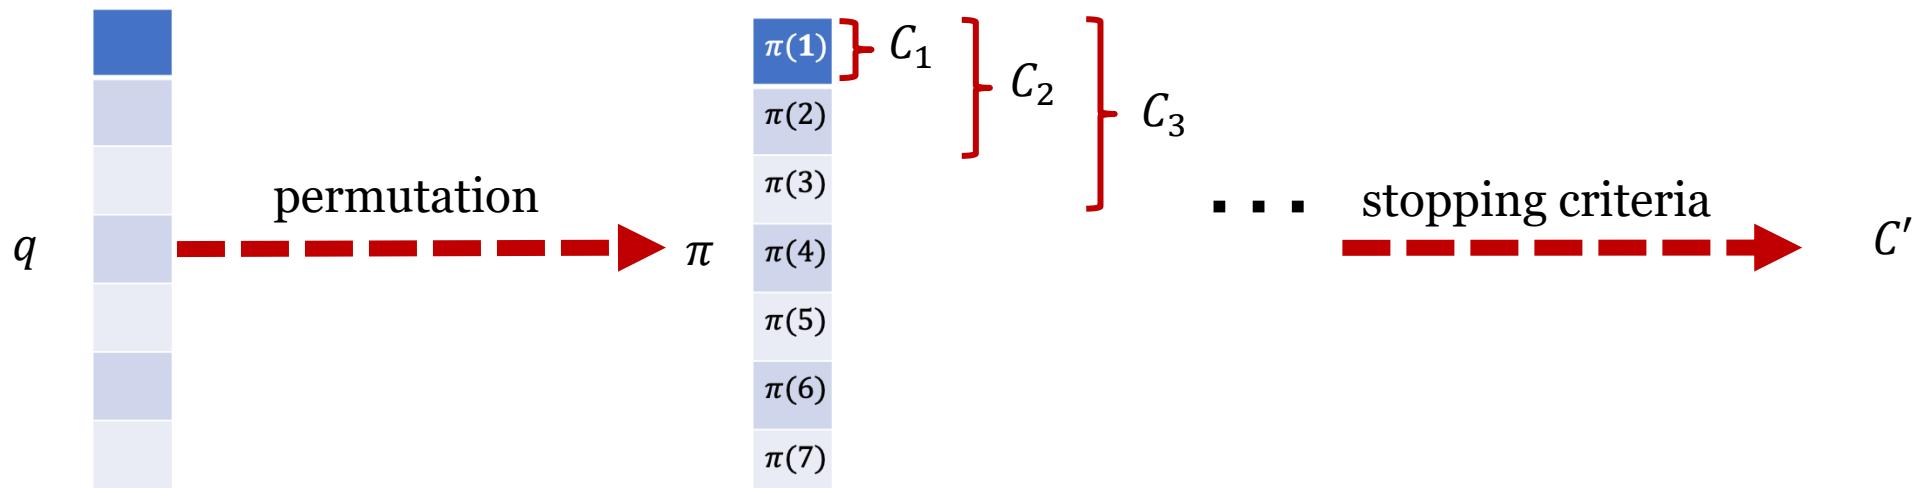
□ Sweep cut procedure

- Obtaining high-order random walk distribution by performing THRW
- Computing the permutation π of vector q in the way that:

$$\frac{q(\pi(1))}{d(\pi(1))} \geq \frac{q(\pi(2))}{d(\pi(2))} \geq \dots \geq \frac{q(\pi(n))}{d(\pi(n))}$$

Node Degree

- Iteratively checking the potential cuts C_1, C_2, \dots, C_{n-1} , where $C_i = \{\pi(1), \dots, \pi(i)\}$, and identify the first cut C' that satisfies the user-defined stopping criteria.



- Zhou, Dawei, et al. "A local algorithm for structure-preserving graph cut." ACM SIGKDD. 2017.
- Spielman, Daniel A., et al. "A local clustering algorithm for massive graphs and its application to nearly linear time graph partitioning." SIAM Journal on computing. 2013.

MAPPR Algorithm

□ Motif-based approximate personalized PageRank algorithm (MAPPR)

- Step 1. Create a weighted graph with weighted edge conductance that equals to the motif conductance in the original graph.

$$w(i, j) = \#\text{motif instances containing nodes } i \text{ and } j.$$

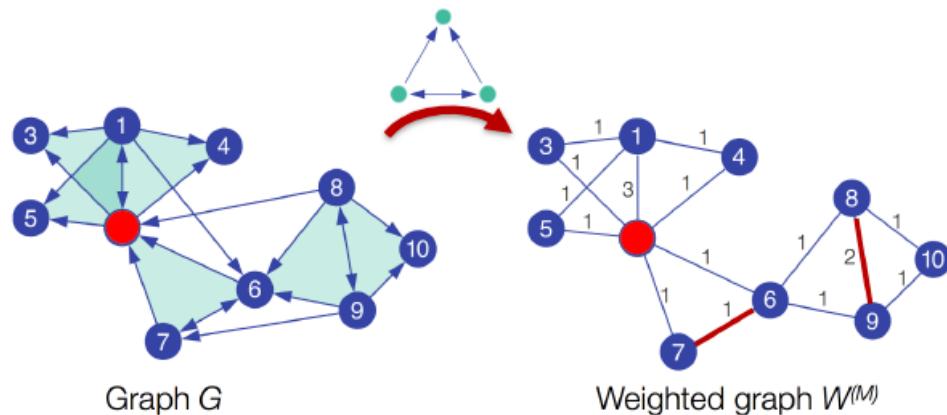


Figure adopted from [Yin et al., KDD 2017]

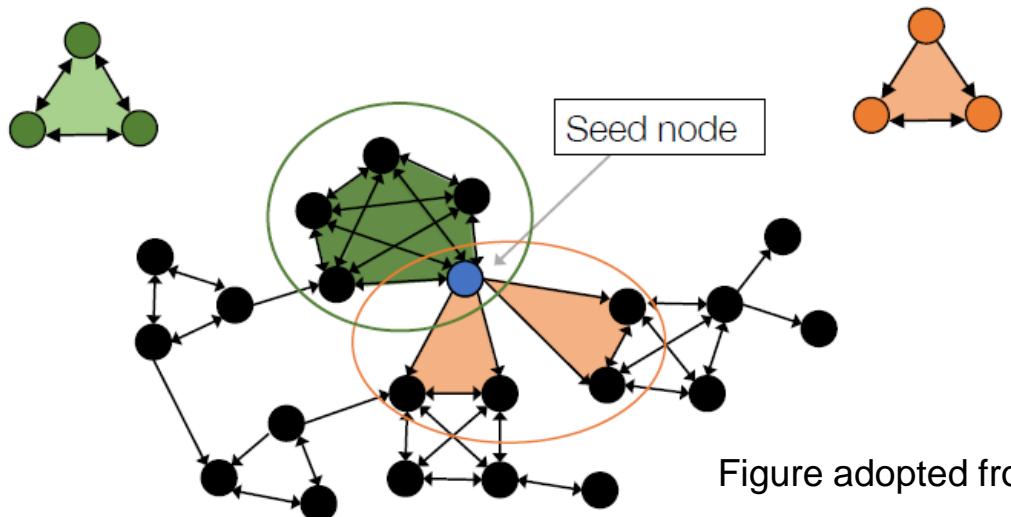
The motif conductance (approximately) equals the weighted edge conductance in this weighted graph [Benson et al., 16].

- Yin, Hao, et al. “Local higher-order graph clustering.” ACM SIGKDD, 2017.
- Benson, Austin, et al. “Higher-order organization of complex networks.” Science, 2016.

MAPPR Algorithm

□ Motif-based approximate personalized PageRank algorithm (MAPPR)

- **Step 1.** Create a weighted graph with weighted edge conductance equals the motif conductance in the original graph.
- **Step 2.** Compute an approximate PPR vector for this weighted graph.
- **Step 3.** Find a cluster of minimal weighted edge conductance.



Different motifs give different local clustering results!

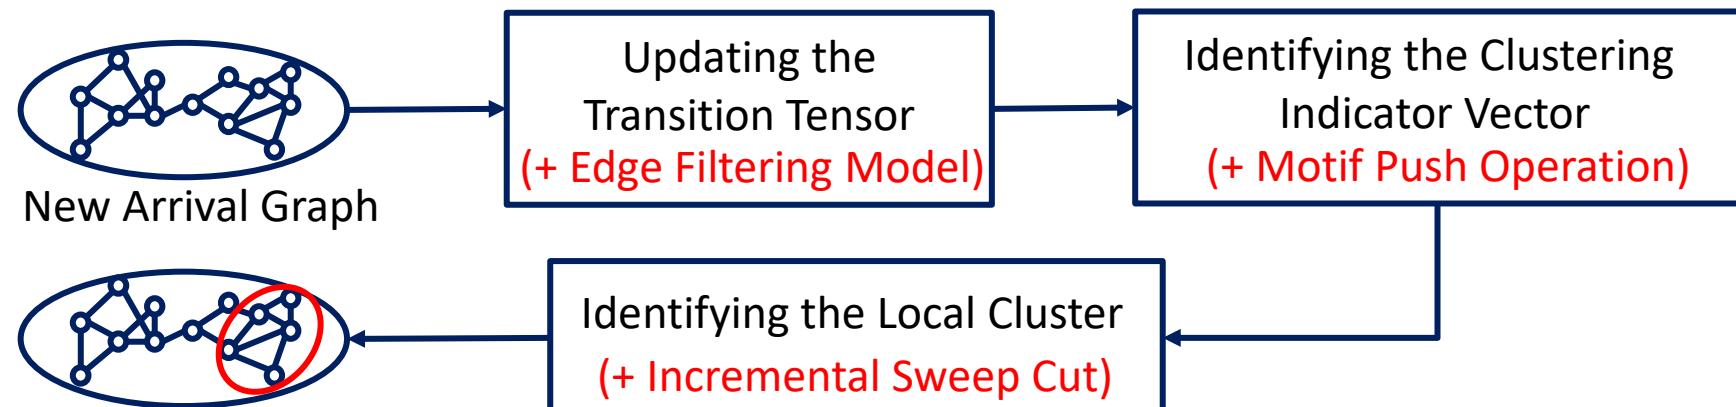
Figure adopted from [Yin et al., KDD 2017]

- Yin, Hao, et al. "Local higher-order graph clustering." ACM SIGKDD 2017.

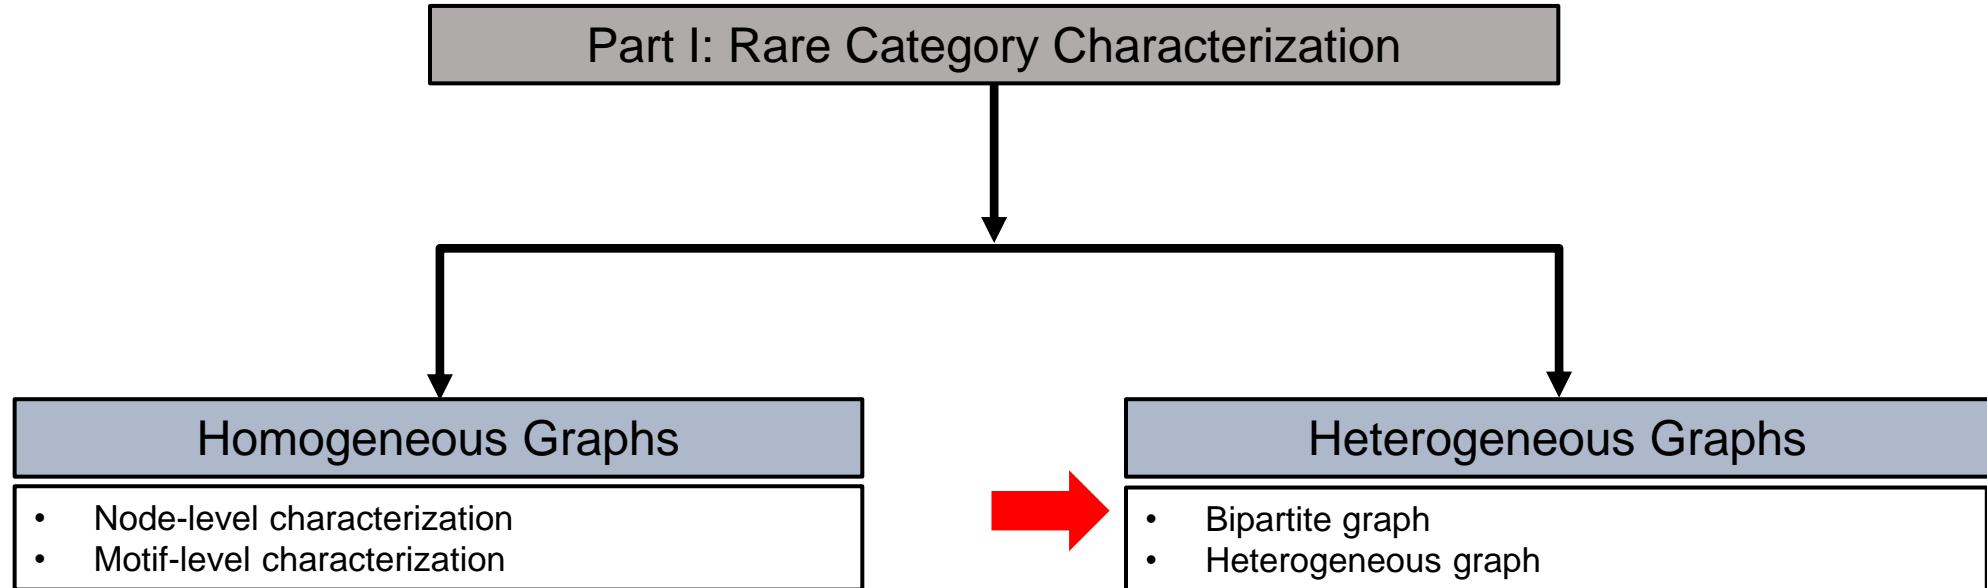
L-MEGA Algorithm

□ Local motif clustering on time-evolving graphs (L-MEGA)

- **Goal:** Identifying the evolution pattern of the local motif cluster effectively and efficiently.
- **Key idea**
 - **Edge filtering model:** Identifying “far-away” edge (v_1, v_2) , filter it out before updating the transition tensor, and save it for the future.
 - **Motif push operation:** Tracking multilinear PageRank $q^{(t+1)}$ from $q^{(t)}$ via incremental updating strategies.
 - **Incremental sweep cut:** Identifying shared sequence between two consecutive time permutations and start from the first different entry.



Rare Category Characterization

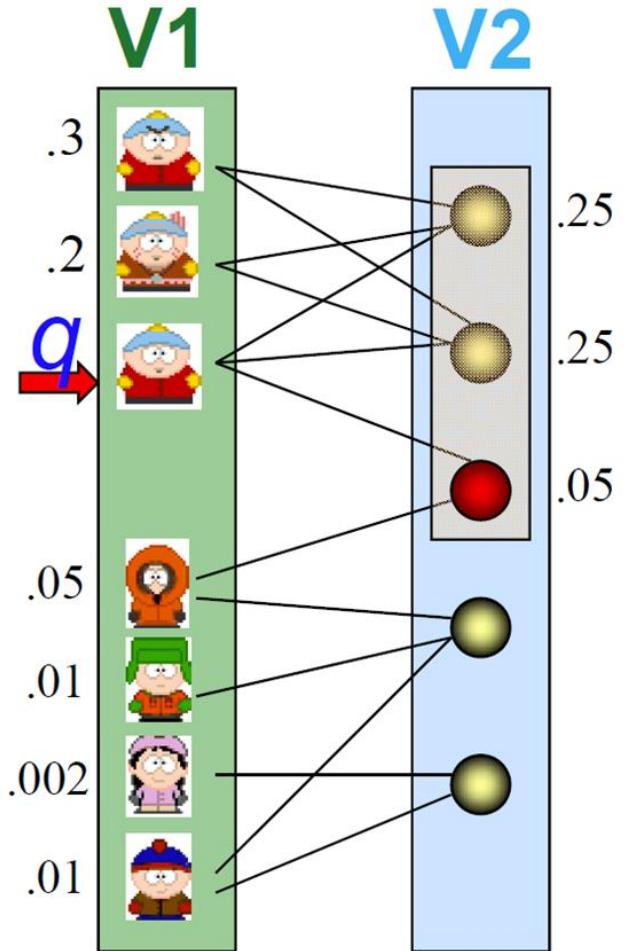


Rare Category in Bipartite Graphs

□ Problem 1.6. Rare category characterization in Bipartite graphs

- Neighborhood formation (NF)
 - Given a query node q in $V1$, what are the relevance scores of all the nodes in $V1$ to q ?
 - Ex: Similar authors in publication networks

- Anomaly detection (AD)
 - Given a query node q in $V1$, what are the normality scores for nodes in $V2$ that link to q ?
 - Ex: Unusual papers in publication networks

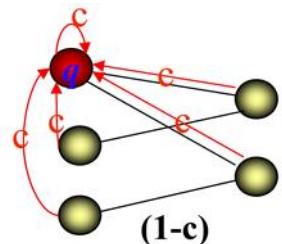


• Sun, Jimeng, et al. “Neighborhood formation and anomaly detection in bipartite graphs.” ICDM, 2005.

Rare Category in Bipartite Graphs

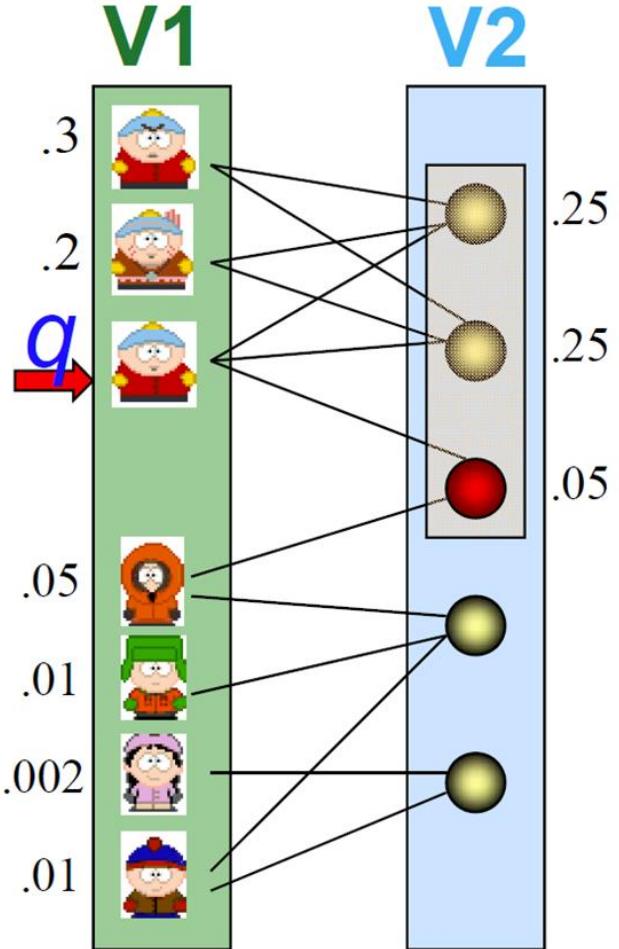
□ Main idea for neighborhood formation

- Conduct random-walk-with-restart from q
- Compute steady-state V_1 as neighborhood relevance
 - Construct transition matrix P
 - Fly-back probability c to q



- Solve for steady state

$$v^{(t+1)} = Pv^{(t)} + cq$$



- Sun, Jimeng, et al. "Neighborhood formation and anomaly detection in bipartite graphs." ICDM, 2005.

Rare Category in Bipartite Graphs

□ Main idea for anomaly detection

- Pairwise “normality” scores of neighbors t
- Function of (e.g., avg) pair-wise scores
 - Find set S of nodes connected to t
 - Compute $|S|^*|S|$ normality matrix R
 - Asymmetric, diagonal reset to 0
 - Apply score function $f(R)$
 - Ex: $f(R) = mean(R)$

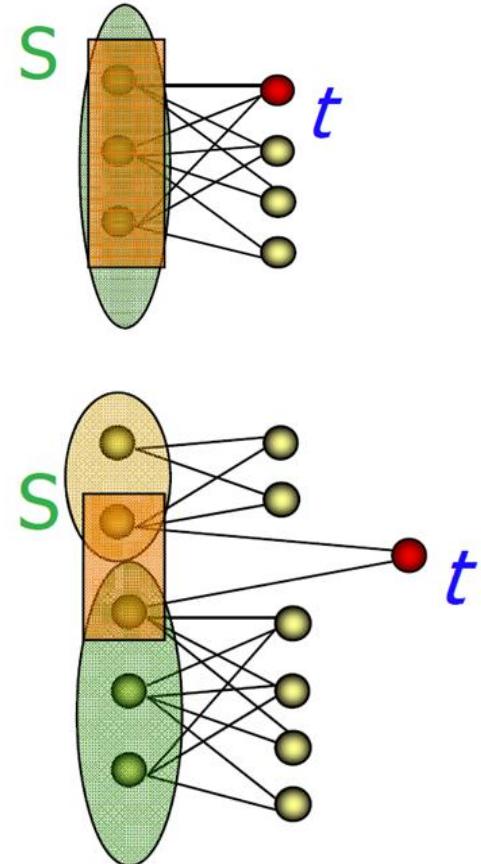
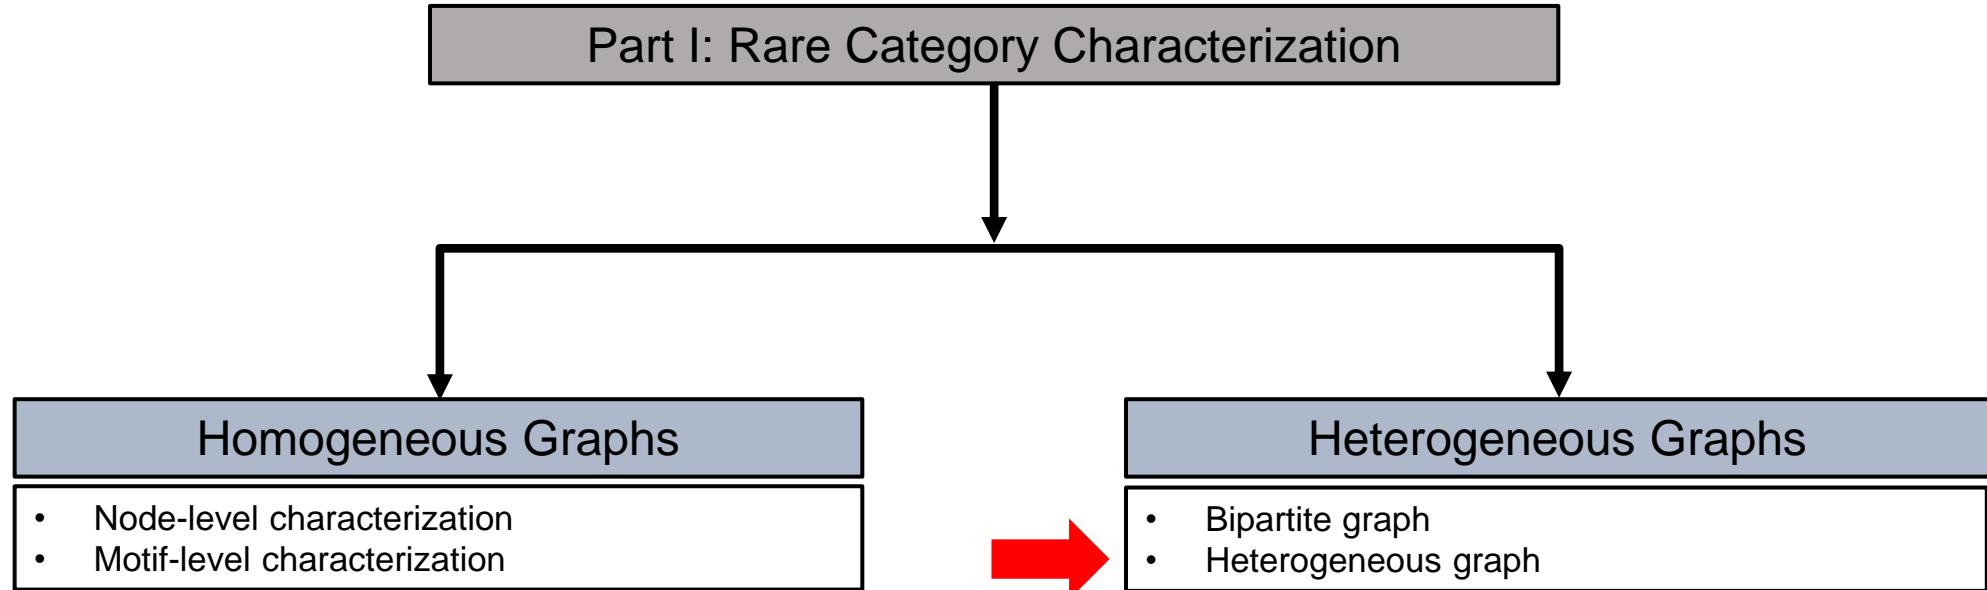


Figure adopted from [Sun et al., ICDM 2005]

- Sun, Jimeng, et al. “Neighborhood formation and anomaly detection in bipartite graphs.” ICDM, 2005.

Rare Category Characterization



Rare Category in Heterogeneous Graphs

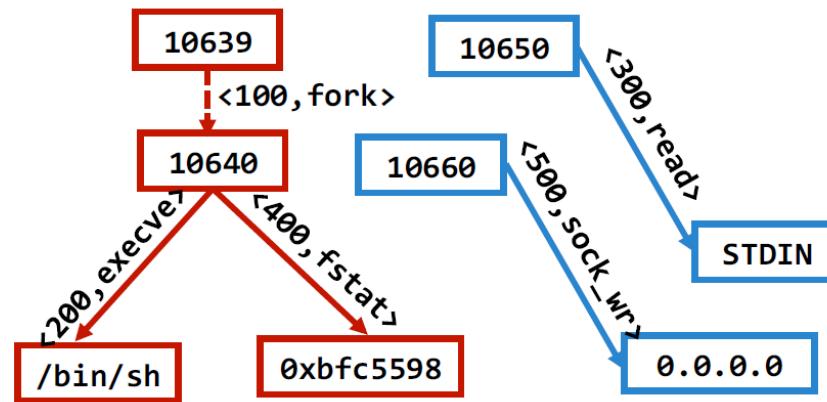
□ Problem 1.7. Rare category characterization in dynamic heterogeneous graphs

- **Given:** A stream of heterogeneous graphs $G^{(t)} = (V, E, T)$ containing different types T of nodes V and edges E .
- **Find:** Rare categories in real-time while consuming bounded memory.

□ Challenges

- Nodes and edges are typed (e.g., fork, read).
- Graph evolves from a steam of typed edges.
- Bounded space and time complexity

Example of two information flow graphs based on system logs

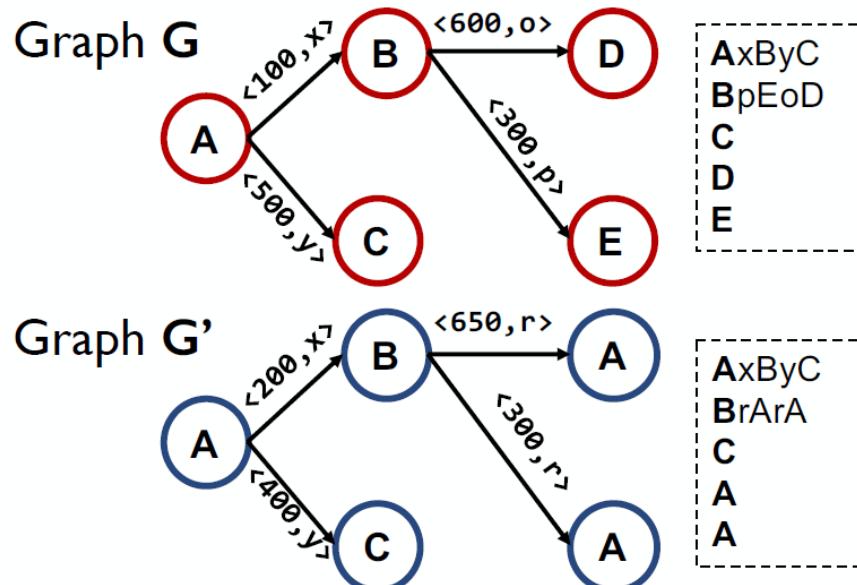


• Manzoor, Emaad, et al. “Fast Memory-efficient Anomaly Detection in Streaming Heterogeneous Graphs.” KDD 2016.

Graph Representation

□ Graph to vectors via shingling

- **Key idea:** encode the heterogeneous graph into a low-dimensional representation.
- Compute the shingle vector for each graph $G^{(t)}$.
- Contain the frequencies of each k -shingle in $G^{(t)}$.



	\mathbf{z}_G	$\mathbf{z}_{G'}$
A	0	2
C	1	1
D	1	0
E	1	0
AxByC	1	1
BrArA	0	1
BpEoD	1	0

Figure adopted from [Manzoor et al., KDD 2016]

- Manzoor, Emaad, et al. “Fast Memory-efficient Anomaly Detection in Streaming Heterogeneous Graphs.” KDD 2016.

Identification of Rare Categories

□ Sketching graphs

- Shingle universe is large and unknown
- Compute L -dimension projection vector from shingle vector via Locality-Sensitive Hashing (LSH)

□ Detection

- Bootstrap K Clusters
- Cluster Centroid: “Average” graph
- Update Clusters: Constant time
- Anomaly Score: Nearest centroid

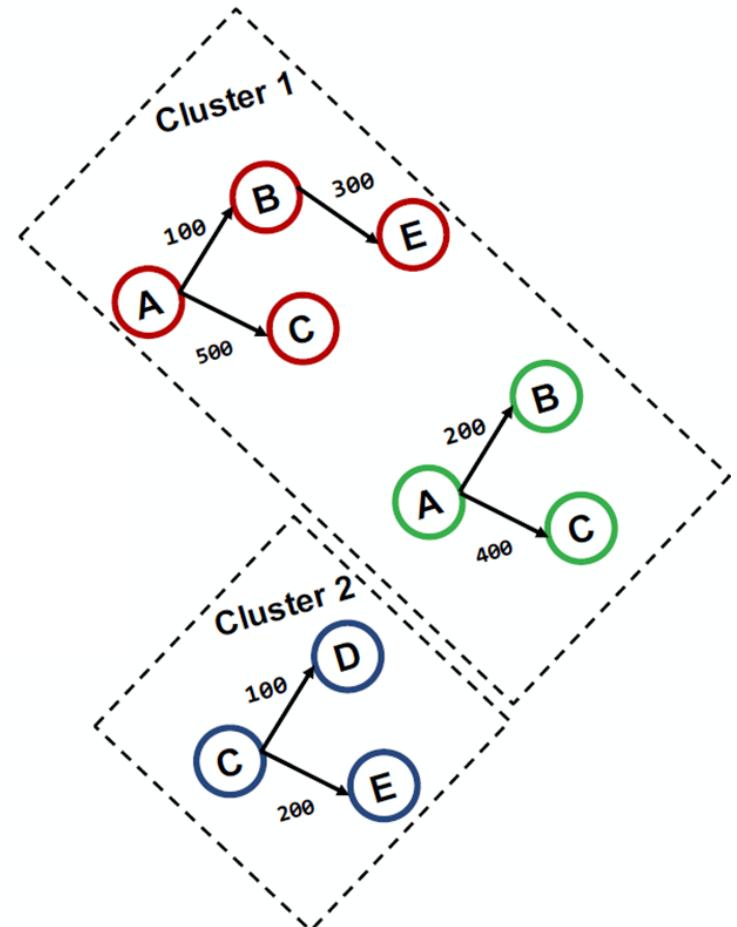


Figure adopted from [Manzoor et al., KDD 2016]

- Manzoor, Emaad, et al. “Fast Memory-efficient Anomaly Detection in Streaming Heterogeneous Graphs.” KDD 2016.

Additional References

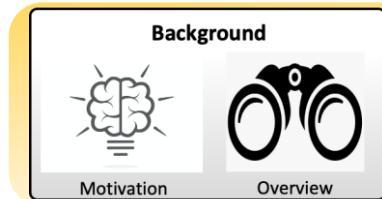
□ Clustering-based approaches

- Nibble [Spielman et al., 2013]
- TSC [Benson et al., 2015]
- TECTONIC [Tsourakakis et al., 2017]
- MOTLOC [Zhou et al., 2018]
- TGS [Garranza et al., 2020]
- ...

□ Graph neural network-based approaches

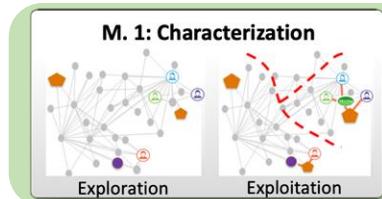
- SPNN [Meng et al., 2018]
- hpGAT [Liu et al., 2019]
- FI-GNNs [Ding et al., 2019]
- Dominant [Ding et al., 2019]
- DevNet2 [Pang et al., 2019]
- ...

Roadmap



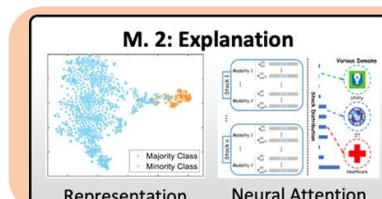
Background

- Motivation
- Research overview



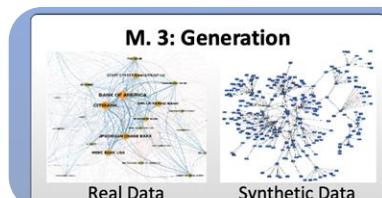
I. Rare Category Characterization

- Rare category characterization on homogeneous graphs
- Rare category characterization on heterogeneous graphs



II. Rare Category Explanation

- Data insights
- Model insights



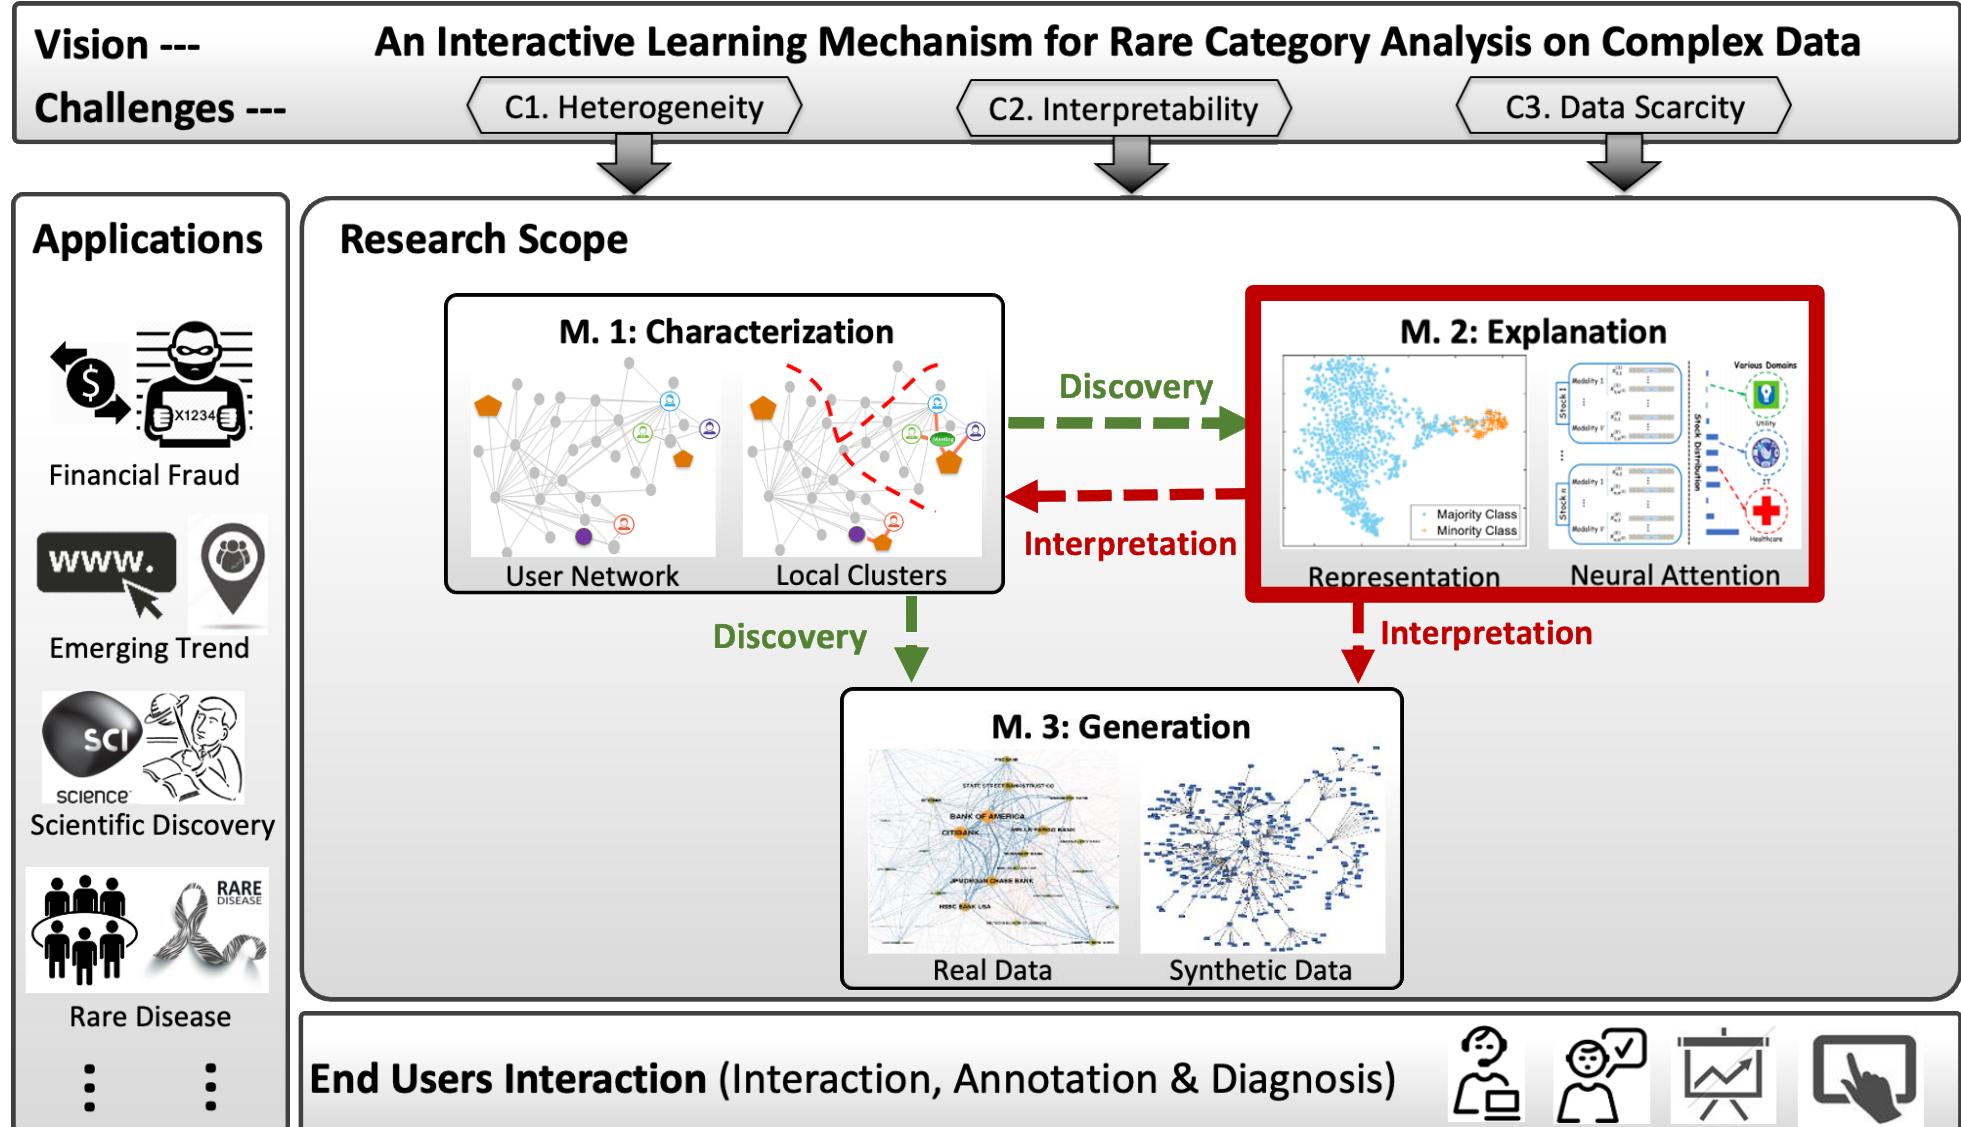
III. Rare Category Generation

- Unsupervised rare category generation
- Supervised rare category generation



IV. Real-world Application

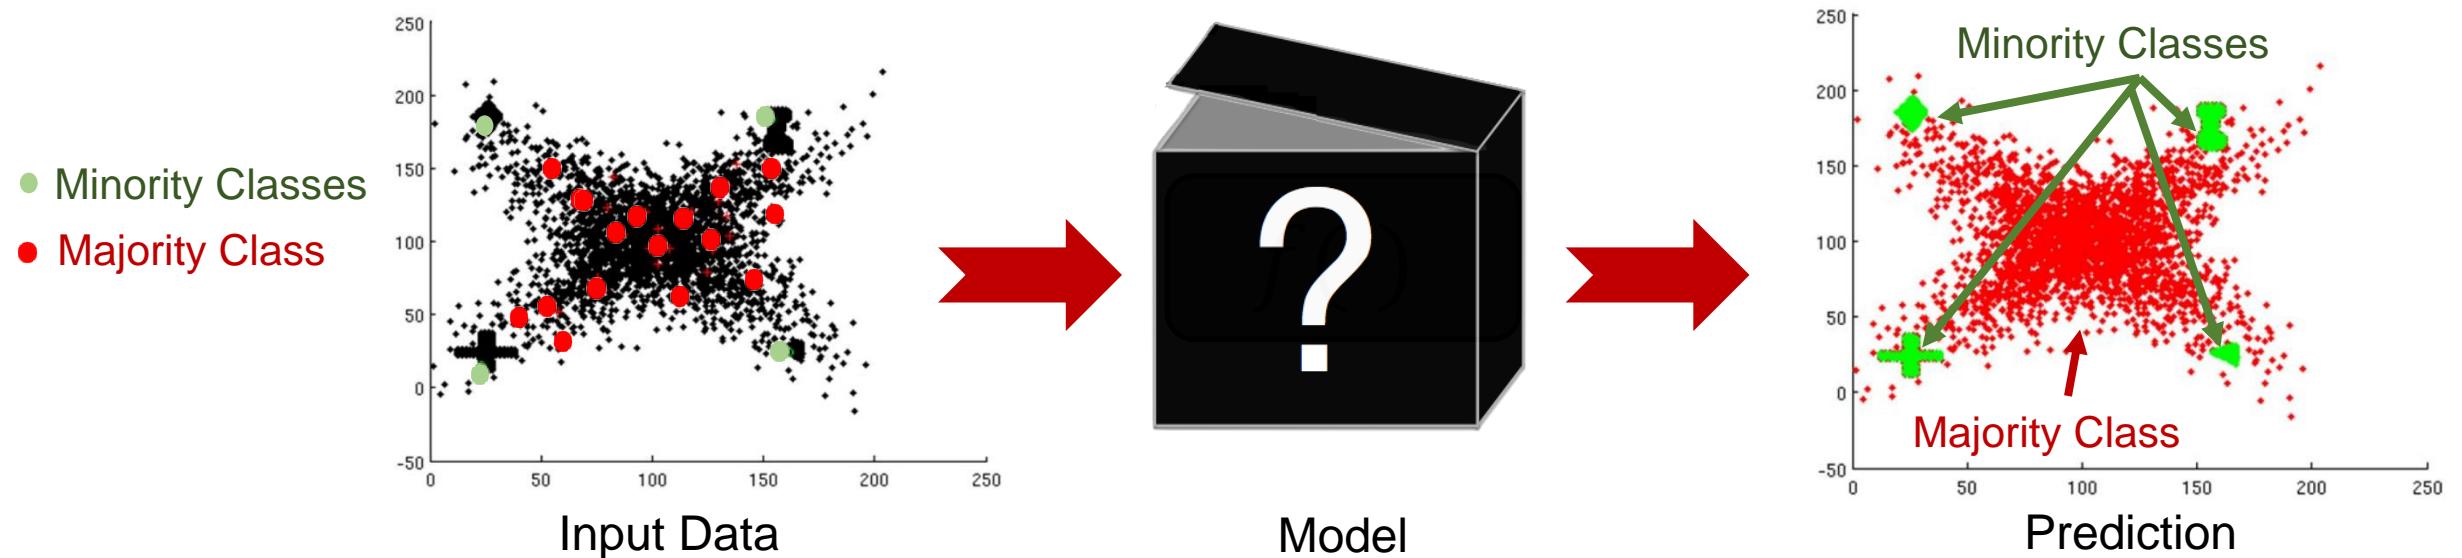
Rare Category Explanation



Rare Category Explanation

□ Problem 2.1. Rare category explanation

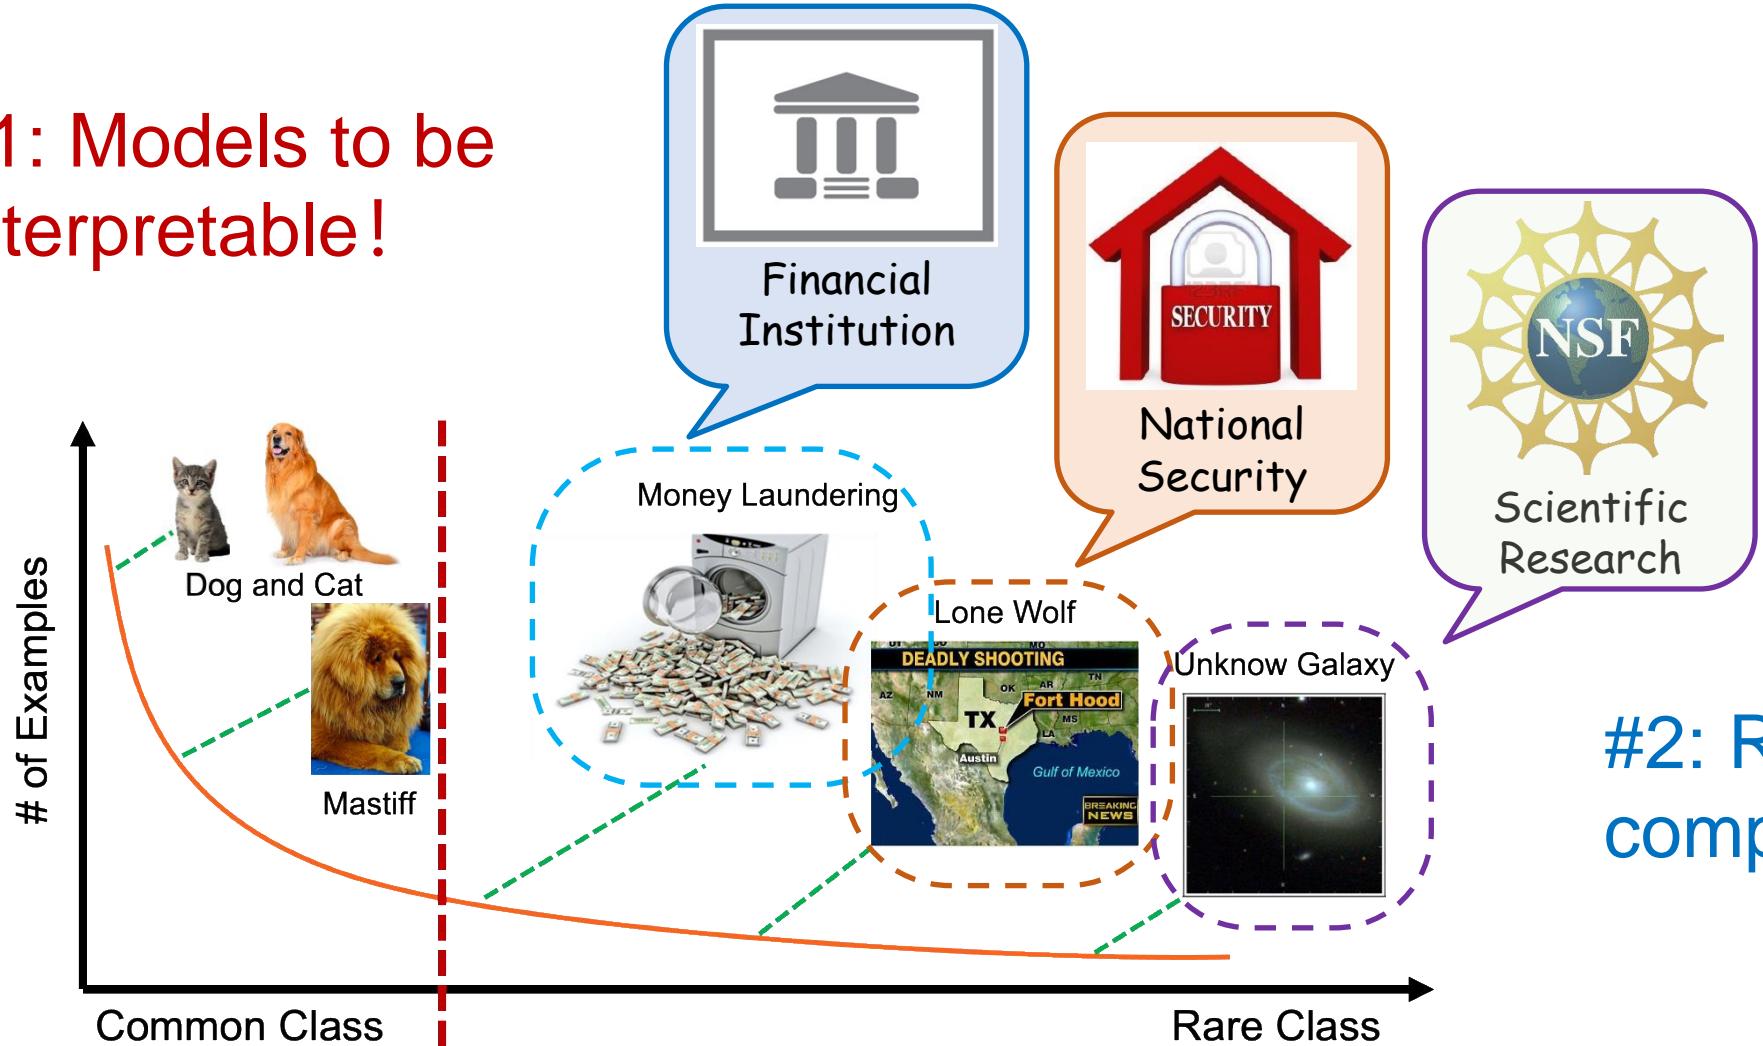
- **Given:** (1) A data set $D = \{x_1, \dots, x_n\}$, (2) a prediction model $f(x)$ for rare category analysis in D .
- **Find:** Interpretation over the observed data D and the prediction model $f(x)$.



- Pelleg, Dan, and Andrew W. Moore. "Active learning for anomaly and rare-category detection." Advances in neural information processing systems. 2005.
- Zhou, Dawei, et al. "A local algorithm for structure-preserving graph cut." ACM SIGKDD. 2017.

Rare Category Explanation

#1: Models to be interpretable!

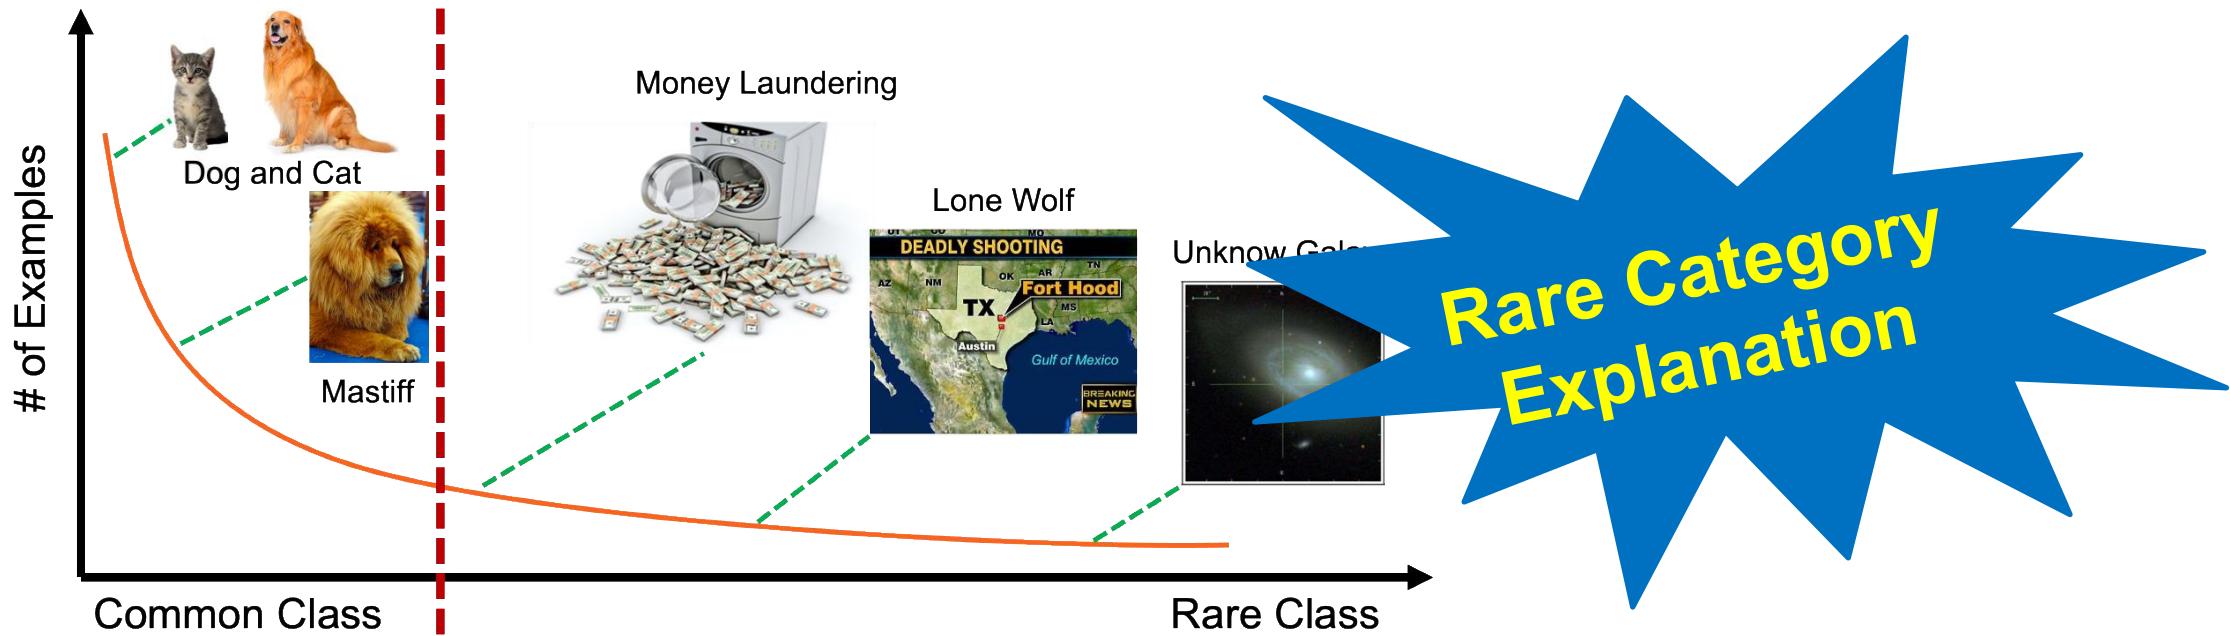


- Abazajian, Kevork N., et al. "The seventh data release of the Sloan Digital Sky Survey." *The Astrophysical Journal Supplement Series*. 2009.
- Spaaij, R.. "The enigma of lone wolf terrorism: An assessment." *Studies in Conflict & Terrorism*. 2010.
- Fich, Eliezer M., and Anil Shivdasani. "Financial fraud, director reputation, and shareholder wealth." *Journal of financial Economics*. 2007.

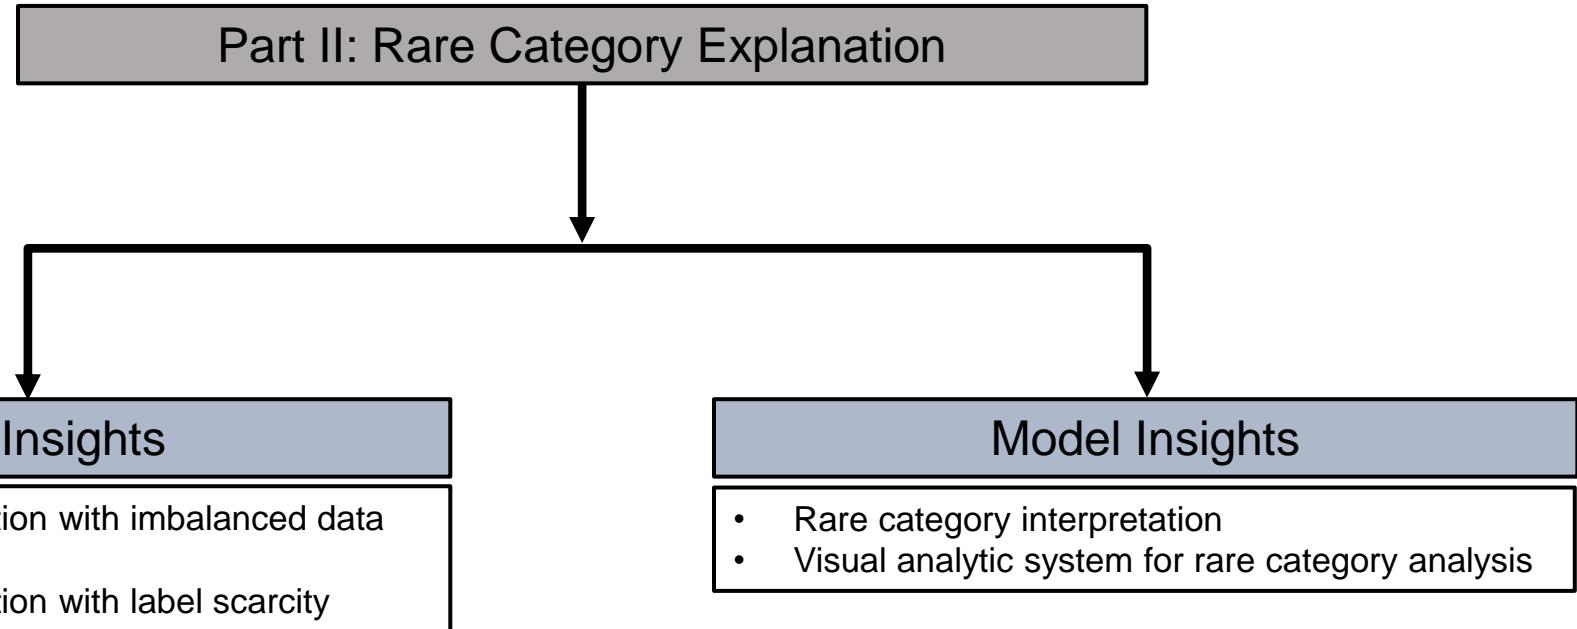
Research Questions

□ Research questions

- **Q 2.1. Data insights**
 - How is the data **distributed**?
- **Q 2.2. Model insights**
 - Why does the model make a **certain prediction** on a particular piece of information?



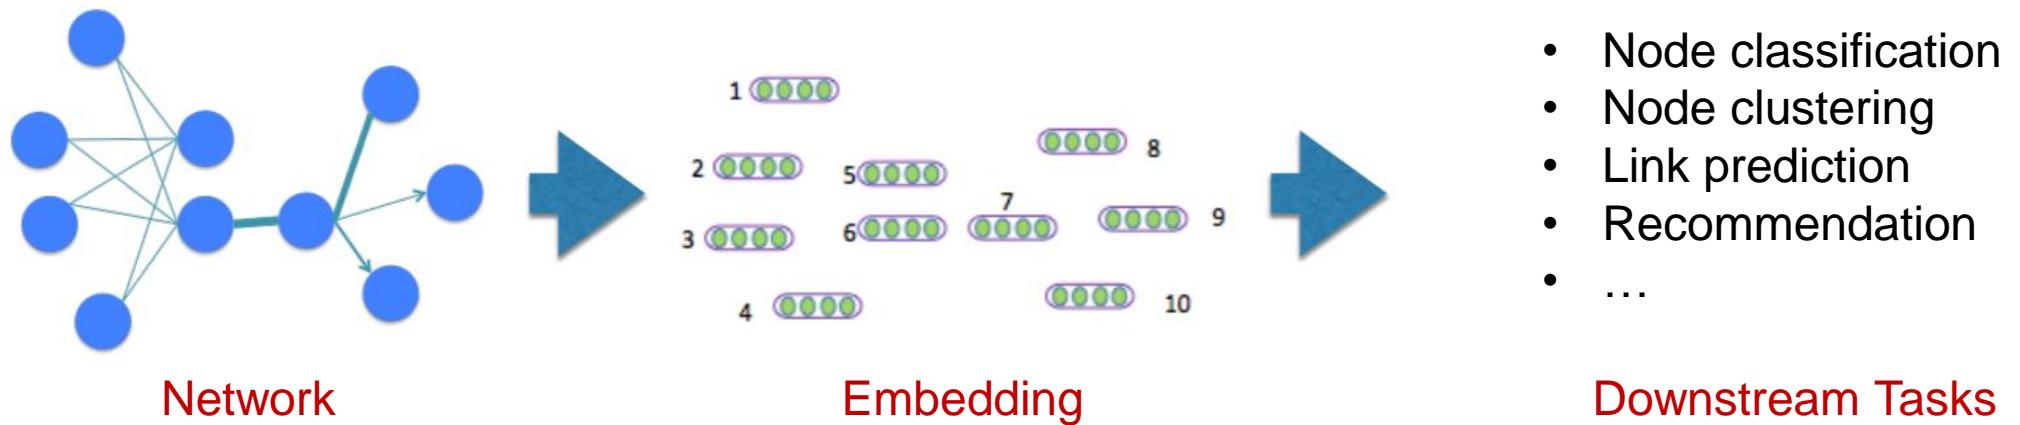
Rare Category Explanation



Network Representation Learning

□ Learning node representation for networks

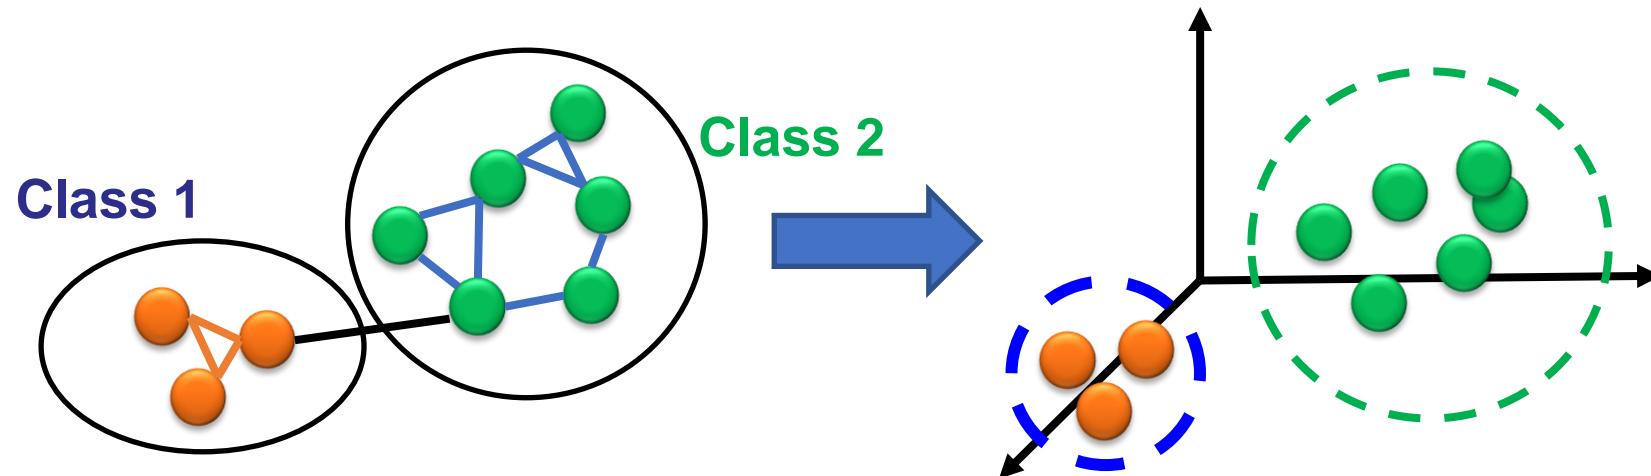
- **Goal:** Represent each node with a vector representation that preserves the structure of networks.
- **Methods:** DeepWalk [Perozzi et al., 2014], Node2vec [Grover et al., 2016], LINE [Tang et al., 2015], TransE [Bordes et al., 2013], GCN [Kipf et al., 2017], GAT [Veličković et al., 2017], GraphSAGE [Hamilton et al, 2016], etc.
- Unable to handle **highly-skewed** data distribution.



Network Representation for Imbalanced Data

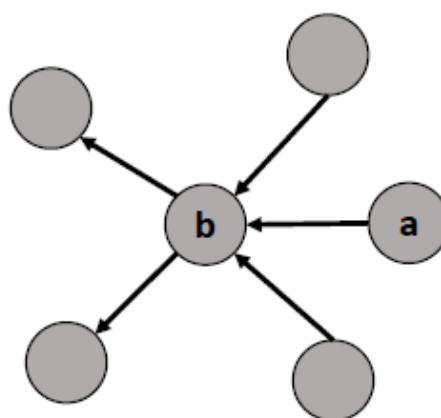
□ Problem 2.2. Imbalanced network representation

- **Given:** (1) a (directed or undirected) network $G=(V,E)$, (2) imbalanced class labels for nodes in V , (3) embedding dimension d .
- **Find:** a low-dimensional vector representation for each node $v \in V$, so that the minority class is separated from majority class in the embedding feature space.

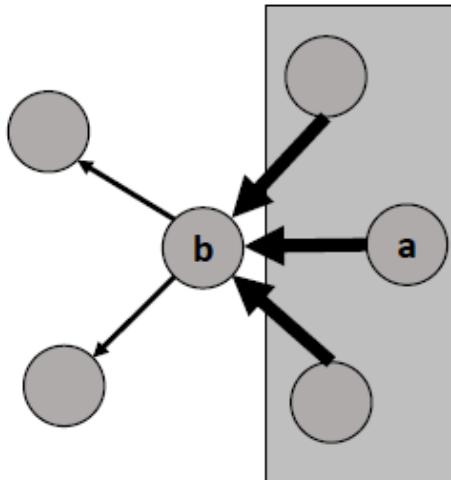


Vertex-Diminished Random Walk

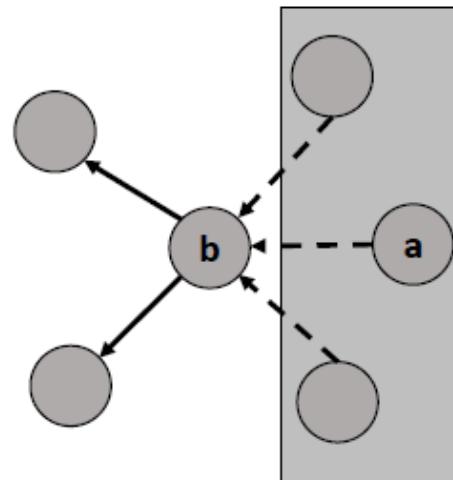
□ **Intuition:** the probability of a transition to one node would decrease each time when it is visited.



(a) Random Walk



(b) VRRW



(c) VDRW

→ *Unchanged transition probability*

→ *Enhanced transition probability*

→ *Weakened transition probability*

Affected area

$$f(S_j(t)) = 1$$

$$f(S_j(t)) = S_j(t)$$

$$f(S_j(t)) = \alpha^{S_j(t)}$$

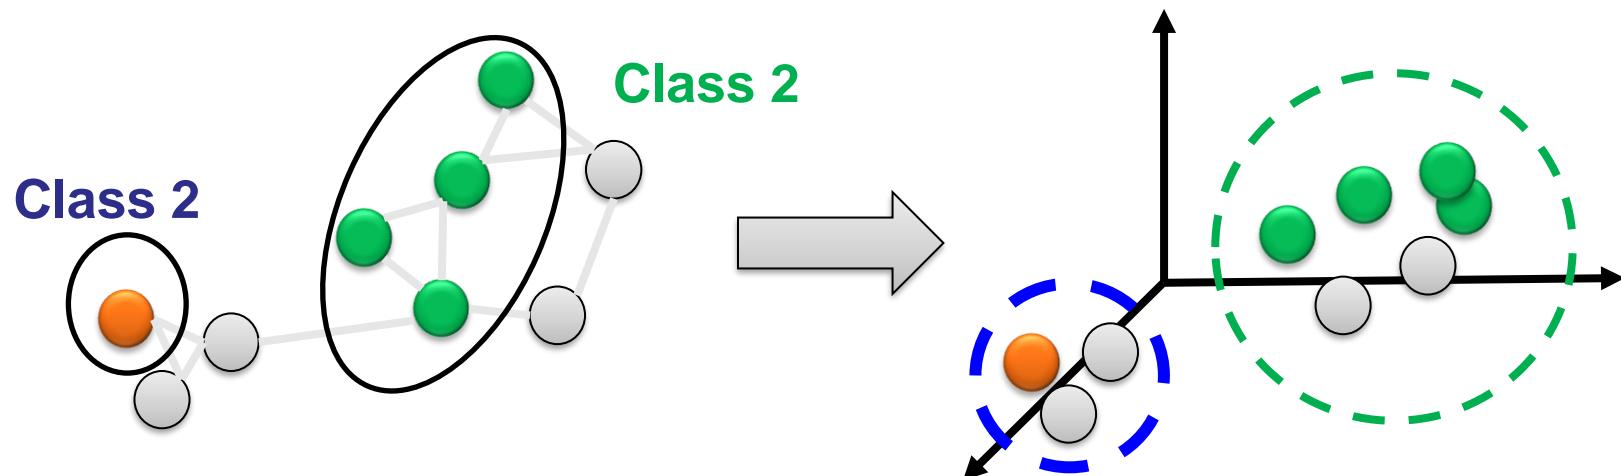
- Wu, Jun, et al. ImVerde: Vertex-Diminished Random Walk for Learning Imbalanced Network Representation. IEEE Big Data, 2018.
- Robin Pemantle. "Vertex-reinforced random walk." Probability Theory and Related Fields, 1992.

Imbalanced Network Embedding

□ The loss of the proposed semi-supervised framework:

$$\mathcal{L} = \mathcal{L}_s + \lambda \mathcal{L}_u$$

- \mathcal{L}_s : supervised loss of predicting the label
- \mathcal{L}_u : unsupervised loss of predicting the graph context
- **Assumption**: nodes with the same context would have the similar embeddings





Imbalanced Network Embedding

□ Unsupervised loss

$$\mathcal{L}_u = - \sum_{(v_i, v_c)} [\log \sigma(w_c^T e_i) + k \cdot \mathbb{E}_{v_n \sim P_n(v_c)} \log \sigma(-w_n^T e_i)]$$

Positive pair Negative pair

(v_i, v_c) :	a node-context pair of node i and its context c sampled from path p using the proposed VDRW
e_i :	embedding vector of node i
w_c :	representation of context c
v_n :	a randomly selected negative sample

□ Supervised loss (cross-entropy error)

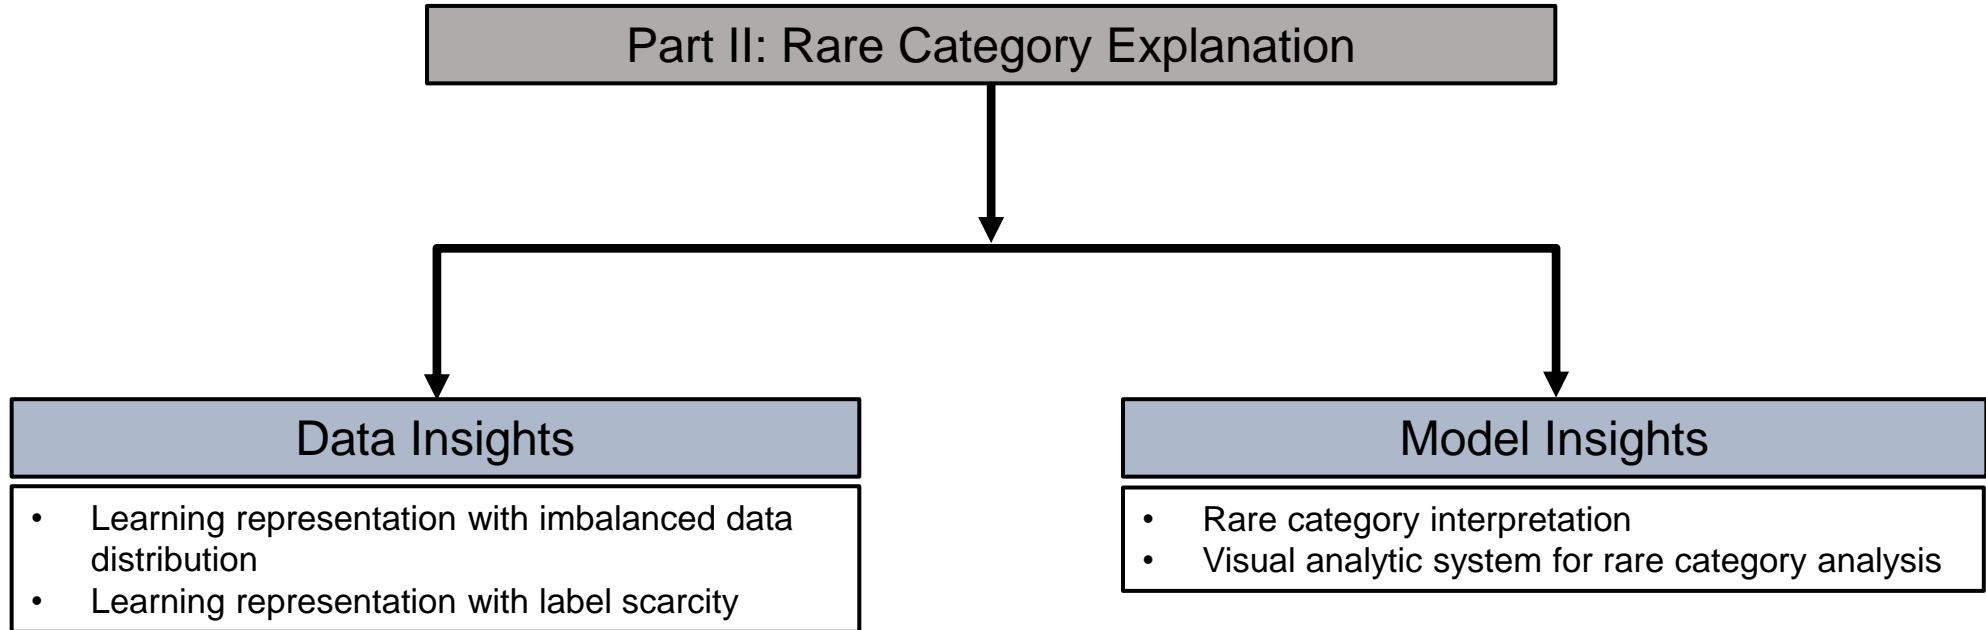
$$\mathcal{L}_s = -\frac{1}{l} \sum_{i=1}^l \log \frac{\exp([h^{l1}(x_i) h^{l2}(e_i)]^T y_i)}{\sum_{y'} \exp([h^{l1}(x_i) h^{l2}(e_i)]^T y')}$$

Label information

l_1, l_2 :	number of hidden layers in neural networks
y_i :	class label
x_i :	feature vector of node i

- Wu, Jun, et al. ImVerde: Vertex-Diminished Random Walk for Learning Imbalanced Network Representation. IEEE Big Data, 2018.

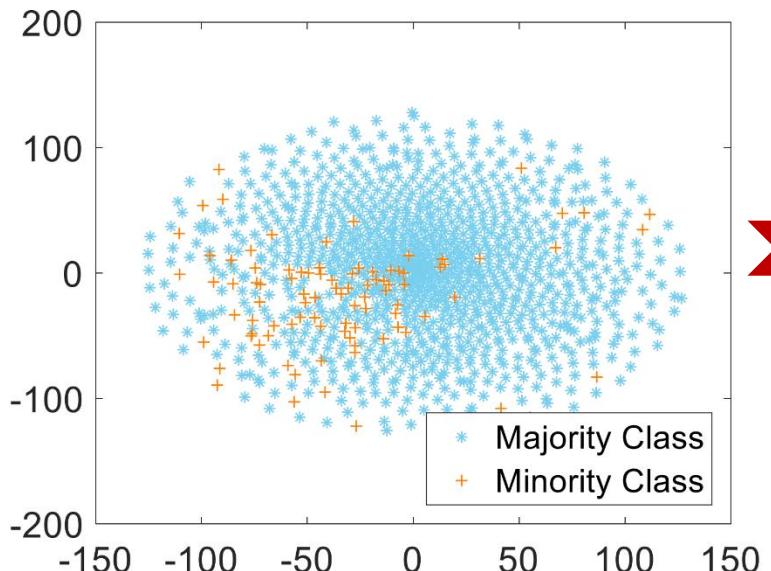
Rare Category Explanation



Learning Representation with Label Scarcity

□ Problem 2.3. Rare-category-oriented network representation

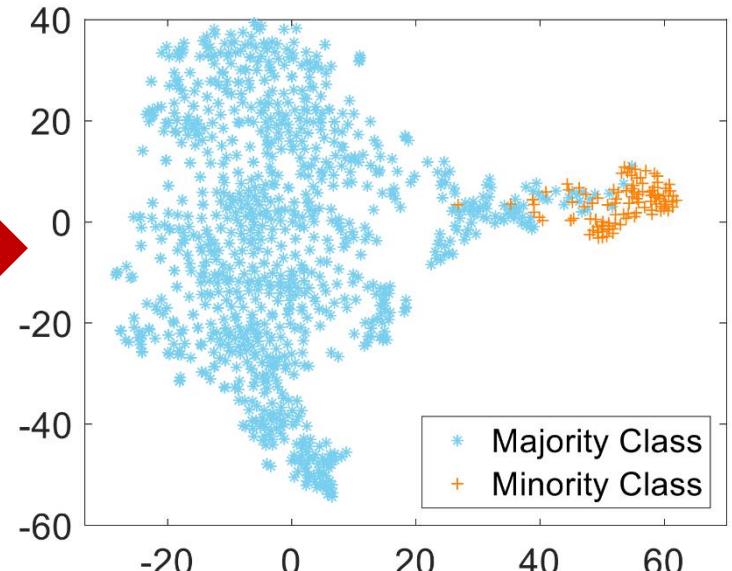
- **Given:** (1) An attributed network $G = (V, E, X)$, (2) one-shot or few-shot labeled rare examples $L = \{x_1, \dots, x_L\}$, and (3) the desired embedding dimension d .
- **Find:** (1) A rare-category-oriented network embedding $E \in R^{n \times d}$, and (2) a list of predicted rare category examples.



(a) Original Feature Space



$$f(\cdot)$$



(b) Embedding Space

- Zhou, Dawei, et al. "Spac: Self-paced network representation for few-shot rare category characterization." ACM SIGKDD. 2018.

SPARC – Algorithm

□ The overall optimization problem

- L denotes the set of labeled examples; U denotes the set of unlabeled examples.
- $X = \{x_1, \dots, x_n\}$ denotes node attributes.
- $E = \{e_1, \dots, e_n\}$ denotes the learned node embedding.
- θ denotes the neural network parameters, and C denotes the number of classes.

$$\arg \min_{E, v, \theta} : \mathcal{L} = \underbrace{\sum_{i=1}^L \sum_{c=1}^C c_{y_i, \hat{y}_i} \log \Pr(\hat{y}_i = c | x_i, e_i)}_{\mathcal{L}_{RCC}: \text{Rare category characterization}} + \underbrace{\sum_{i=1}^{L+U} \left[\sum_{c=1}^C \lambda^{(c)} v_i^{(c)} + \lambda^{(0)} v_i^{(0)} \right]}_{\mathcal{L}_{co}: \text{Self-paced learning}} - \alpha \sum_{i=1}^{L+U} \sum_{c=1}^C v_i^{(c)} v_i^{(0)} + \underbrace{\sum_{i=1}^{L+U} \sum_{c=1}^C v_i^{(c)} \log \Pr(\hat{y}_i = c | x_i, e_i) + v_i^{(0)} \log E_{(i,c,\gamma)} \log \sigma(\gamma \theta_c^T e_i)}_{\mathcal{L}_{RCE}: \text{Rare category embedding}}$$

SPARC – Algorithm

□ Intuitions

- Rare category characterization maximizes the probability of assigning true label to the prediction for both labeled data and unlabeled data.
- Self-paced learning starts with the easy concept and gradually learn the hard concept.
- Rare category embedding minimizes the prediction loss regarding the sampled graph context pairs.

$$\arg \min_{E, v, \theta} \mathcal{L} = \sum_{i=1}^L \sum_{c=1}^C c_{y_i, \hat{y}_i} \log \Pr(\hat{y}_i = c | \mathbf{x}_i, \mathbf{e}_i) + \mathcal{L}_{co}: \text{Self-paced learning}$$

$$+ \sum_{i=1}^{L+U} \left[\sum_{c=1}^C \lambda^{(c)} v_i^{(c)} + \lambda^{(0)} v_i^{(0)} \right] - \alpha \sum_{i=1}^{L+U} \sum_{c=1}^C v_i^{(c)} v_i^{(0)}$$

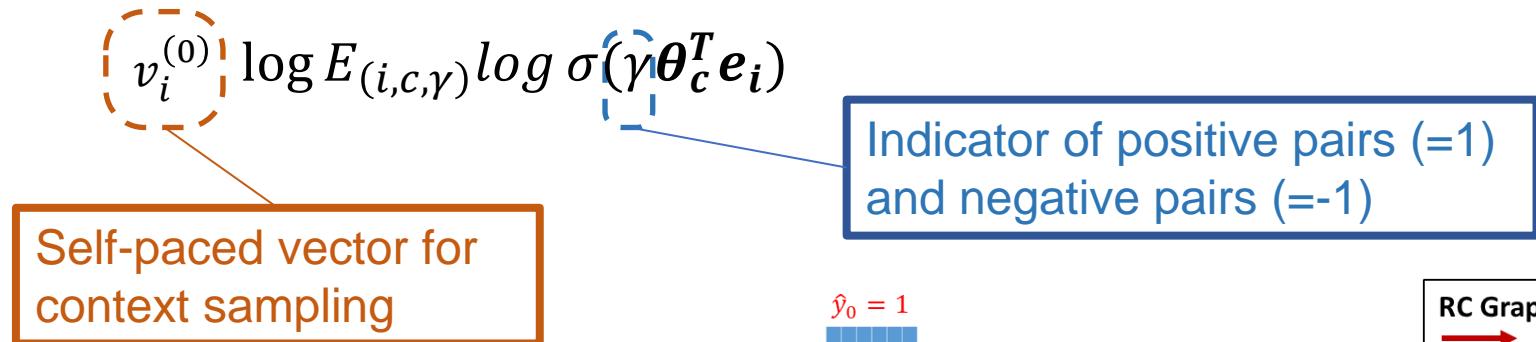
\mathcal{L}_{RCC} : Rare category characterization \mathcal{L}_{RCE} : Rare category embedding

$$+ \sum_{i=1}^{L+U} \sum_{c=1}^C v_i^{(c)} \log \Pr(\hat{y}_i = c | \mathbf{x}_i, \mathbf{e}_i) + v_i^{(0)} \log E_{(i,c,\gamma)} \log \sigma(\gamma \boldsymbol{\theta}_c^T \mathbf{e}_i)$$

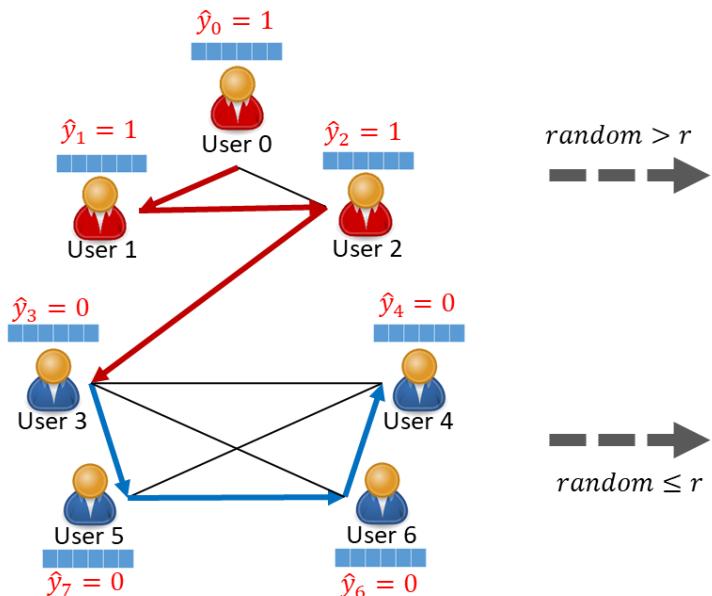
SPARC – Algorithm

□ Rare category embedding

- Minimize the cross-entropy loss of predicting context pairs (i, c) .



- Self-paced dual network context sampling.
 - Extracting two types of context pairs (i, c) .
 - With probability p , extracting general context pairs.
 - With probability $1-p$, extracting rare-category oriented context pairs via the non-zero elements in $v^{(0)}$.



RC Graph Context

Positive Pairs:

- $(u_0, u_2, +1)$

Negative Pairs:

- $(u_0, u_5, -1)$
- $(u_2, u_6, -1)$

General Graph Context

Positive Pairs:

- $(u_3, u_5, +1)$

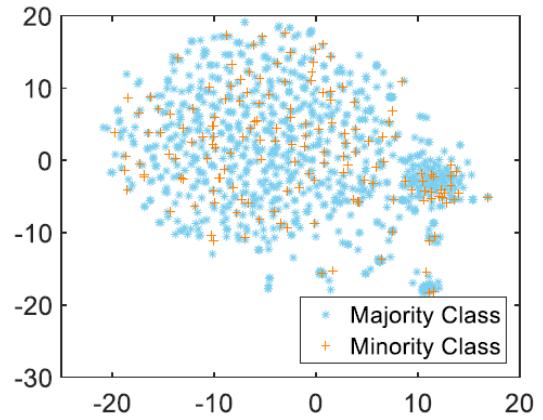
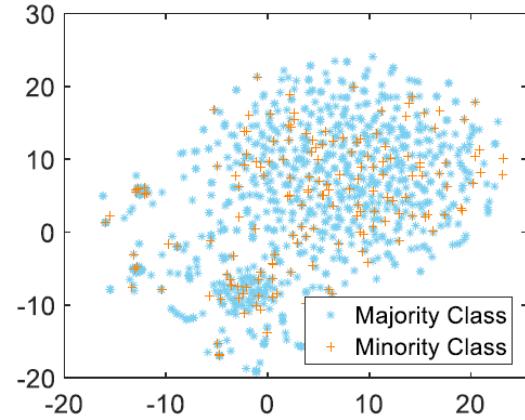
Negative Pairs:

- $(u_5, u_0, -1)$
- $(u_6, u_1, -1)$

Experimental Results

□ Network layout visualization

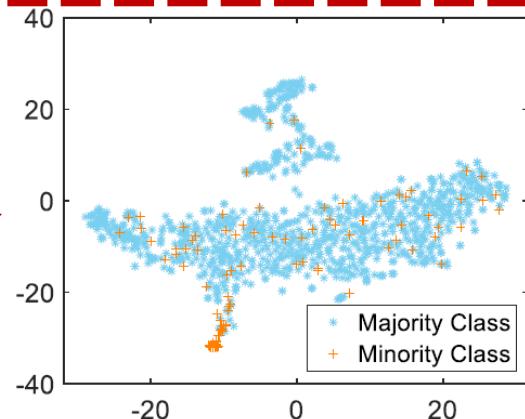
Unsupervised network embedding



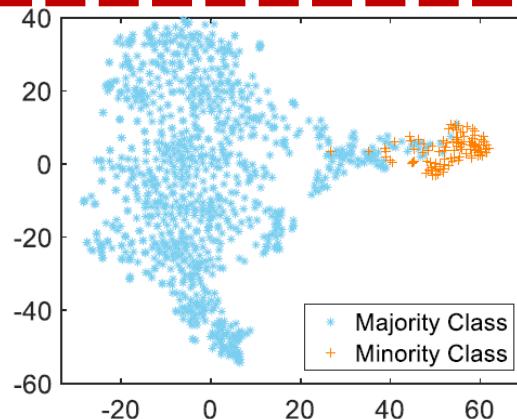
(a) DeepWalk

Label-informed network embedding

Rare category is not well characterized.



(c) PLANETOID

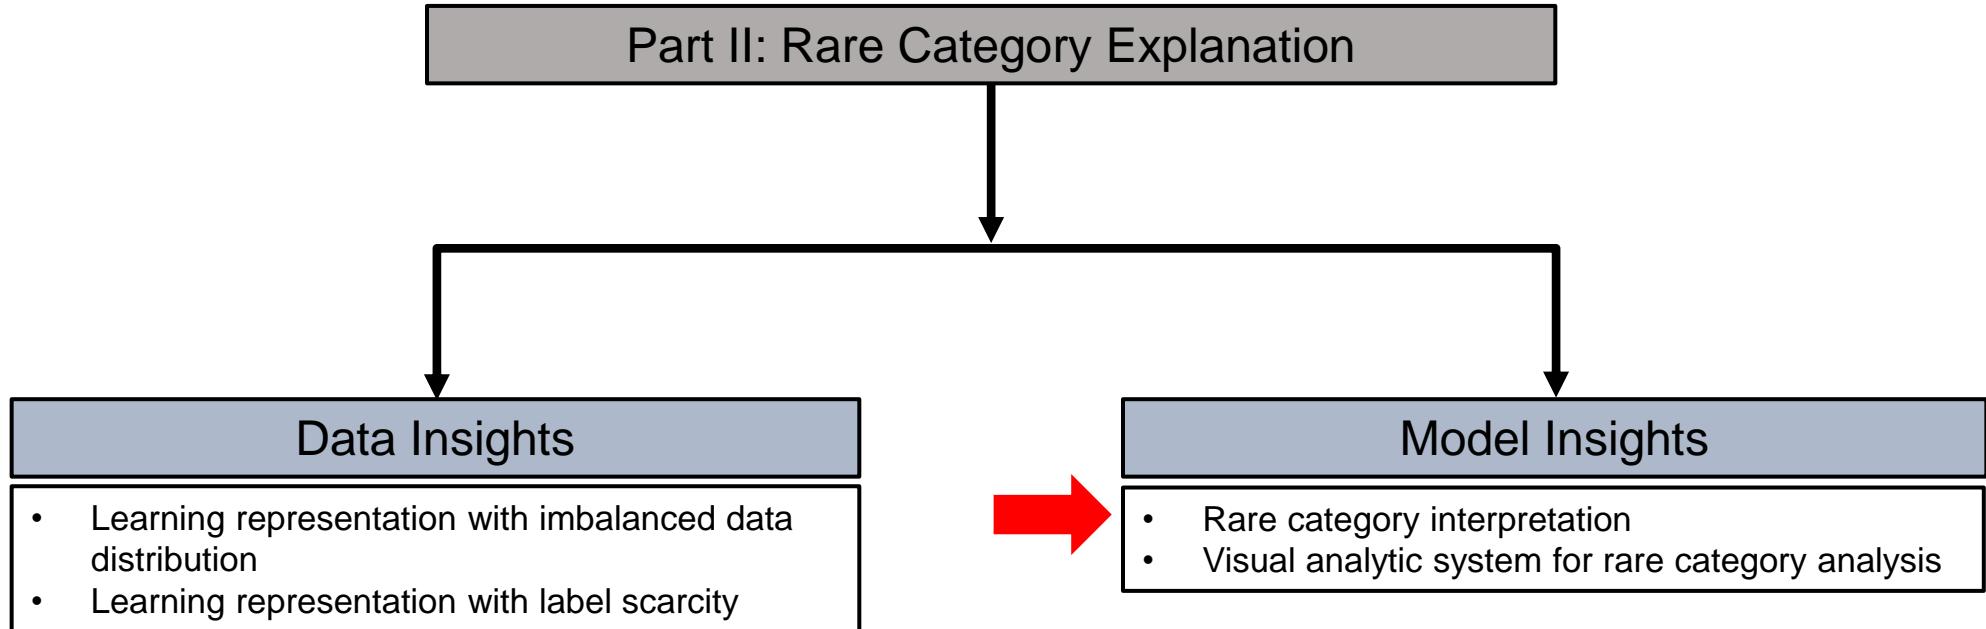


(d) SPARC

Rare category is well separated from the background!

- Zhou, Dawei, et al. "Spac: Self-paced network representation for few-shot rare category characterization." ACM SIGKDD. 2018.
- Yang, Zhilin, et al. "Revisiting semi-supervised learning with graph embeddings." ICML. 2016.

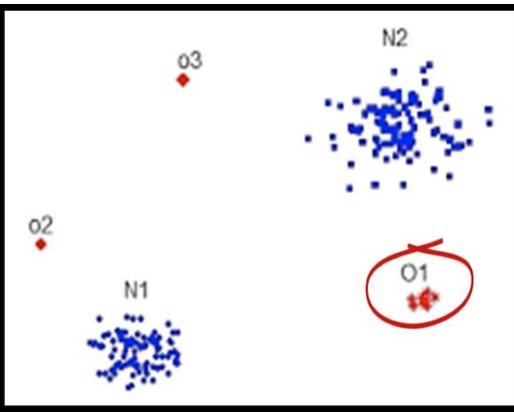
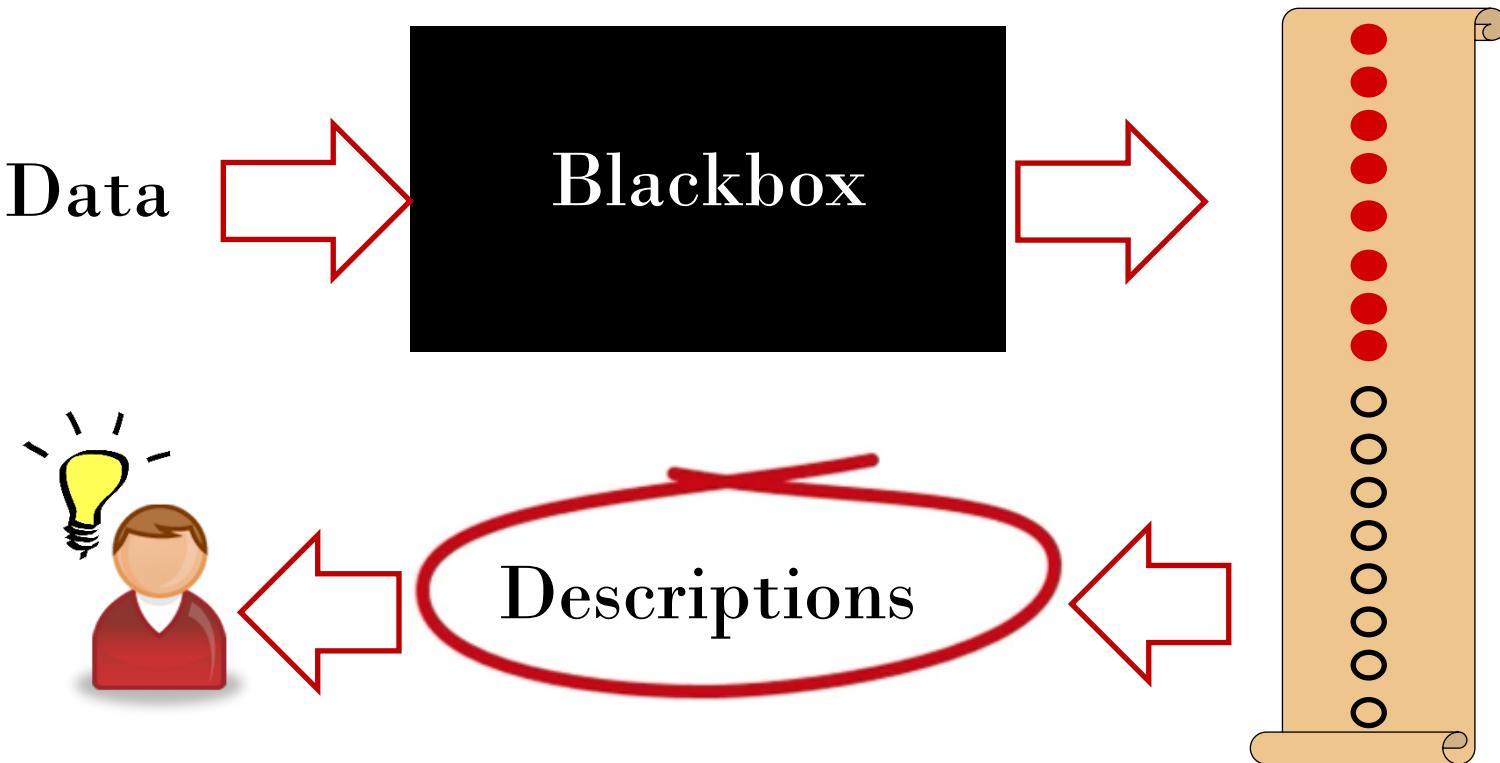
Rare Category Explanation



Rare Category Interpretation

□ Problem 2.4. Rare category description

- **Given:** data samples $X = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$.
- **Find:** The interpretation for each rare categories.

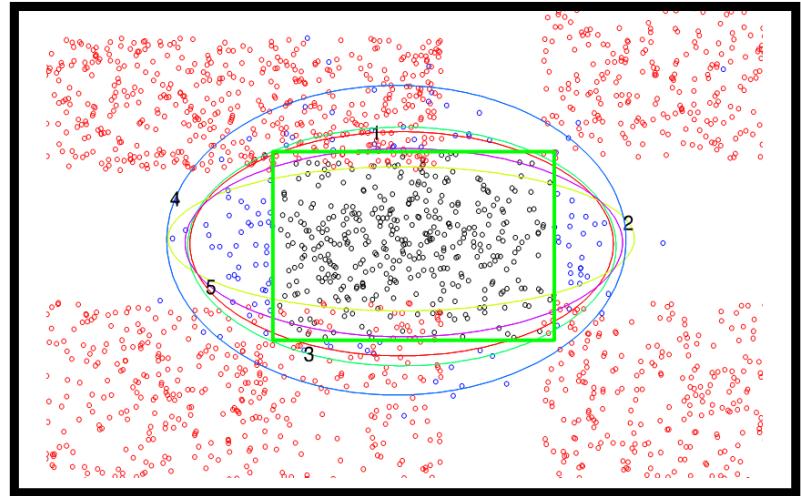


Grouped Outliers =
Rare Category

x-PACS – Algorithm

□ Step 1: Subspace Clustering

- Hierarchically combine dense clusters.
 - **Threshold mass**: min # of anomalies in cluster.
- Output dense & pure subspace clusters.
 - **Threshold purity**: max # of normal pts in the clusters.



□ Step 2: Rectangles to ellipsoids

- Hyper-rectangles are box-like, can refine into flexible shapes: hyper-ellipsoids.
- Recovering other nearby anomalies.

black: anomalies
in subspace cluster
blue: “nearby”
anomalies
red: normal pts

$$p(\mathbf{c}, \mathbf{M}) = \{\mathbf{x} \mid (\mathbf{x} - \mathbf{c})^T \mathbf{M}^{-1} (\mathbf{x} - \mathbf{c}) \leq 1\}$$

• Macha, Meghanath et al. "Explaining Anomalies in Groups with Characterizing Subspace Rules." DMKD. 2018.

x-PACS – Algorithm

□ Step 3: Transmission

- An Encoding Scheme:



Goal: transmit information of
which samples are anomalies,
using as few bits as possible

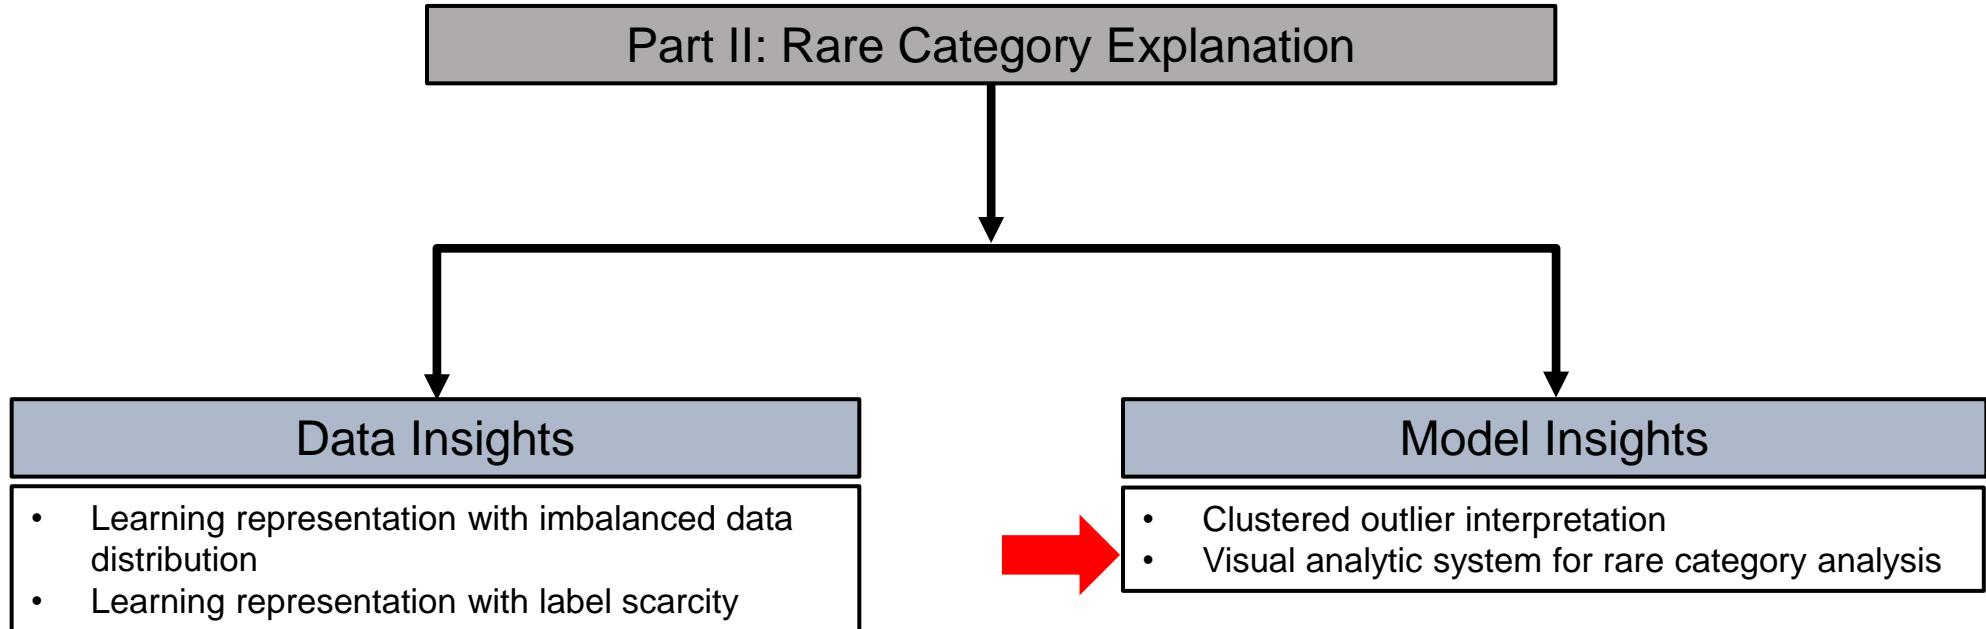


- Generating Rules to Output

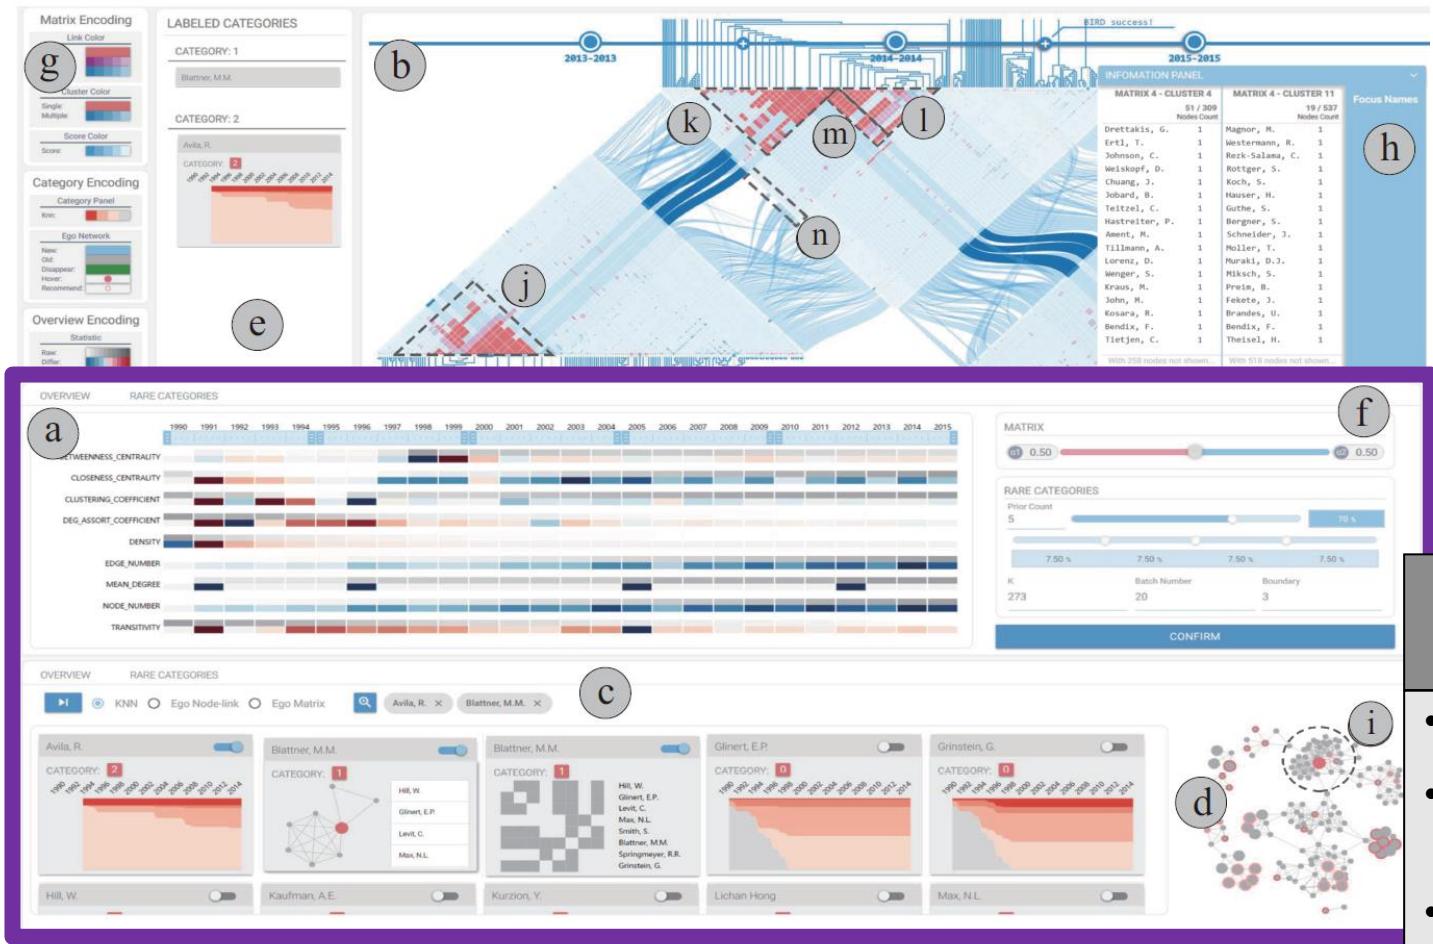
DEFINITION . (FEATURE RULES) *Given a axis-aligned ellipsoid $p(\mathbf{c}, \mathbf{M})$ in a subspace $f_{t_1} \times \dots \times f_{t_d}$, a rule on feature t_z is an interval $(\mathbf{c}[z] - \text{radius}_z, \mathbf{c}[z] + \text{radius}_z)$*

- Macha, Meghanath et al. "Explaining Anomalies in Groups with Characterizing Subspace Rules." DMKD. 2018.

Rare Category Explanation



RCANALYZER Prototype System



Data Exploration Module

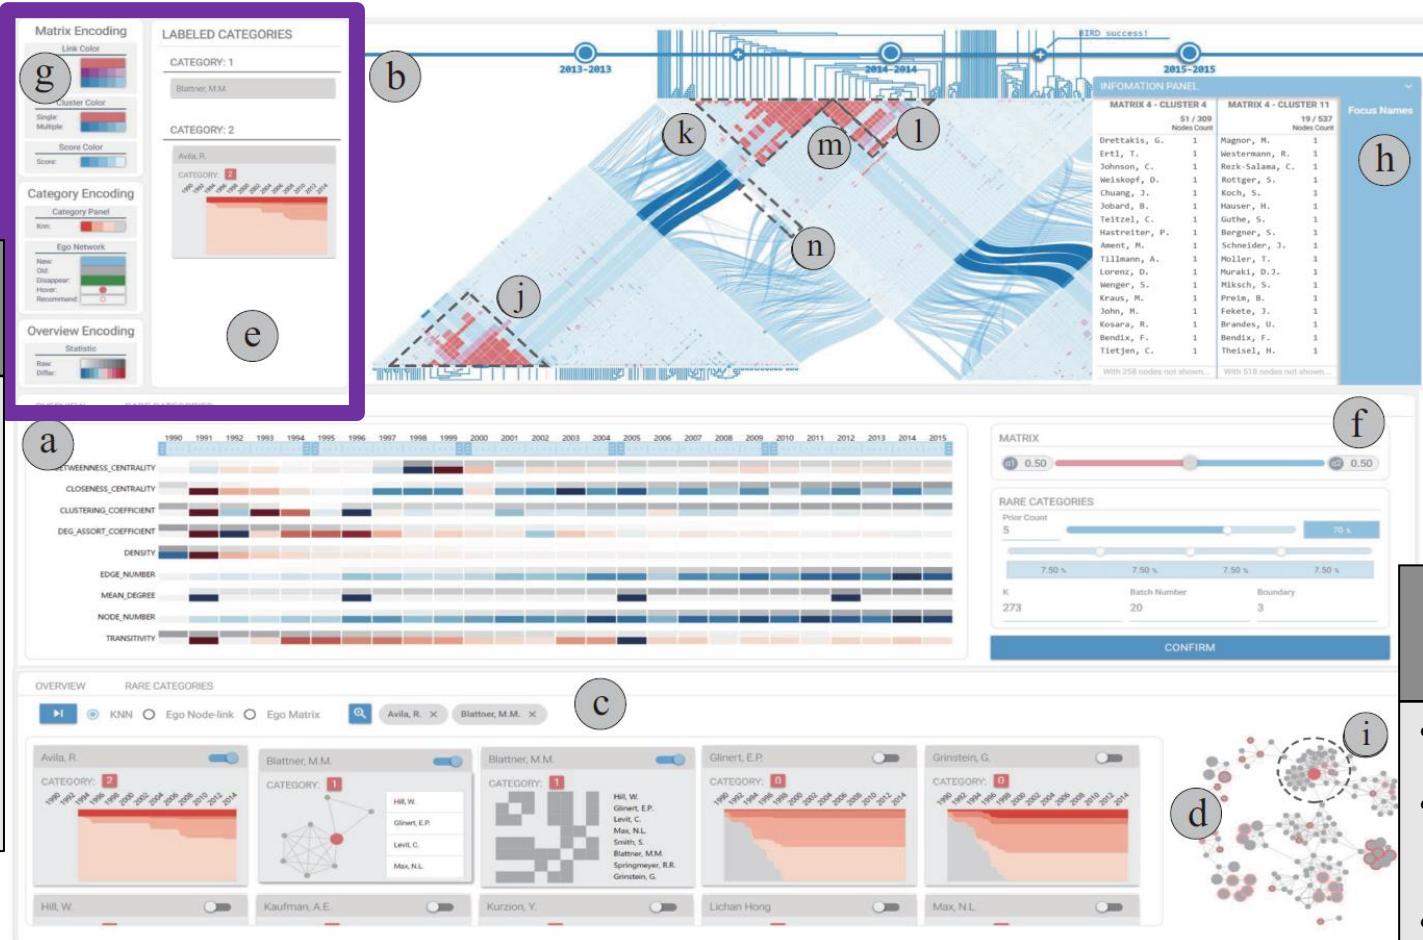
- Represent raw data
- Interactive visualization for data querying
- Support data filtering

- Pan, Jia-cheng, et al. "RCAnalyzer: visual analytics of rare categories in dynamic networks." Frontiers of Information Technology & Electronic Engineering. 2020.
- Zhou, Dawei, et al. "Discovering rare categories from graph streams." Data mining and knowledge discovery. 2017.

RCANALYZER Prototype System

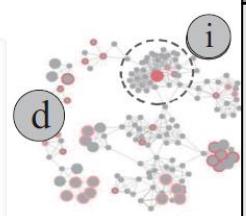
Feature Selection Module

- Visualize the variance of data
- Visualize the correlation of data
- Guide the feature selection and subspace investigation process



Data Exploration Module

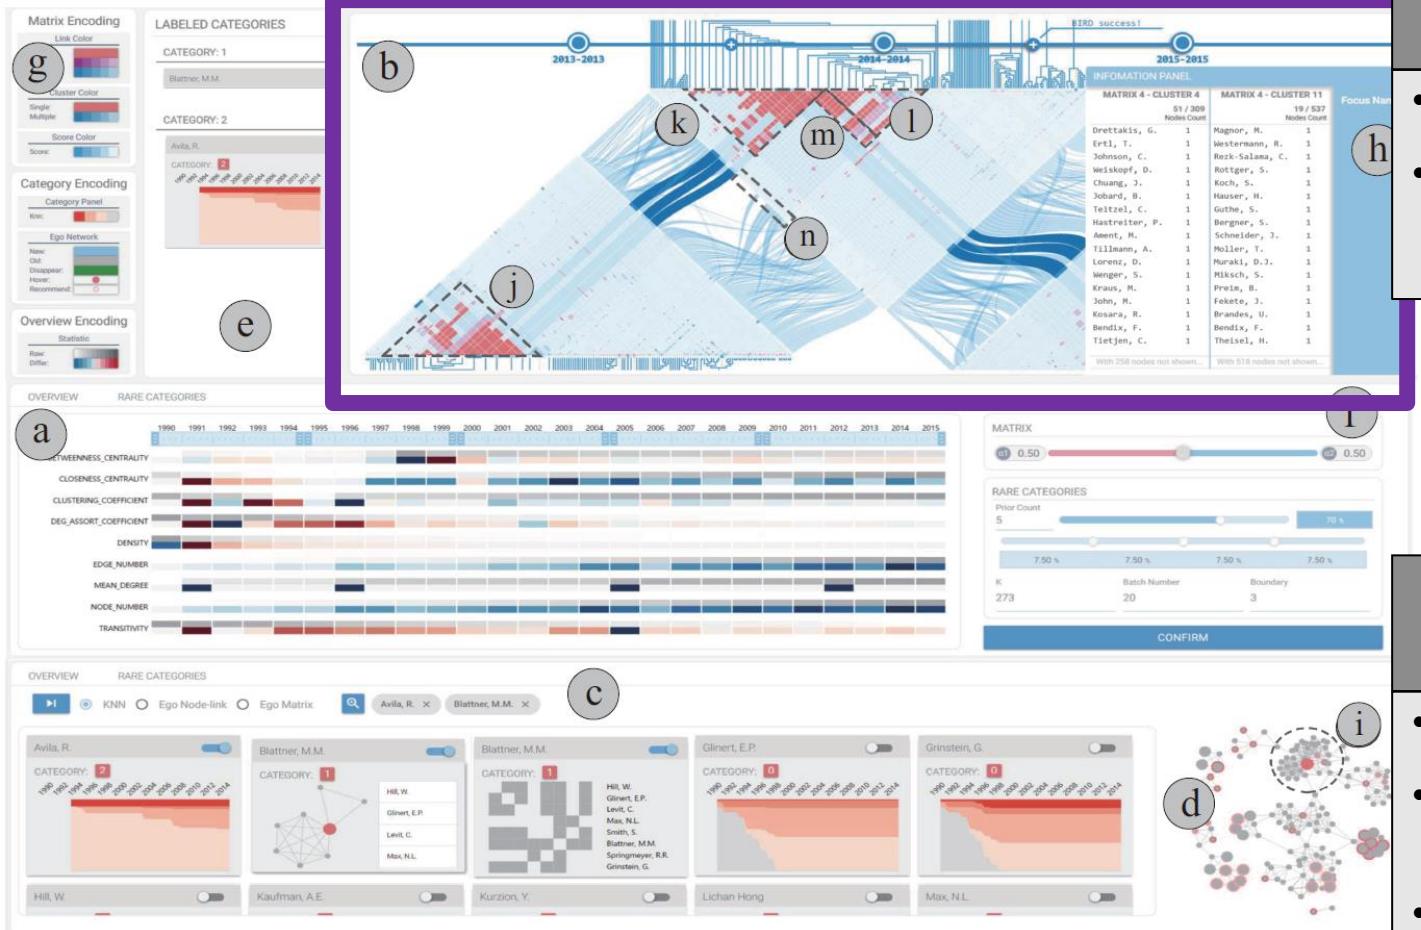
- Represent raw data
- Interactive visualization for data querying
- Support data filtering



RCANALYZER Prototype System

Feature Selection Module

- Visualize the variance of data
- Visualize the correlation of data
- Guide the feature selection and subspace investigation process



Rare Category Analysis Module

- Interactive active learning
- Visualize rare examples in a salient representation

Data Exploration Module

- Represent raw data
- Interactive visualization for data querying
- Support data filtering

Additional References

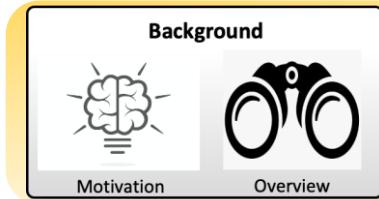
□ Representation-based approaches

- LMLE-kNN [Huang et al., 2016]
- PLANETOID [Yang et al., 2016]
- XGBOD [Zhao et al., 2018]
- s2sL [Dumpala et al., 2018]
- CRT [Chen et al., 2020]
- ...

□ Interpretation-based approaches

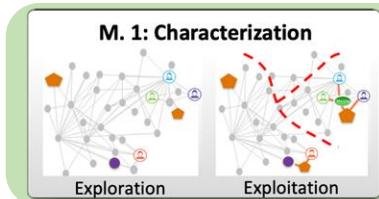
- COIN [Liu et al., 2018]
- SAIFE [Rajendran et al., 2018]
- LookOut [Gupta et al., 2019]
- SIF [Li et al., 2020]
- InFoRM [Kang et al., 2020]
- ...

Roadmap



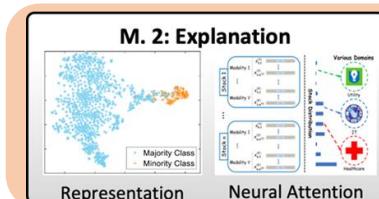
Background

- Motivation
- Research overview



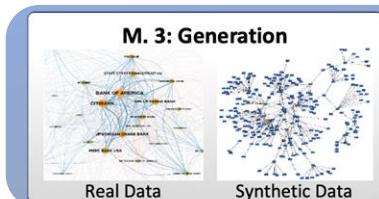
I. Rare Category Characterization

- Rare category characterization on homogeneous graphs
- Rare category characterization on heterogeneous graphs



II. Rare Category Explanation

- Data insights
- Model insights

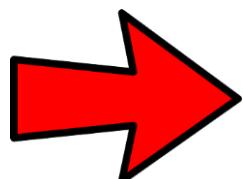


III. Rare Category Generation

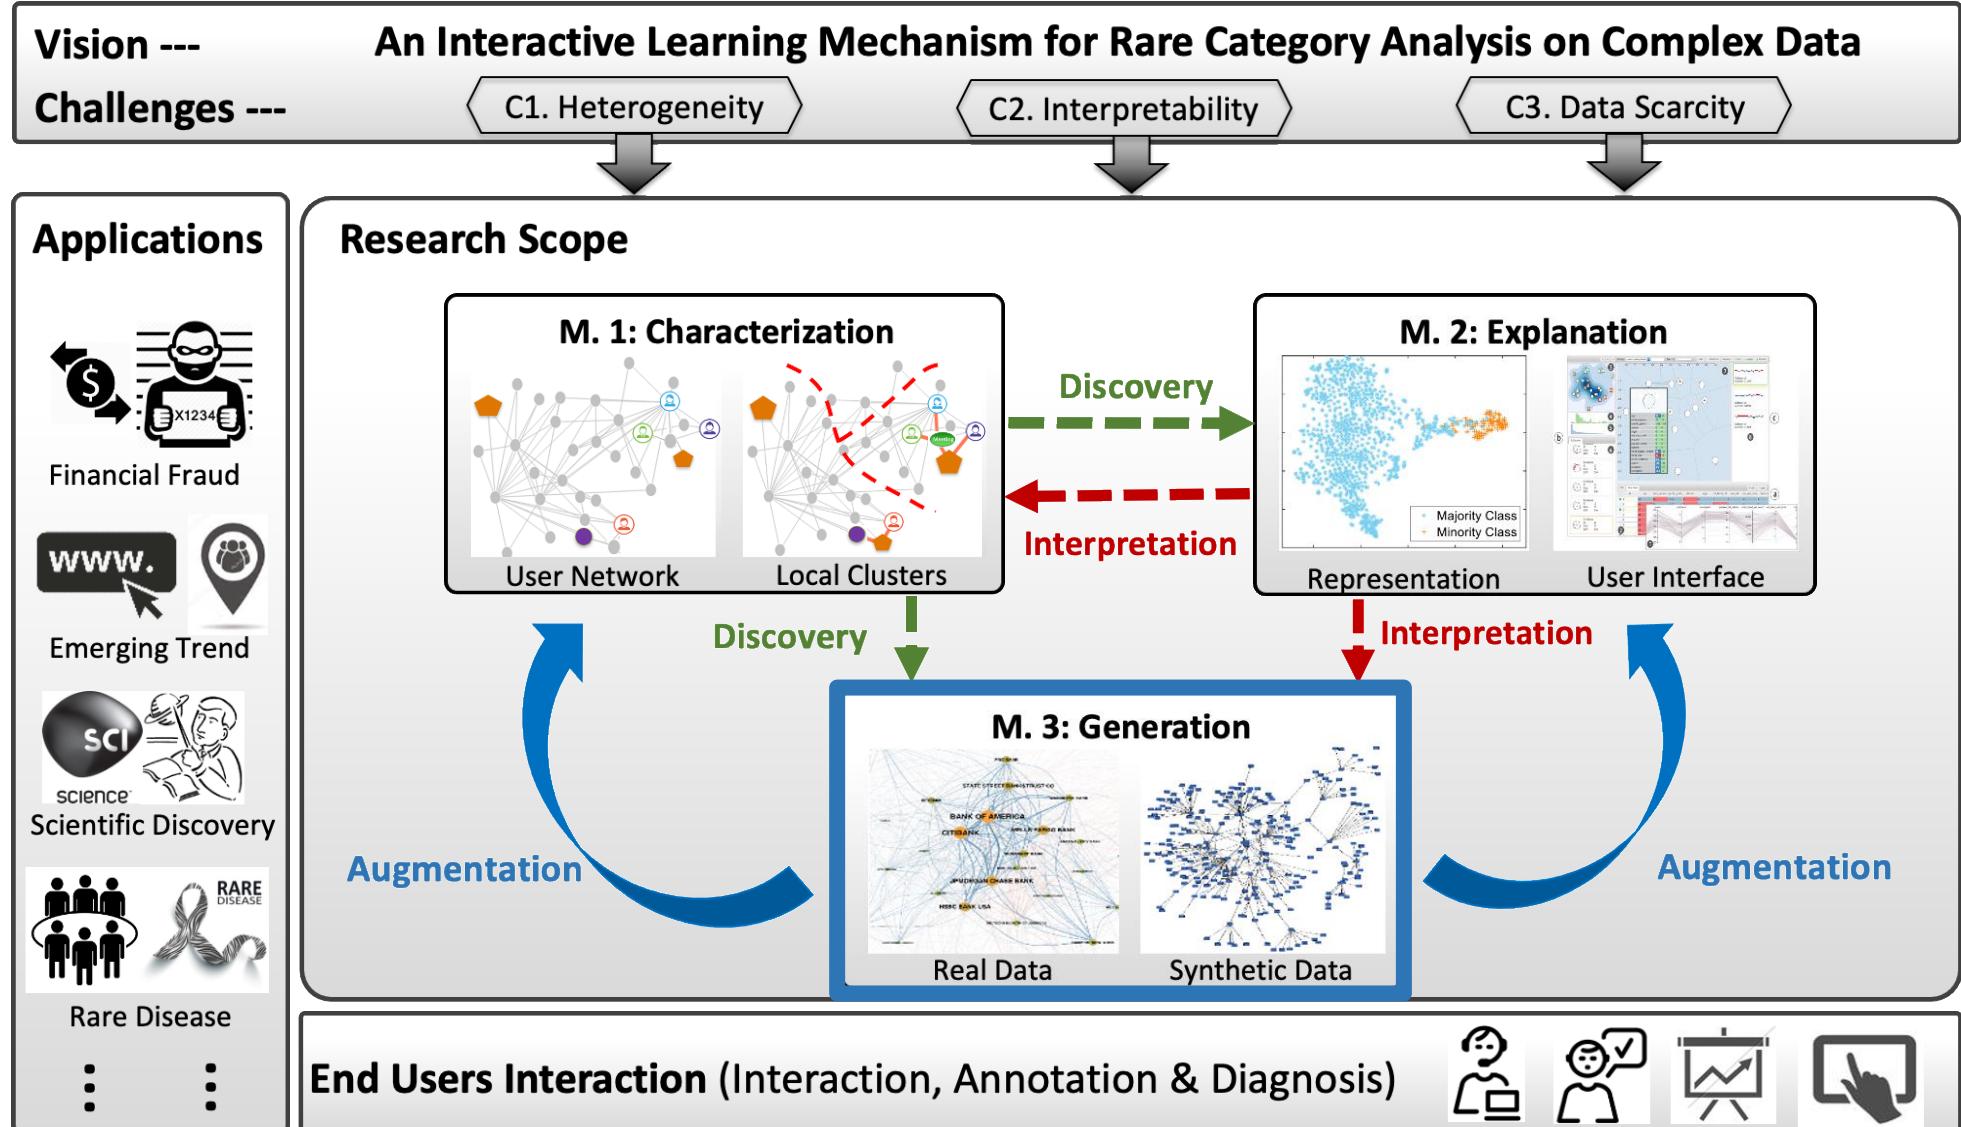
- Unsupervised rare category generation
- Supervised rare category generation



IV. Real-world Application



Rare Category Generation



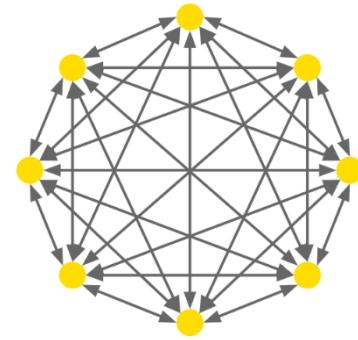
Rare Category Generation

□ Problem 3.1. Rare category generation

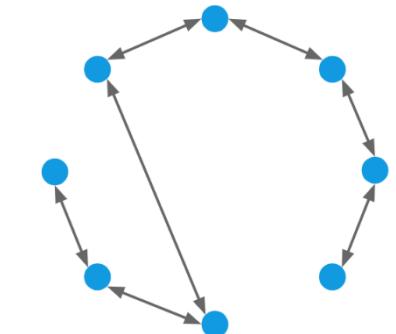
- **Given:** (1) A data set $D = \{x_1, \dots, x_n\}$, (2) a rare category analysis model $f(x)$ trained in D .
- **Find:** (1) A generative model $g(x)$ tailored for D , (2) an augmented data set D' that can boost the performance of $f(x)$ in downstream rare category analysis tasks.

□ Previous work

- RCA models suffer from the data sparsity.
 - Sparse networks.
 - Rare category (**local cluster**) is not prominent.
- Synthetic minority oversampling technique (SMOTE).
 - Synthesizing new examples for the **minority class**.
 - Not for graph-structured data.



A subgraph



A subgraph
with missing links

Injecting rich connections

• Chawla, Nitesh V., et al. "SMOTE: synthetic minority over-sampling technique." Journal of artificial intelligence research. 2002.



Rare Category Generation

□ Research questions

- **Q1. Unsupervised rare category generation**
 - Can we augment the **sparse** networks by injecting synthetic connections?
- **Q2. Supervised rare category generation**
 - Can we generate **task-specific** data sets while preserving key properties?

□ Observations

- **Unsupervised rare category generation**
 - Sparsity issue exists in graph data.
 - **Augmenting the original network** boosts the performance of rare category detection (Zhou et al., 2020).
- **Supervised rare category generation**
 - Label-informed Molecular graph generative model **preserving key properties** (Jin et al., 2018);
 - Generating molecular graphs maintaining **chemical validity** and **high potency** (Jin et al., 2018);
 - Generating molecular graphs from a specific class (Simonovsky et al., 2018);
 - Considering the desired chemical properties as the label for graph generation.

• Simonovsky, Martin et al. "Graphvae: Towards generation of small graphs using variational autoencoders." ICANN 2018.

- 77 - • Jin, Wengong, et al. "Junction tree variational autoencoder for molecular graph generation." ICML 2018.

• Zhou, Dawei, et al. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." ACM SIGKDD. 2020.

Generative Models

□ Variational autoencoder based methods

- GraphVAE (Simonovsky et al., 2018), JTVAE (Jin et al., 2018), Graphite-VAE (Grover et al., 2019), Reg-VAE (Ma et al., 2018), NEVAE (Samanta et al., 2019),

□ Generative adversarial network based methods

- NetGAN (Bojchevski et al., 2018), MolGAN (De Cao et al., 2018), GCPN (You et al., 2018), Misc-GAN (Zhou et al., 2019), LGGAN (Fang et al., 2019),

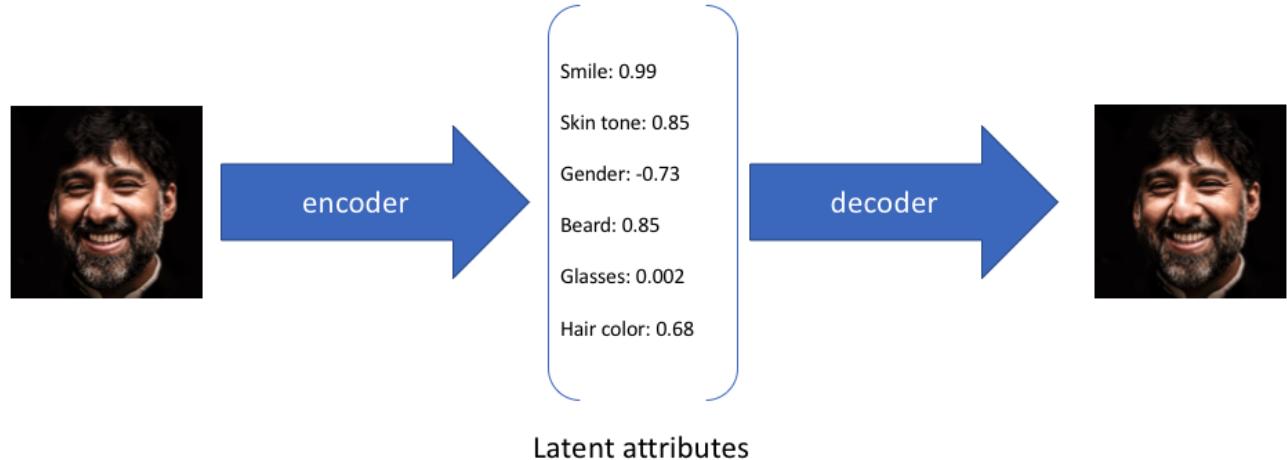
□ Autoregressive based methods

- GraphRNN (You et al., 2018), TagGen (Zhou et al., 2020), MolecularRNN (Popova et al., 2019),

Variational Autoencoder

□ Autoencoder

- An encoder model $q_\phi(z|x)$,
- A decoder model $p_\theta(x|z)$.



□ Maximizing the variational lower-bound

- Reparameterization trick to approximate intractable posterior

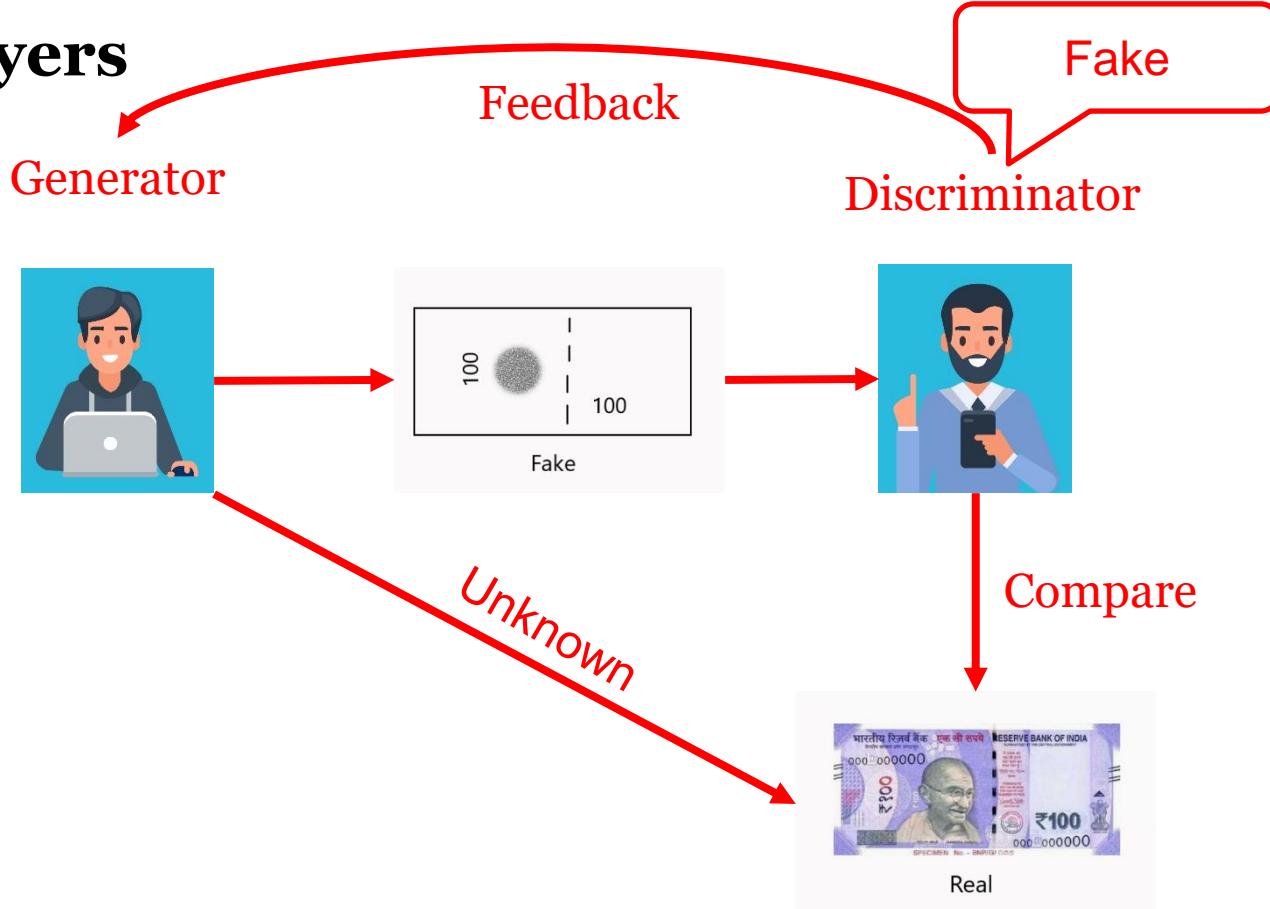
$$\mathcal{L}(\theta, \phi, x) = \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) || p_\theta(z))$$

Reconstruction Capability Regularization

Generative Adversarial Nets (GANs)

□ GAN involves two competing players

- Generator G
- Discriminator D



- Goodfellow, Ian, et al. "Generative adversarial nets." NIPS 2014.

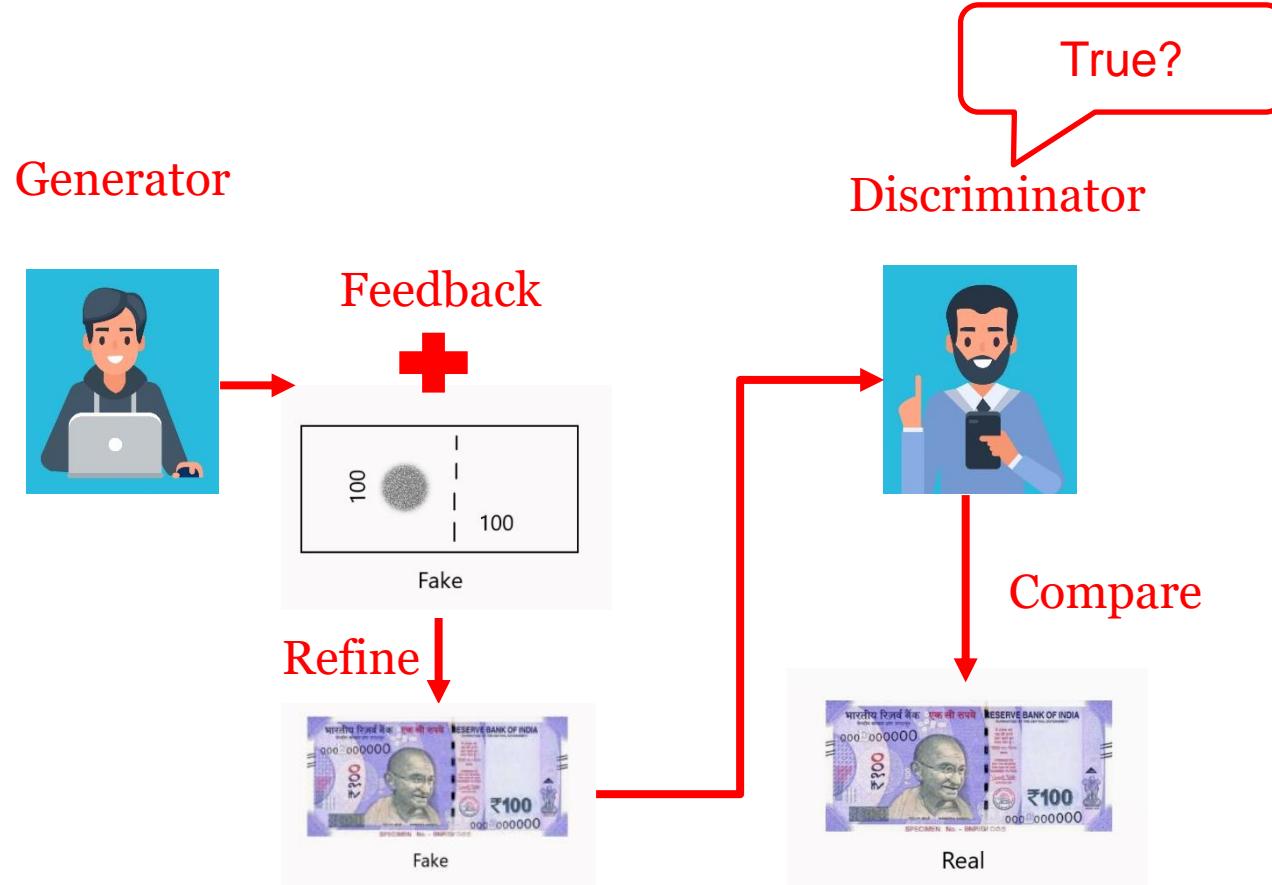
Generative Adversarial Nets (GANs)

□ Generator G

- Takes random variable as an input
- Generates synthetic samples
- Aims fool the discriminator D .

□ Discriminator D

- Aims to distinguish between **real data** and **generated data**.



• Goodfellow, Ian, et al. "Generative adversarial nets." NIPS 2014.

Generative Adversarial Nets (GANs)

□ Generator G

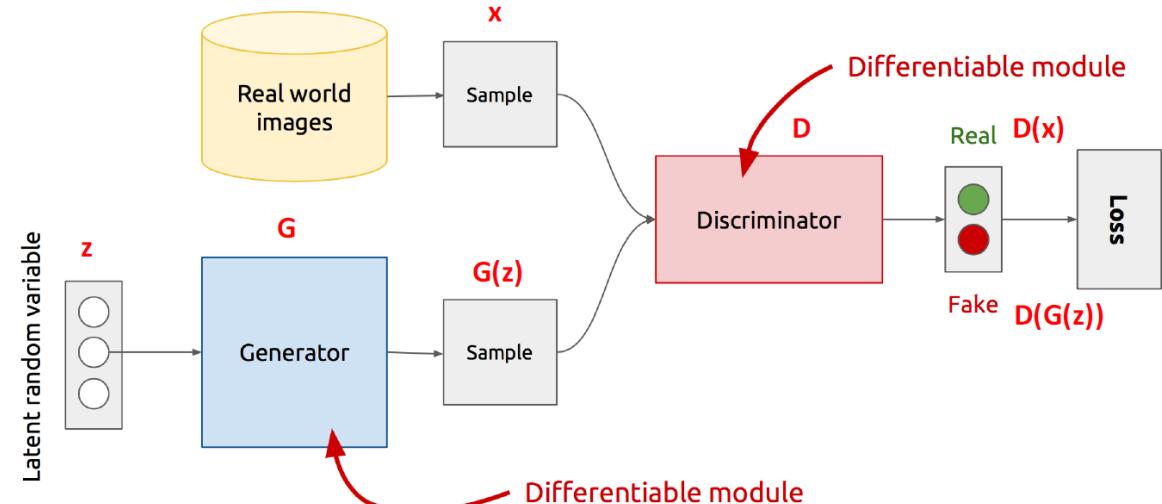
- Takes random variable as an input
- Generates synthetic samples
- Aims fool the discriminator D .

□ Discriminator D

- Aims to distinguish between **real** data and **generated data**.

□ The **min-max game** is formulated as

$$\min_G \max_D V(G, D) = \mathbb{E}_{x \sim p_{data}(x)} \log D(x) + \mathbb{E}_{z \sim p_z(z)} \log(1 - D(G(z)))$$



• Goodfellow, Ian, et al. "Generative adversarial nets." NIPS 2014.

- 82 - • Image source: <https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

Deep Autoregressive Models

□ Key idea:

- Predicts future behavior based on past behavior

□ Example

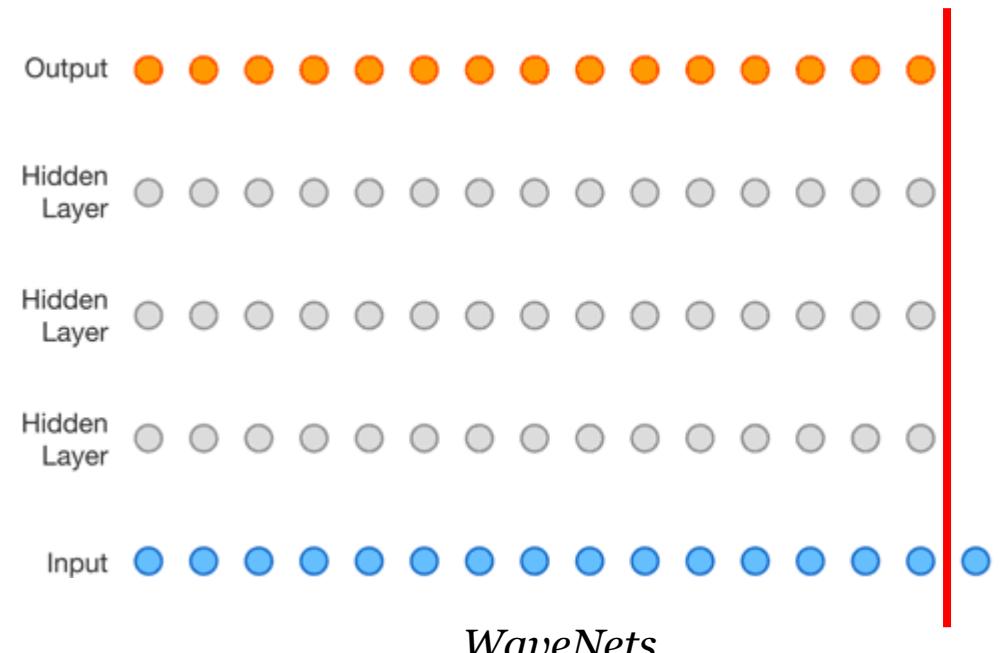
- Recurrent Neural Network (RNN)

□ WaveNets (Oord et al. 2016)

- A generative model of audio waveforms
- $p(x) = \sum_{t=1}^T p(x_t|x_1, \dots, x_{t-1})$

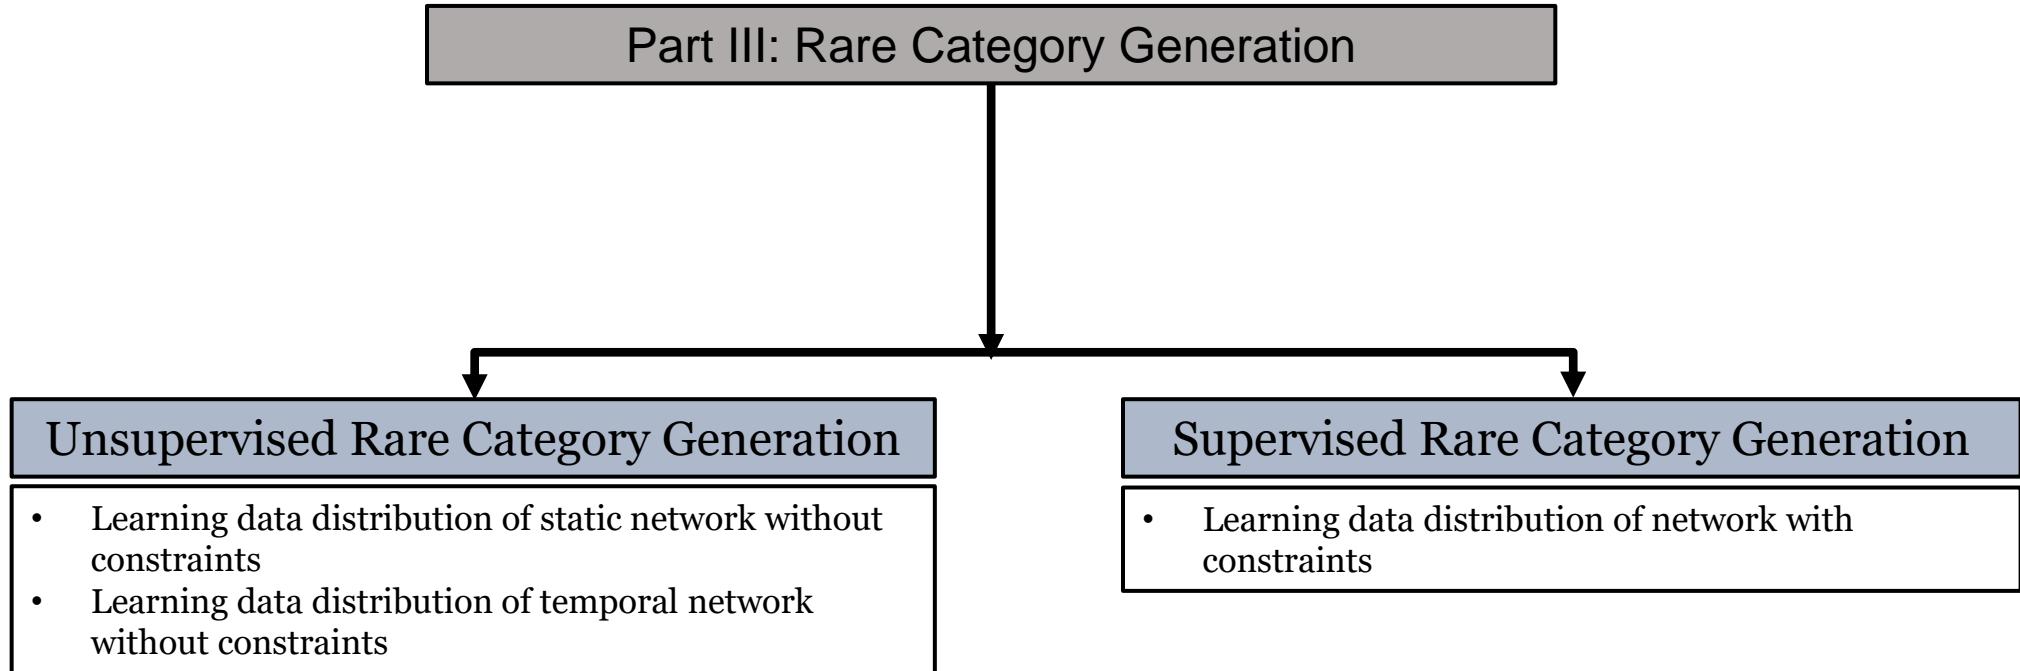
□ XLNet (Yang, et al. 2019)

- A generalized autoregressive language model
- $\max_{\theta} \sum_{t=1}^T \log p_{\theta}(x_t|x_1, \dots, x_{t-1})$



- Oord, Aaron van den, et al. "Wavenet: A generative model for raw audio." arXiv preprint arXiv:1609.03499 (2016).
- Yang, Zhilin, et al. "Xlnet: Generalized autoregressive pretraining for language understanding." NeurIPS 2019.
- Image source: <https://deepmind.com/blog/article/wavenet-generative-model-raw-audio>

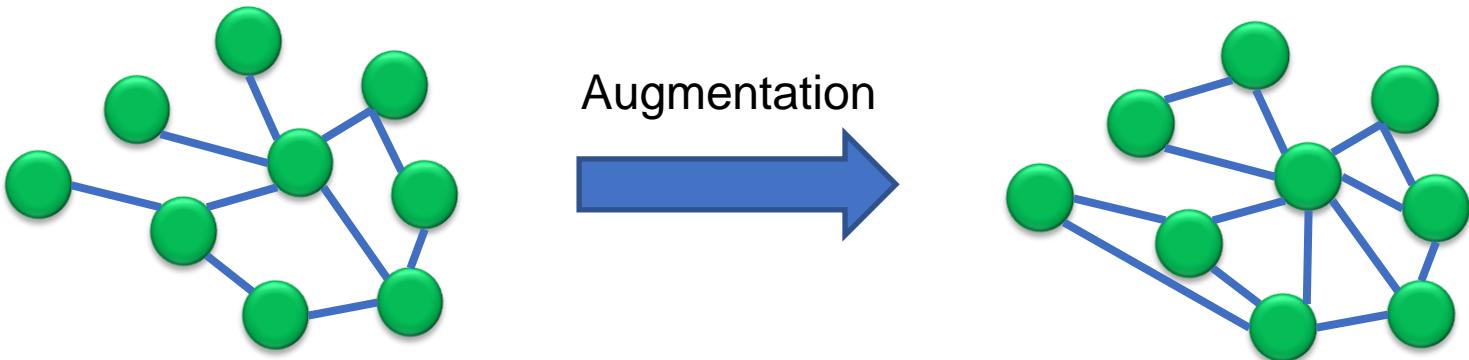
Rare Category Generation



Unsupervised Rare Category Generation

□ Problem 3.2. Static network generation

- **Given:**
 - (1) A static network G ,
 - (2) A rare category analysis model $f(G)$ trained in G .
- **Find:**
 - (1) A generative model $g(G)$ tailored for G ,
 - (2) An augmented graph G' that can boost the performance of $f(G')$ in downstream rare category analysis tasks.



GraphRNN

□ Autoregressive Generative Model

- Not learns the distribution of graph $p(G)$ directly;
- Aims to learn the distribution of graph sequence $p(S^\pi)$;

□ Deep Autoregressive model

- One process that generates **a sequence of nodes** (graph-level RNN);
- One process that generates **a sequence of edges** for each newly added node (edge-level RNN).
- $p(S_i^\pi | S_{<i}^\pi)$ captures how current node is connected to previous nodes based on how the previous nodes are connected to each other.

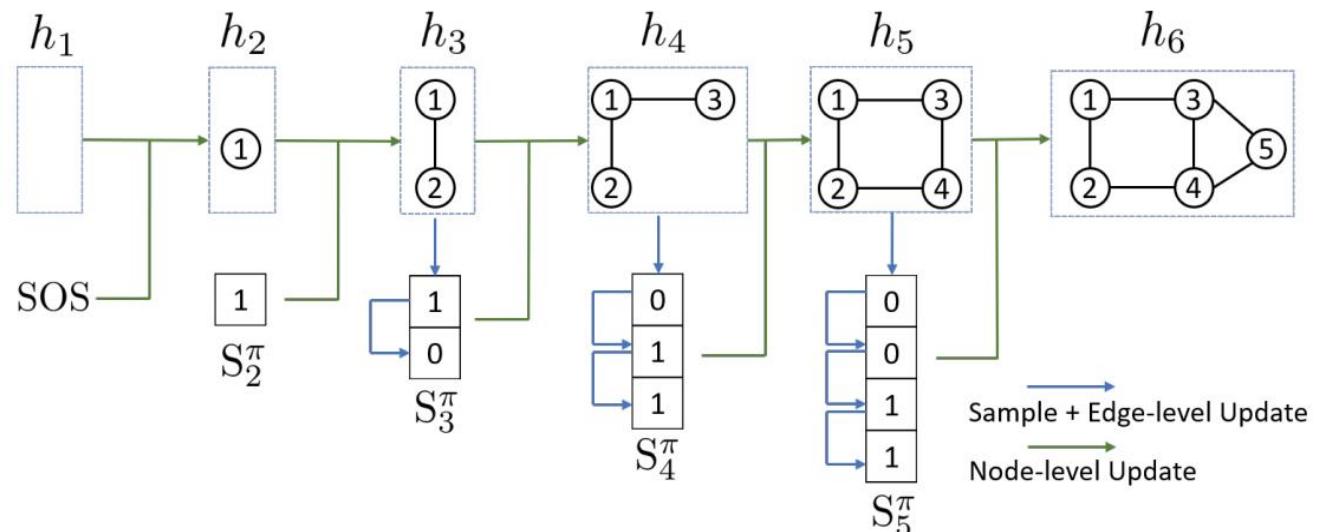
$$p(S^\pi) = \prod_{i=1}^{n+1} p(S_i^\pi | S_1^\pi, \dots, S_{i-1}^\pi) = \prod_{i=1}^{n+1} p(S_i^\pi | S_{<i}^\pi)$$

- You, Jiaxuan, et al. "Graphrnn: Generating realistic graphs with deep auto-regressive models." ICML 2018.

GraphRNN

□ A Hierarchical RNN

- An RNN that consists of a state-transition function f_{trans} and an output function f_{out} to parameterize $p(S_i^\pi | S_{<i}^\pi)$.
- Graph-level: $h_i = f_{trans}(h_{i-1}, S_{i-1}^\pi)$ encodes the state of the graph generated so far.
- Edge-level: $f_{out}(h_i)$ generates the edges of a given node.
- Graph is assembled from $S^\pi = \{S_1^\pi, S_2^\pi, \dots, S_N^\pi\}$.



- You, Jiaxuan, et al. "Graphrnn: Generating realistic graphs with deep auto-regressive models." ICML 2018.

NetGAN

□ Key idea:

- Aims to learn the distribution of graph sequence

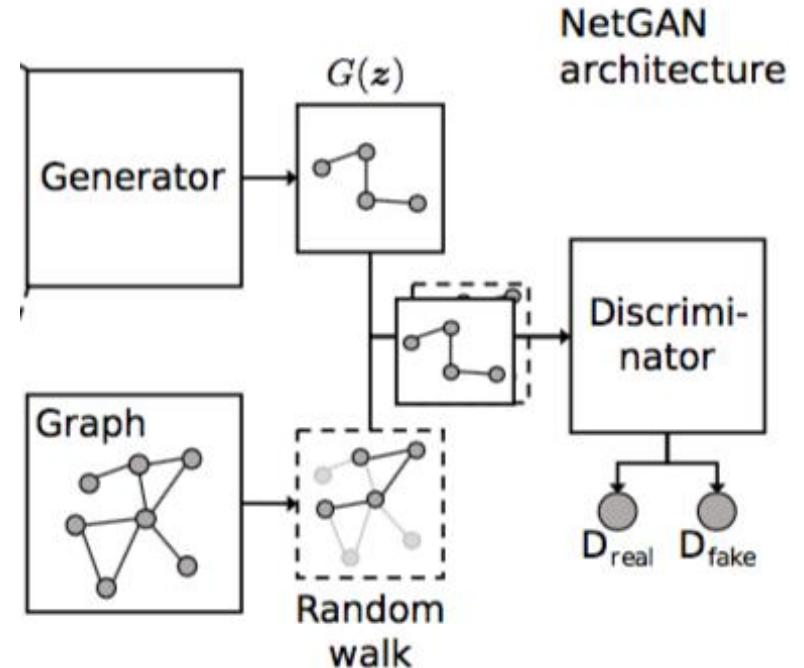
□ Step 1: Learning how to generate realistic random walks

○ Generator G

- Generate fake random walks;
- Learns to generate random walks (v_1, \dots, v_T) .

○ Discriminator D

- Takes random walks as inputs;
- Distinguishes whether random walks are real or fake.

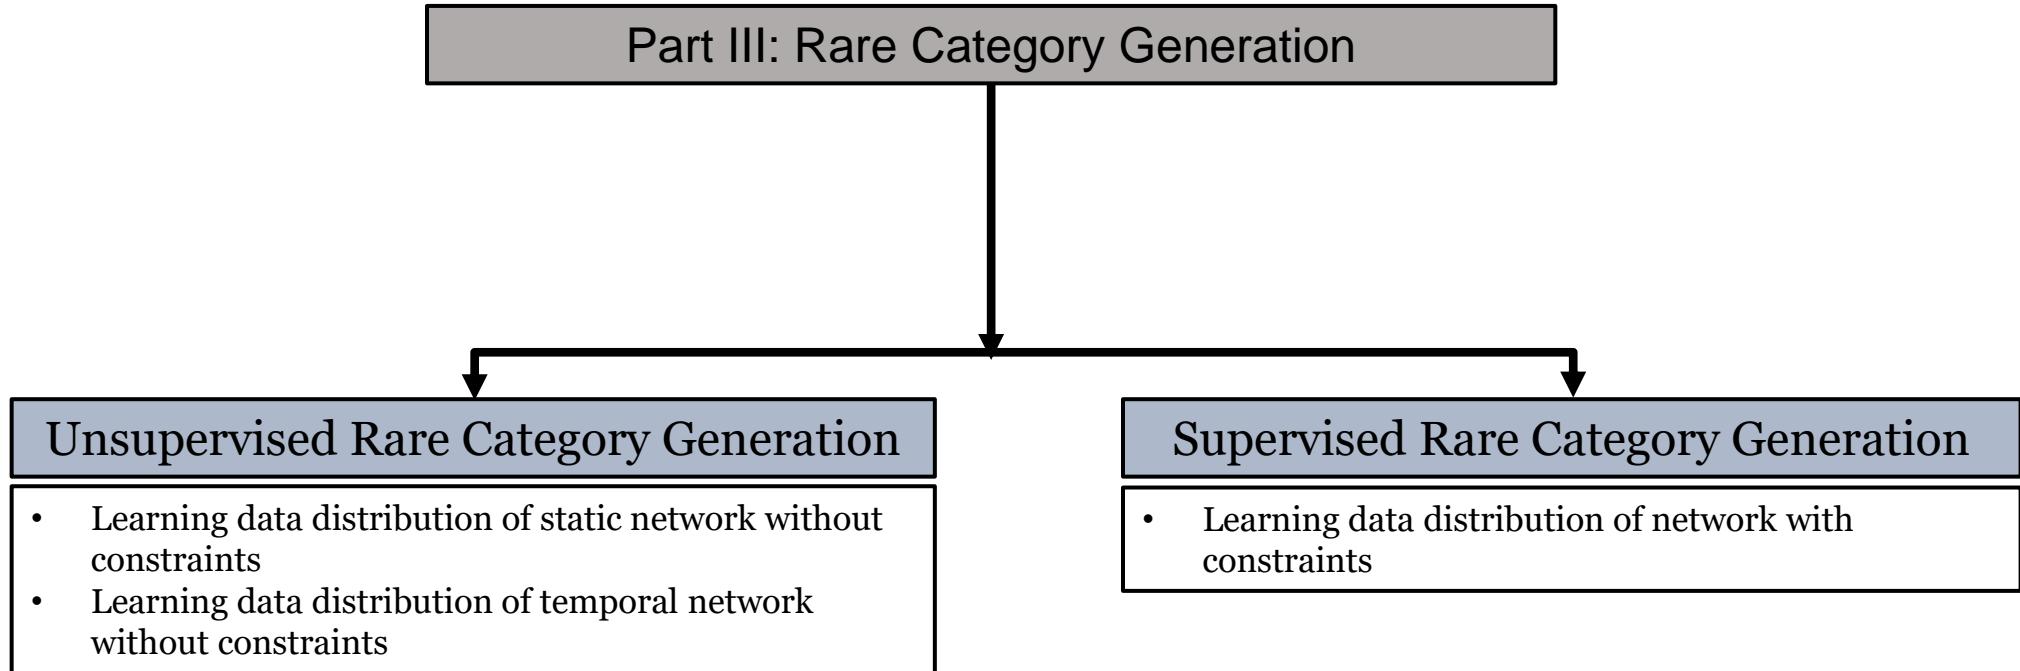


- Bojchevski, Aleksandar, et al. "Netgan: Generating graphs via random walks." ICML 2018.

□ Step 2: Assembling the adjacency matrix

- Constructs a score matrix S with edge counts;
- Converts the matrix S into binary adjacency matrix A in two steps.
 - (1). Ensure that every node i has at least one edge by sampling a neighbor j with probability $p_{ij} = \frac{s_{ij}}{\sum_v s_{iv}}$.
 - (2). Continue sampling edges without replacement using the probability $p_{ij} = \frac{s_{ij}}{\sum_{u,v} s_{uv}}$ for each edge (i,j) , until we reach the desired number of edges.

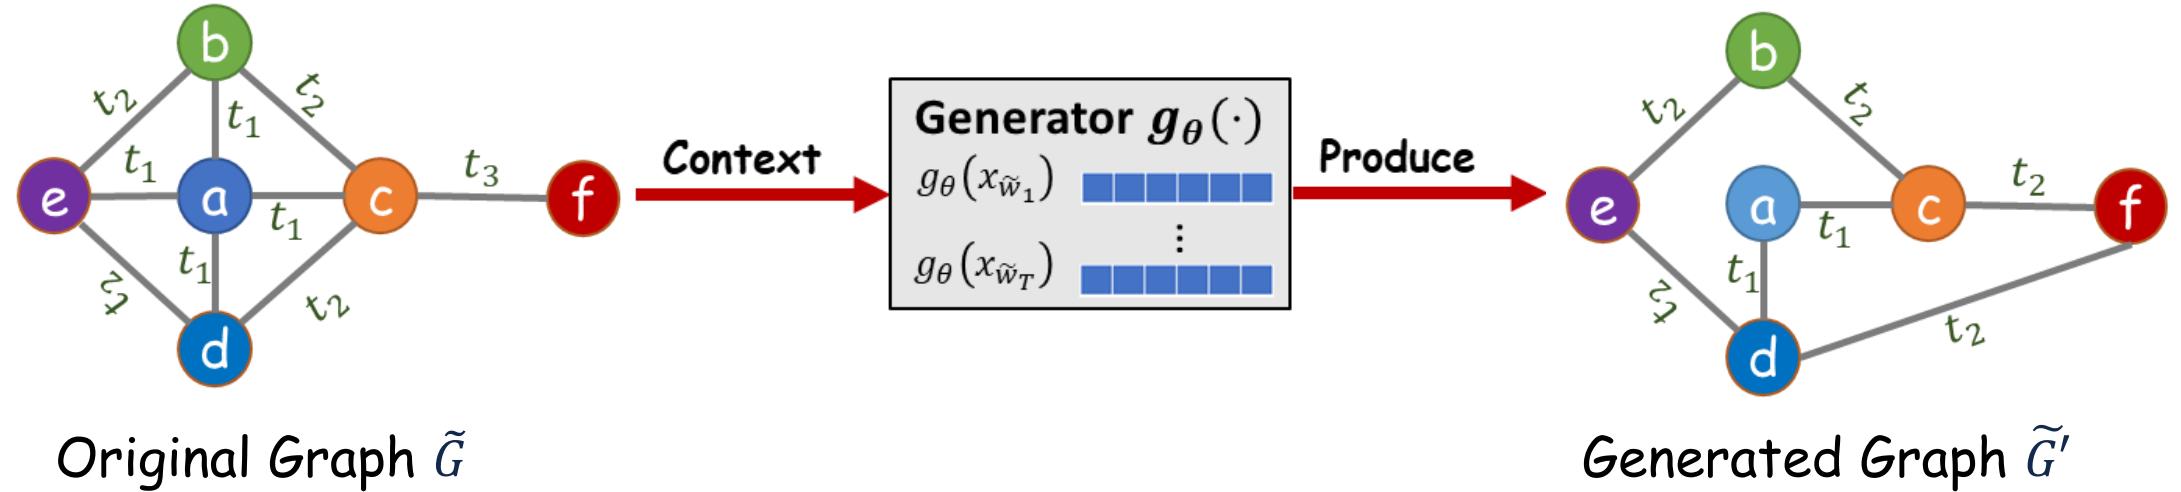
Rare Category Generation



Unsupervised Rare Category Generation

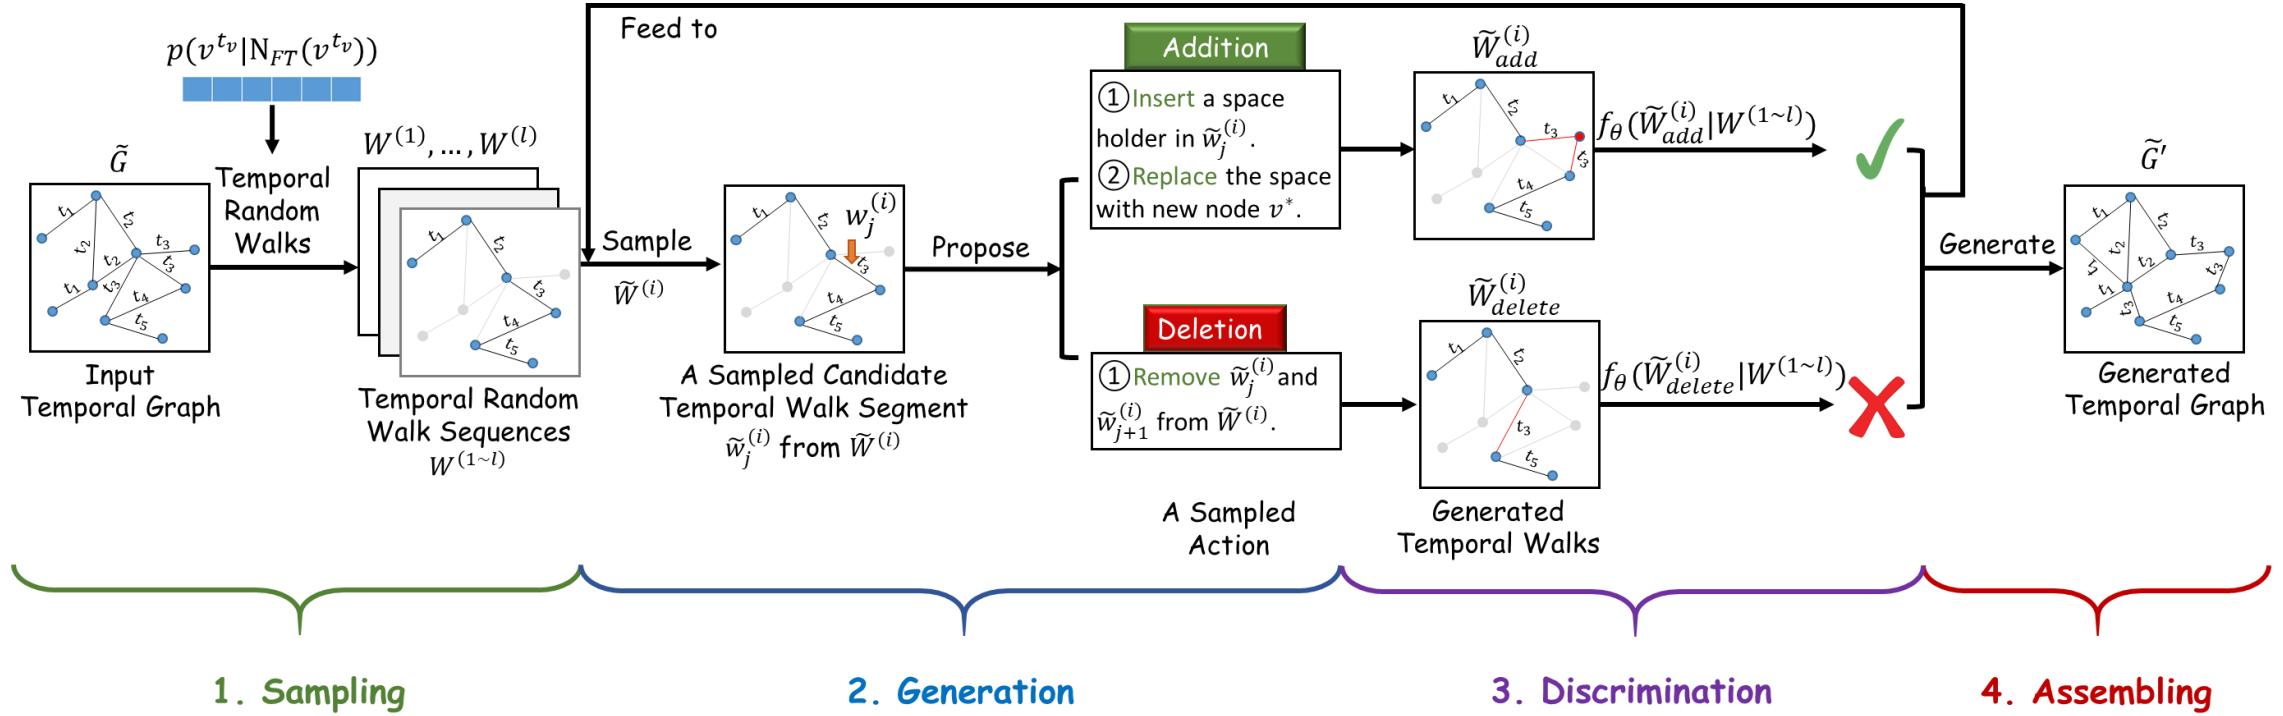
□ Problem 3.3. Temporal network generation

- **Given:**
 - (1) A temporal network \tilde{G} , which is represented as a collection of timestamped edges $\tilde{E} = (e_1^{t_{e_1}}, e_2^{t_{e_2}}, \dots, e_m^{t_{e_m}})$.
 - (2) A rare category analysis model $f(\tilde{G})$ trained in \tilde{G} .
- **Find:**
 - A synthetic temporal network \tilde{G}' that can boost the performance of $f(\tilde{G}')$ in downstream rare category analysis tasks.



- Zhou, Dawei, et al. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." ACM SIGKDD. 2020.

TAGGen – Temporal Graph Generation



1. A novel context extraction strategy for temporal networks.

2. A family of local operations to perform addition and deletion of nodes and edges.

3. A bi-level self-attention mechanism.

- Zhou, Dawei, et al. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." ACM SIGKDD. 2020.

TAGGen – Algorithm

□ S1: Context sampling

- **Goal:** selecting initial nodes for conducting temporal random walks.
- **Assumption:** weak dependence.

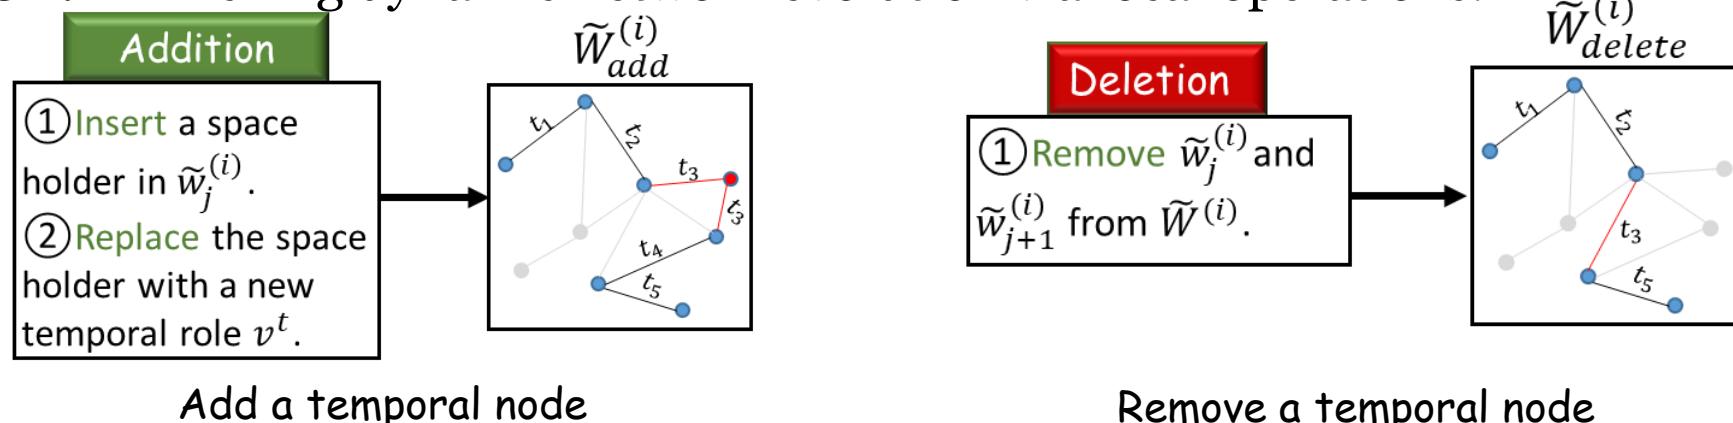
Temporal neighborhood of v^{t_v}

$$p(v^{t_v} | N_{FT}(v^{t_v})) > \delta [p(v^{t_v} | N_S(v^{t_v}))p(v^{t_v} | N_T(v^{t_v}))]$$

Weak dependence

□ S2: Sequence generation

- **Goal:** generating synthetic temporal random walks.
- **Solution:** mimicking dynamic network evolution via local operations.



• Zhou, Dawei, et al. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." ACM SIGKDD. 2020.

TAGGen – Algorithm

□ S3: Sequence discrimination

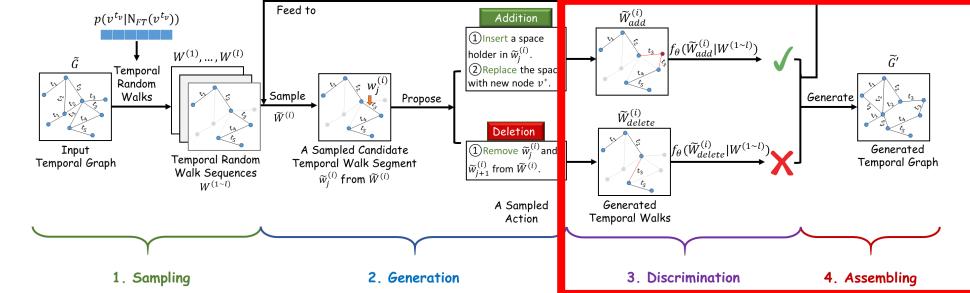
- **Goal:** selecting synthetic random walks that are plausible in the input graph.
- **Solution:** a bi-level self-attention mechanism.
 - maximizing the action likelihood via the deep autoregressive model $f_\theta(\cdot)$.

$$p(\tilde{W}^{(i)}_{action} | W^{(1 \sim l)})$$

- \tilde{W}_{action} : generated random walk sequence after a sampled action.

□ S4: Graph Assembling

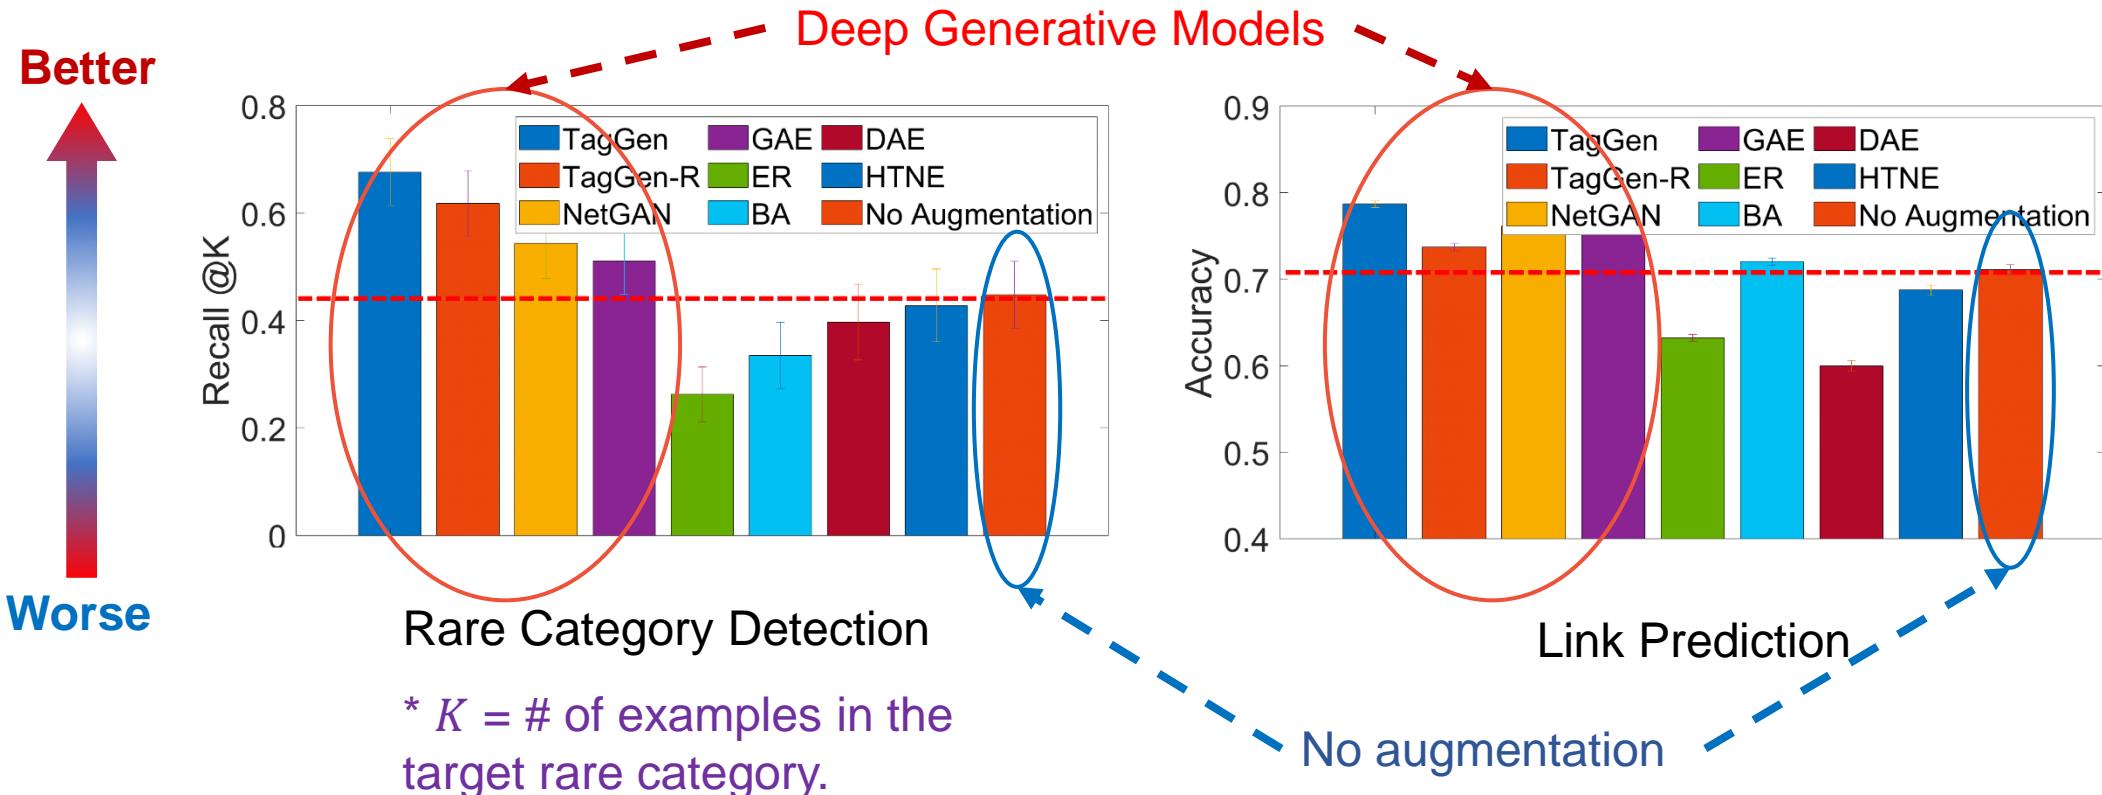
- **Goal:** assemble all the generated temporal random walks and generate the temporal networks.
- **Solution:** assembling rules to avoid some rare temporal occurrences (i.e., with a small degree) are not sampled.
 - Sample at least one temporal edge starting from each node with probability $p(v^{t_v})$.
 - Sample at least one temporal edge at each timestamp with probability $p(e^{t_e})$.
 - Stop until the generated graph has the same edge density as the input graphs.



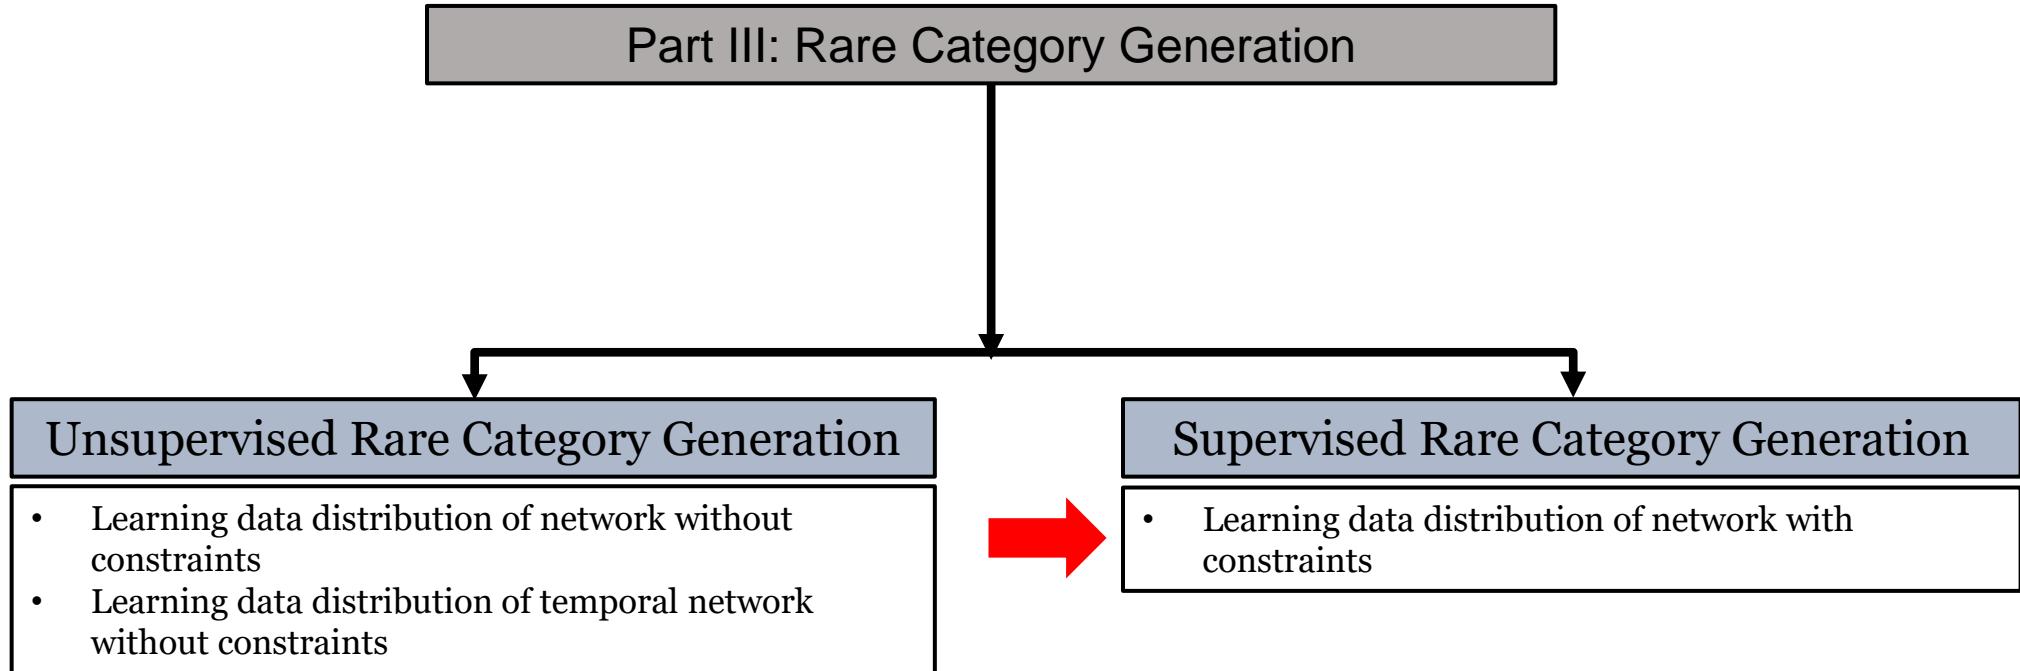
Unsupervised Graph Generation Experiment

□ Data augmentation in downstream tasks.

- All augmented graphs/original graph are evaluated by logistic regression.
- 10% more edges are injected to augment the original graph.



Rare Category Generation





Supervised Rare Category Generation

□ Problem 3.4. Molecular graph generation

- **Given:**
 - (1) A set of molecular graphs $D = \{G_1, \dots, G_N\}$,
 - (2) Labels or constraints for graph generation,
 - (3) A rare category analysis model $f(x)$ trained in D .
- **Find:**
 - (1) A generative model $g(G)$ tailored for D ,
 - (2) A set of generated graphs D' that can boost the performance of $f(x)$ in downstream rare category analysis tasks.

GraphVAE

□ **Key idea:** generate a probabilistic fully-connected graph and use a graph matching algorithm to align it to the ground truth.

□ **Objective function:**

- **Reconstruction loss** enforces high similarity of generated graphs to the input graph G .
- **KL-divergence** regularizes the space to allow for sampling of z directly from $p(z)$ instead from $q_\phi(z|G)$.

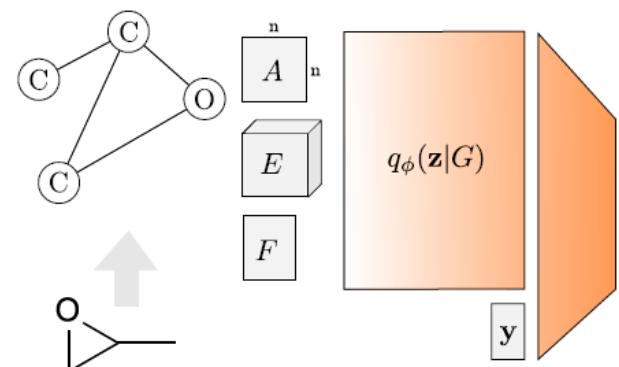
$$\mathcal{L}(\theta, \phi, G) = \mathbb{E}_{q_\phi(z|G)} [-\log p_\theta(G|z)] + D_{KL}(q_\phi(z|G)||p_\theta(z))$$

Reconstruction loss

Regularization

□ **Encoder:**

- Adjacency matrix A ,
- Edge attribute tensor E ,
- Node attribute matrix F ,
- Embeds the graph into continuous representation z .

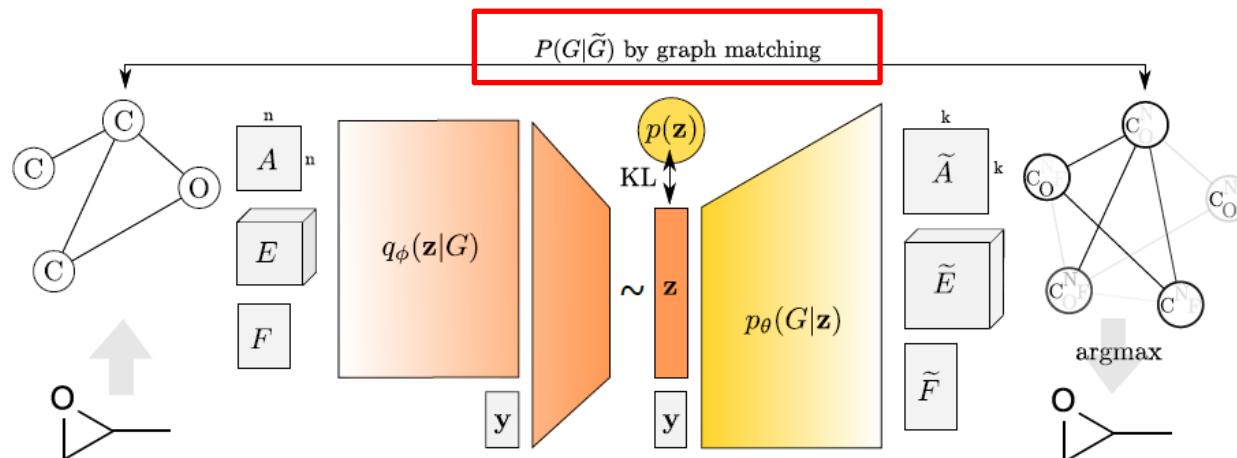


• Simonovsky, Martin et al. "Graphvae: Towards generation of small graphs using variational autoencoders." ICANN 2018.

GraphVAE

□ Decoder:

- Outputs a probabilistic fully-connected graph $\tilde{G} = (\tilde{A}, \tilde{E}, \tilde{F})$.
- $\tilde{A} \in [0,1]^{k \times k}$ contains both node probabilities $\tilde{A}_{a,a}$ and edge probabilities $\tilde{A}_{a,b}$ ($a \neq b$).
- $\tilde{E} \in \mathbb{R}^{k \times k \times d_e}$ indicates class probabilities for edges.
- $\tilde{F} \in \mathbb{R}^{k \times d_n}$ contains class probabilities for nodes.
- Max-pooling graph matching aligns the reconstructed graph to the ground truth.
- At test time, the process can be **conditioned on label y** for controlled sampling.



- Simonovsky, Martin et al. "Graphvae: Towards generation of small graphs using variational autoencoders." ICANN 2018.

□ **Key idea:** consider each subgraph as a node, generate a tree-structured scaffold over chemical substructures, and then combine them into a molecule.

□ **Graph variational autoencoder**

- **Aims to capture the fine-grained connectivity**

- **Graph encoder**

- Encodes graph structure with **message passing**.

$$\nu_{uv}^{(t)} = \tau(\mathbf{W}_1^g \mathbf{x}_u + \mathbf{W}_2^g \mathbf{x}_{uv} + \mathbf{W}_3^g \sum_{w \in N(u) \setminus v} \nu_{wu}^{(t-1)})$$

Messages Node Edge
feature feature feature

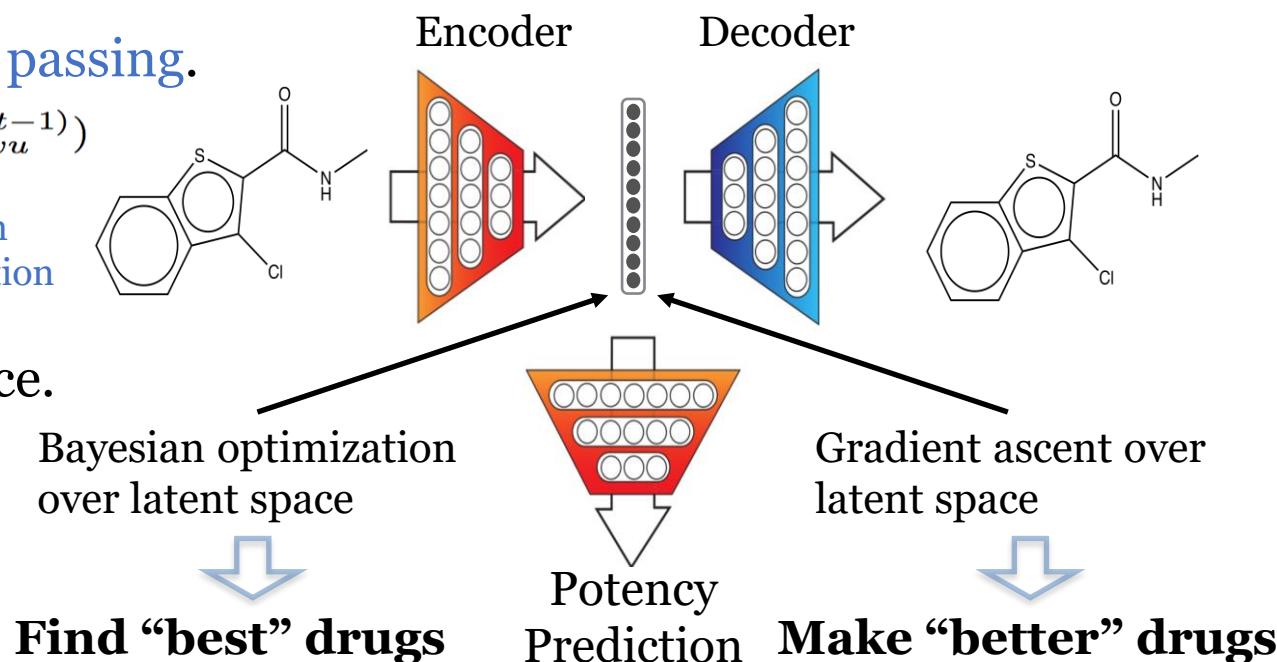
Messages from previous iteration

- **Graph decoder**

- Reconstructs molecule from latent space.

- **Classifier**

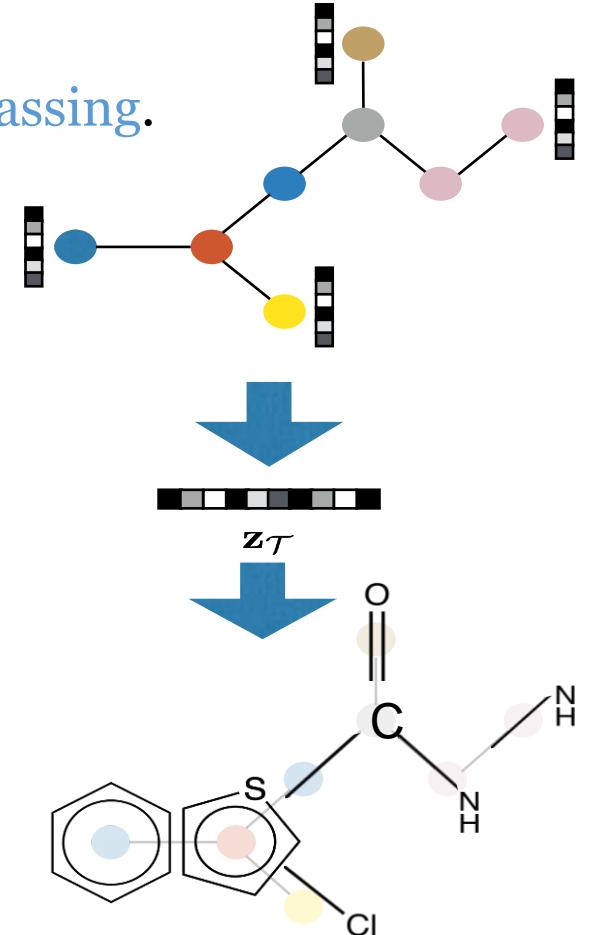
- Predicts the potency of a molecule.



• Jin, Wengong, et al. "Junction tree variational autoencoder for molecular graph generation." ICML 2018.

□ Junction tree graph generation

- **Aims to capture the tree structure and clusters**
- **Tree encoder**
 - Vectorizes the junction tree into latent space z_T with **message passing**.
 - $m_{ij} = GRU(x_i, \{m_{ki}\}_{k \in N(i) \setminus j})$
 - Encode longer range interactions between nodes.
- **Tree decoder**
 - Reconstruct the junction tree from latent space z_T .



□ Assemble two substructures into a coherent molecular graph

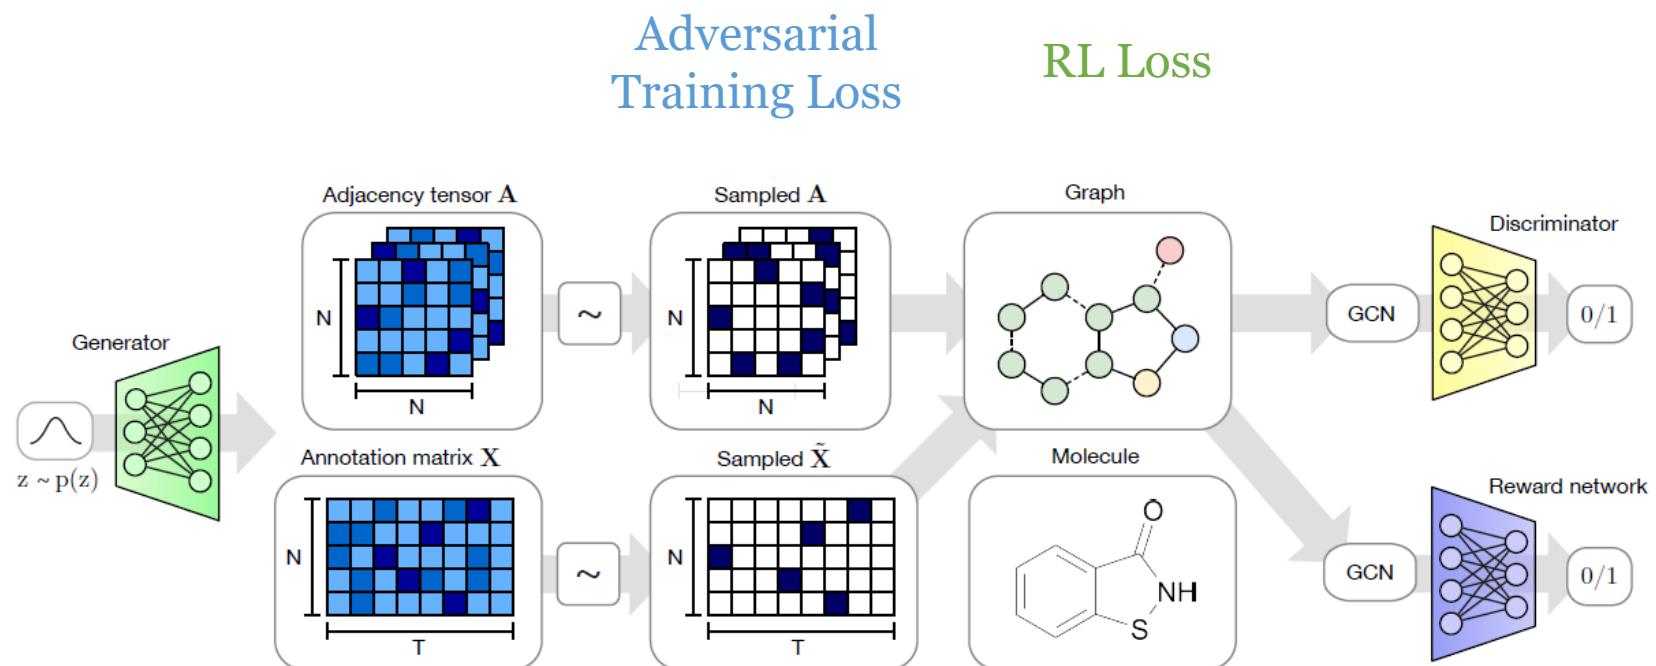
- Jin, Wengong, et al. "Junction tree variational autoencoder for molecular graph generation." ICML 2018.

□ A likelihood-free molecular graph generative model

- Deterministic policy gradient is utilized to achieve useful property (e.g., to be easily synthesizable).

□ Objective function:

$$L(\theta) = \lambda L_{WGAN} + (1 - \lambda) L_{RL}$$



- De Cao, Nicola, et al. "MolGAN: An implicit generative model for small molecular graphs." ICML 2018.

□ Generator:

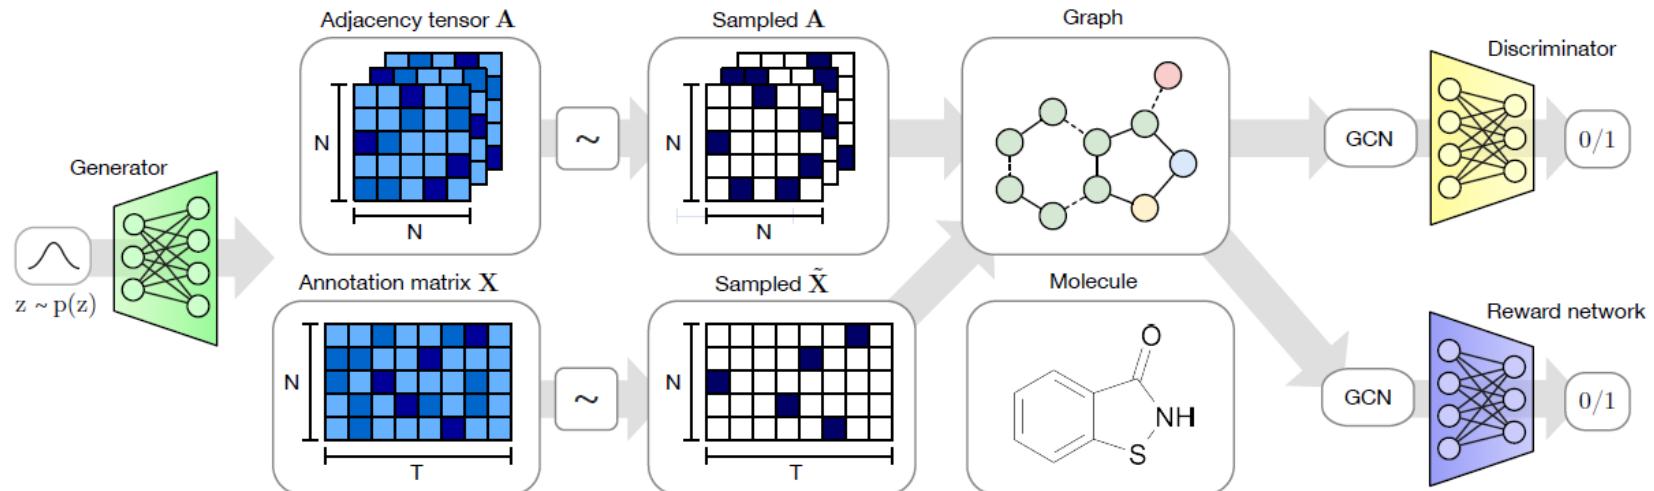
- Generates atom types X and bond types A from a prior distribution.

□ Discriminator:

- Distinguishes if the graph is real or synthesized.

□ Reward network

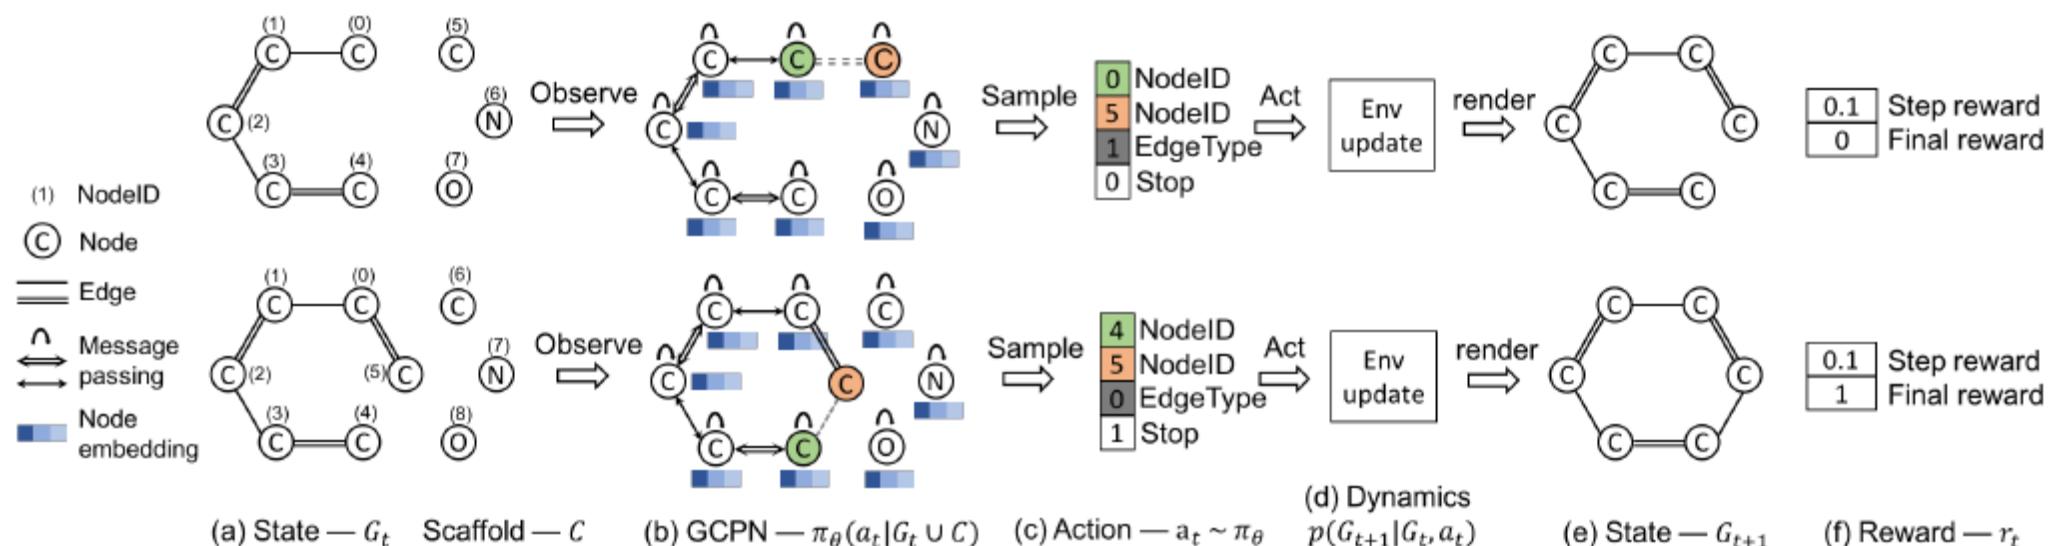
- Assigns a reward to the graph provided by external software;
- **Assigns zero to invalid molecular graph.**



- De Cao, Nicola, et al. "MolGAN: An implicit generative model for small molecular graphs." ICML 2018.

Graph Convolutional Policy Network (GCPN)

- **Key idea:** learn a reinforcement learning agent that iteratively adds substructures and edges to the molecular graph in a chemistry-aware environment.
 - **Goal:** generate graphs that maximize a given property function $S(G') \in R$.
- $$\max \mathbb{E}_{G'}[S(G')]$$



• You, Jiaxuan, et al. "Graph convolutional policy network for goal-directed molecular graph generation." Advances in neural information processing systems 2018.

Supervised Graph Generation Experiment

□ **Dataset:** ZINC250k molecule dataset.

□ **Constrained property optimization:** focus on generating molecules with the highest possible penalized logP and QED scores.

Method	Penalized logP				QED			
	1st	2nd	3rd	Validity	1st	2nd	3rd	Validity
ZINC	4.52	4.30	4.23	100.0%	0.948	0.948	0.948	100.0%
Hill Climbing	—	—	—	—	0.838	0.814	0.814	100.0%
ORGAN	3.63	3.49	3.44	0.4%	0.896	0.824	0.820	2.2%
JT-VAE	5.30	4.93	4.49	100.0%	0.925	0.911	0.910	100.0%
GCPN	7.98	7.85	7.80	100.0%	0.948	0.947	0.946	100.0%

• You, Jiaxuan, et al. "Graph convolutional policy network for goal-directed molecular graph generation." Advances in neural information processing systems 2018.

Addition Algorithms

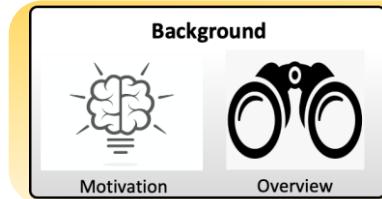
□ Unsupervised Rare Category Generation

- GraphGAN [Wang et al., 2018]
- Graph-Nets [Li et al., 2018]
- Graphite-VAE [Grover et al., 2019]
- Misc-GAN [Zhou et al., 2019]
- ...

□ Supervised Rare Category Generation

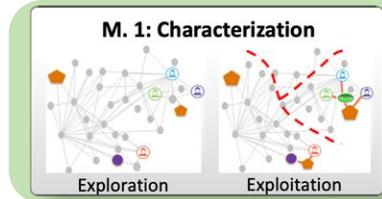
- Reg-VAE [Ma et al., 2018]
- MolecularRNN [Popova et al., 2019]
- NEVAE [Samanta et al., 2019]
- LF-MolGAN [Polsterl et al., 2019]
- LGGAN [Fang et al., 2019]
- CondGEN [Yang et al., 2019]
- ...

Roadmap



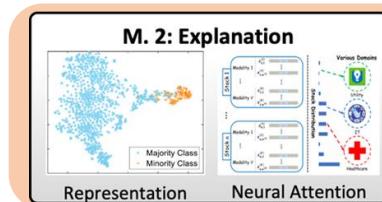
Background

- Motivation
- Research overview



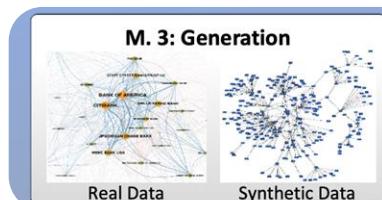
I. Rare Category Characterization

- Rare category characterization on homogeneous graphs
- Rare category characterization on heterogeneous graphs



II. Rare Category Explanation

- Data insights
- Model insights

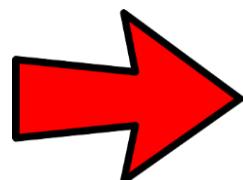


III. Rare Category Generation

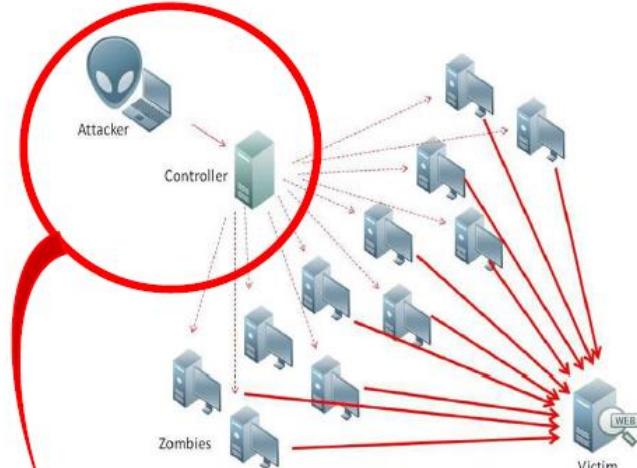
- Unsupervised rare category generation
- Supervised rare category generation



IV. Real-world Application



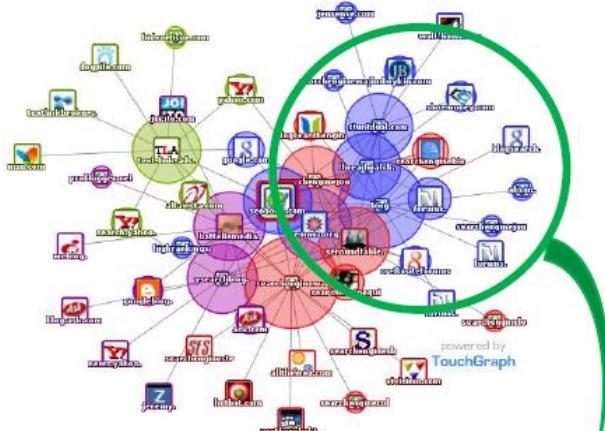
Real-World Applications



Computer Network



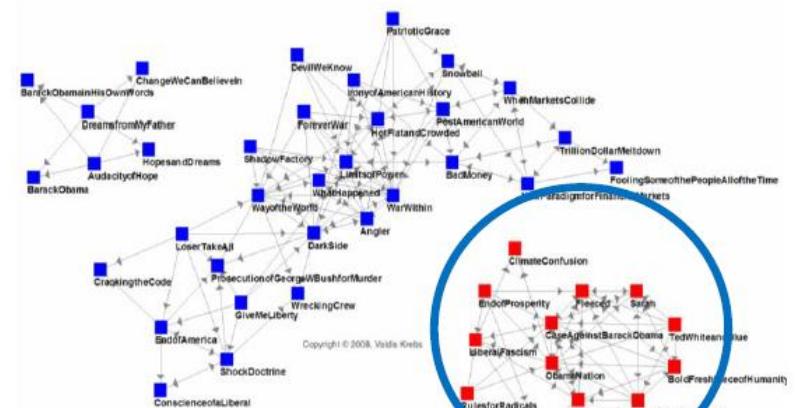
Network Intrusion



Social Network



Emerging Trends

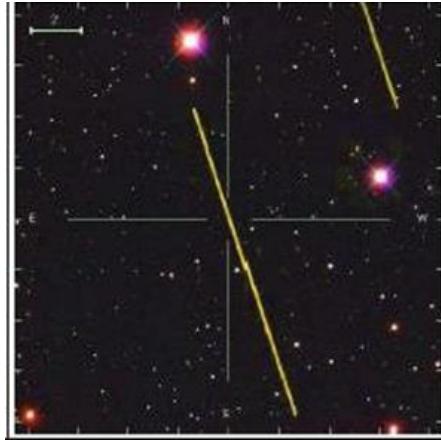


Online Transaction Network



Money Laundering

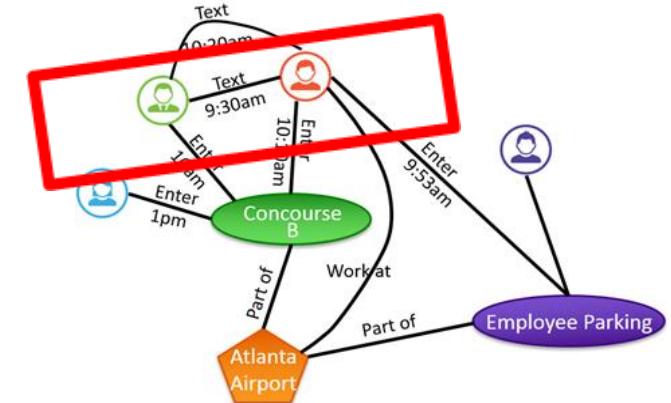
Real-World Applications



1. Out Space Discovery



2. Racial Hatred Discovery



3. Airport Insiders



4. Gene Disease



5. Identity Theft

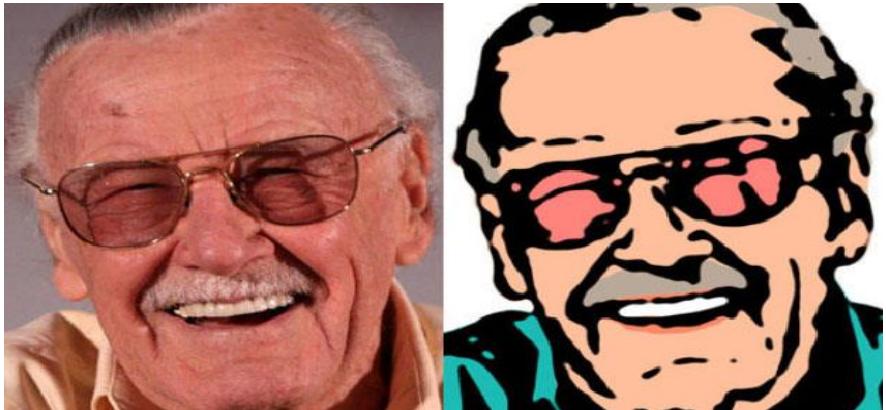


6. Lone Wolf Terrorism

Real-World Applications

□ Image style transfer

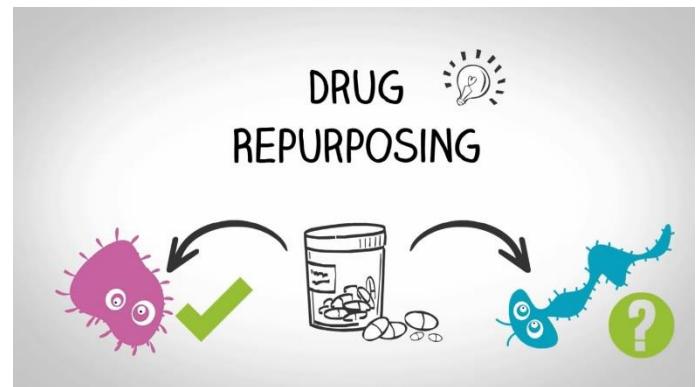
- GAN-based method
- Application: TikTok
- Real image → animation style



Style Transfer

□ Graph Generation and Drug Repurposing

- Identify the similar symptoms of two diseases
- Test a therapy approved for one disease for its effect on another
- Select drugs with likelihood of success



Drug Repurposing

• Zhu, Jun-Yan, et al. "Unpaired image-to-image translation using cycle-consistent adversarial networks." *Proceedings of the IEEE international conference on computer vision*. 2017.

- 110 - • Image source: <https://www.linkedin.com/pulse/drug-repurposing-hopes-hurdles-abirami-santhanam-phd/>

References

- Y. X. Wang, D. Ramanan, and M. Hebert. Learning to model the tail. In Advances in Neural Information Processing Systems, 2017.
- R. Spaaij. "The enigma of lone wolf terrorism: An assessment." Studies in Conflict & Terrorism, 2010.
- E. W. Ngai, Y. Hu, Y. H. Wong, Y. Chen, and X. Sun. The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature. Decision support systems, 2010.
- D. Pelleg, and A. W. Moore: Active Learning for Anomaly and Rare-Category Detection. NIPS, 2004.
- J. He. Rare category analysis (Doctoral dissertation, Carnegie Mellon University, School of Computer Science, Machine Learning Department), 2010.
- H. He, and A. ., Edwardo. "Learning from imbalanced data." IEEE Transactions on Knowledge & Data Engineering, 2008.
- N. V. Chawla, J. Nathalie, and K. Aleksander. "Special issue on learning from imbalanced data sets." ACM Sigkdd Explorations Newsletter, 2004
- Workshop on Learning from Imbalanced Data Sets, Co-located with AAAI 2000.
- Workshop on Learning with Imbalanced Domains: Theory and Applications, Co-located with ECML/PKDD 2018.
- Workshop on Active Learning for Decision-Making from Imbalanced Observational Data, Co-located with ICML 2019.

References

Part I

- J. He, and J. Carbonell. Co-Selection of Features and Instances for Unsupervised Rare Category Analysis. SDM 2010.
- J. Sun, H. Qu, D. Chakrabarti, and C. Faloutsos. Neighborhood formation and anomaly detection in bipartite graphs. ICDM 2005.
- H. Tong, and C. Lin. Non-Negative Residual Matrix Factorization with Application to Graph Anomaly Detection. SDM 2011.
- E. A. Manzoor, S. M. Milajerdi, and L. Akoglu: Fast Memory-efficient Anomaly Detection in Streaming Heterogeneous Graphs. KDD 2016.
- D. Zhou, J. He, K. S. Candan, and H. Davulcu. MUVR: Multi-View Rare Category Detection. IJCAI 2015.
- L. Akoglu, M. McGlohon, and C. Faloutsos. OddBall: Spotting Anomalies in Weighted Graphs. PAKDD, 2010.
- D. Zhou, S. Zhang, M. Y. Yildirim, S. Alcorn, H. Tong, H. Davulcu, and J. He: A Local Algorithm for Structure-Preserving Graph Cut. KDD 2017.
- S. Zhang, D. Zhou, M. Y. Yildirim, S. Alcorn, J. He, H. Davulcu, and H. Tong. HiDDen: Hierarchical Dense Subgraph Detection with Application to Financial Fraud Detection. SDM 2017.
- G. Pang, C. Shen, and A. van den Hengel. Deep anomaly detection with deviation networks. ACM SIGKDD 2019.
- K. Ding, J. Li, R. Bhanushali, and H. Liu. Deep anomaly detection on attributed networks. SDM 2019.

References

Part I

- V. Pavan, and W.-K. Wong. "Category detection using hierarchical mean shift." ACM SIGKDD 2009.
- K. Ding, Y. Li, J. Li, C. Liu, and H. Liu. Graph Neural Networks with High-order Feature Interactions. arXiv preprint arXiv:1908.07110.
- Z. Liu, W. Liu, P.Y. Chen, C. Zhuang, and C. Song. hpGAT: High-Order Proximity Informed Graph Attention Network. IEEE Access, 2019.
- C. Meng, S.C. Mouli, B. Ribeiro, and J. Neville. Subgraph pattern neural networks for high-order graph evolution prediction. AAAI 2018.
- A.G. Carranza, R.A. Rossi, A. Rao, and E. Koh. Higher-order spectral clustering for heterogeneous graphs. arXiv preprint arXiv:1810.02959, 2018.
- D. Zhou, J. He, H. Davulcu, and R. Maciejewski. Motif-Preserving Dynamic Local Graph Cut. BigData-2018 2018.
- E. T. Charalampos, P. Jakub, and M. Michael. Scalable Motif-aware Graph Clustering. WWW 2017.
- A.R. Benson, D.F. Gleich, and J. Leskovec. Tensor spectral clustering for partitioning higher-order network structures. SDM 2015.
- D.A. Spielman, and S.H. Teng. A local clustering algorithm for massive graphs and its application to nearly linear time graph partitioning. SIAM Journal on computing, 2013.

References

Part II

- D. Zhou, J. He, H. Yang, and W. Fan. SPARC: Self-Paced Network Representation for Few-Shot Rare Category Characterization. KDD 2018.
- J. Wu, J. He, and Y. Liu. ImVerde: Vertex-Diminished Random Walk for Learning Network Representation from Imbalanced Data. IEEE Bigdata, 2018.
- N. Liu, D. Shin, and X. Hu. Contextual Outlier Interpretation. IJCAI 2018.
- M., Meghanath, and L. Akoglu. "Explaining anomalies in groups with characterizing subspace rules." DMKD 2018.
- H. Lin, S. Gao, D. Gotz, F. Du, J. He, and N., Cao. Rclens: Interactive rare category exploration and identification. IEEE transactions on visualization and computer graphics, 2018.
- J. Pan, D. Han, F. Guo, D. Zhou, N. Cao, J. He, M. Xu, and W. Chen. "RCAnalyzer: visual analytics of rare categories in dynamic networks." Frontiers of Information Technology & Electronic Engineering, 2020.
- J. Kang, J. He, R. Maciejewski, and H. Tong. InFoRM: Individual Fairness on Graph Mining. ACM SIGKDD 2020.
- N. Gupta, D. Eswaran, N. Shah, L. Akoglu, and C. Faloutsos. Beyond outlier detection: Lookout for pictorial explanation. ECML 2018.
- S. Rajendran, W. Meert, V. Lenders, and S. Pollin. SAIFE: Unsupervised wireless spectrum anomaly detection with interpretable features. IEEE DySPAN 2018.

References

Part II

- J. Chen, Z. Xiu, B. Goldstein, R. Henao, L. Carin, and C. Tao. Supercharging Imbalanced Data Learning With Causal Representation Transfer. arXiv preprint arXiv:2011.12454, 2020.
- S.H. Dumpala, R. Chakraborty, S.K. Kopparapu, and T.C.S. Reseach. A Novel Data Representation for Effective Learning in Class Imbalanced Scenarios. In IJCAI 2018.
- Y. Zhao, Y. and M.K. Hryniewicki. XGBOD: improving supervised outlier detection with unsupervised representation learning. IJCNN2018.
- Z. Yang, W. Cohen, and R. Salakhudinov. Revisiting semi-supervised learning with graph embeddings. ICML 2018.
- C. Huang, Y. Li, C.C. Loy, and X. Tang. Learning deep representation for imbalanced classification. CVPR 2016.

References

- N. Chawla, K. Bowyer, L. Hall, and W. Kegelmeyer. "SMOTE: synthetic minority over-sampling technique." *Journal of artificial intelligence research*. 2002.
- J. You, R. Ying, X. Ren, W. Hamilton, and J. Leskovec. "Graphrnn: Generating realistic graphs with deep auto-regressive models." *ICML* 2018.
- A. Bojchevski, O. Shchur, D. Zügner, and S. Günnemann. "Netgan: Generating graphs via random walks." *ICML* 2018.
- D. Zhou, L. Zheng, J. Han, and J. He. "A Data-Driven Graph Generative Model for Temporal Interaction Networks." *ACM SIGKDD*. 2020.
- M. Simonovsky, and N. Komodakis. "Graphvae: Towards generation of small graphs using variational autoencoders." In *International Conference on Artificial Neural Networks*, pp. 412-422. Springer, Cham, 2018.
- W. Jin, R. Barzilay, and T. Jaakkola. "Junction tree variational autoencoder for molecular graph generation." *ICML* 2018.
- N. Cao, and T. Kipf. "MolGAN: An implicit generative model for small molecular graphs." *ICML* 2018.
- J. You, B. Liu, Z. Ying, V. Pande, and J. Leskovec. "Graph convolutional policy network for goal-directed molecular graph generation." In *Advances in neural information processing systems* 2018.
- H. Wang, J. Wang, J. Wang, M. Zhao, W. Zhang, F. Zhang, X. Xie, and M. Guo, "Graphgan: Graph representation learning with generative adversarial nets." *AAAI* 2018.
- D. Zhou, L. Zheng, J. Xu, and J. He, "Misc-GAN: A Multi-scale Generative Model for Graphs." *Frontiers in Big Data*, 2019.

References

Part III

- Y. Li, O. Vinyals, C. Dyer, R. Pascanu, and P. Battaglia, “Learning deep generative models of graphs.” arXiv preprint arXiv:1803.03324. 2018
- A. Grover, A. Zweig, and S. Ermon, “Graphite: Iterative generative modeling of graphs.” In International conference on machine learning (pp. 2434-2444). PMLR 2019.
- C. Yang, P. Zhuang, W. Shi, A. Luu, and P. Li, “Conditional structure generation through graph variational generative adversarial nets.” In Advances in Neural Information Processing Systems 2019.
- S. Fan, and B. Huang, “Labeled graph generative adversarial networks.” arXiv preprint arXiv:1906.03220. 2019.
- T. Ma, J. Chen, and C. Xiao, “Constrained generation of semantically valid graphs via regularizing variational autoencoders.” In Advances in Neural Information Processing Systems 2018.
- M. Popova, M. Shvets, J. Oliva, and O. Isayev, “MolecularRNN: Generating realistic molecular graphs with optimized properties.” arXiv preprint arXiv:1905.13372. 2019.
- S. Pölsterl, and C. Wachinger, “Likelihood-Free Inference and Generation of Molecular Graphs.” arXiv preprint arXiv:1905.10310. 2019
- B. Samanta, A. De, G. Jana, V. Gómez, P. Chattaraj, N. Ganguly, and M. Gomez-Rodriguez, “Nevae: A deep generative model for molecular graphs.” Journal of Machine Learning Research, 2020.

